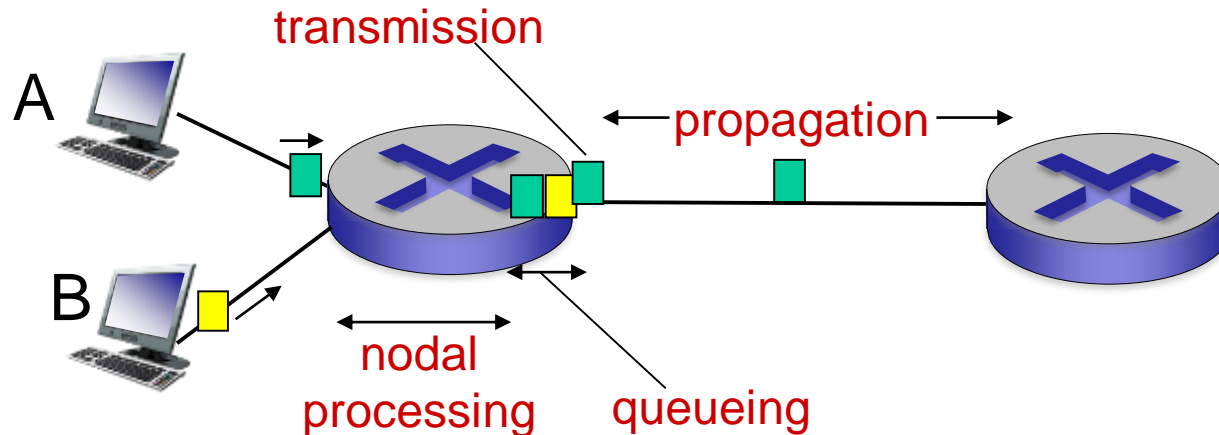


# **CS 3800: Computer Networks**

**Final Exam Review**

Instructor: John Korah

# Four sources of packet delay



$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$

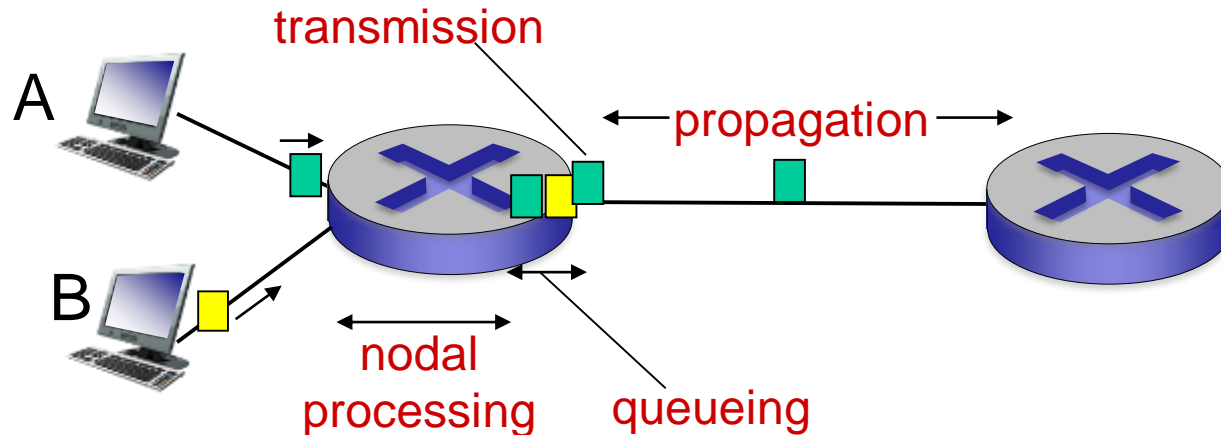
**$d_{\text{proc}}$ : nodal processing**

- check bit errors
- determine output link
- typically < msec

**$d_{\text{queue}}$ : queueing delay**

- time waiting at output link for transmission
- depends on congestion level of router

# Four sources of packet delay



$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$

$d_{\text{trans}}$ : transmission delay:

- $L$ : packet length (bits)
- $R$ : link bandwidth (bps)
- $d_{\text{trans}} = L/R$

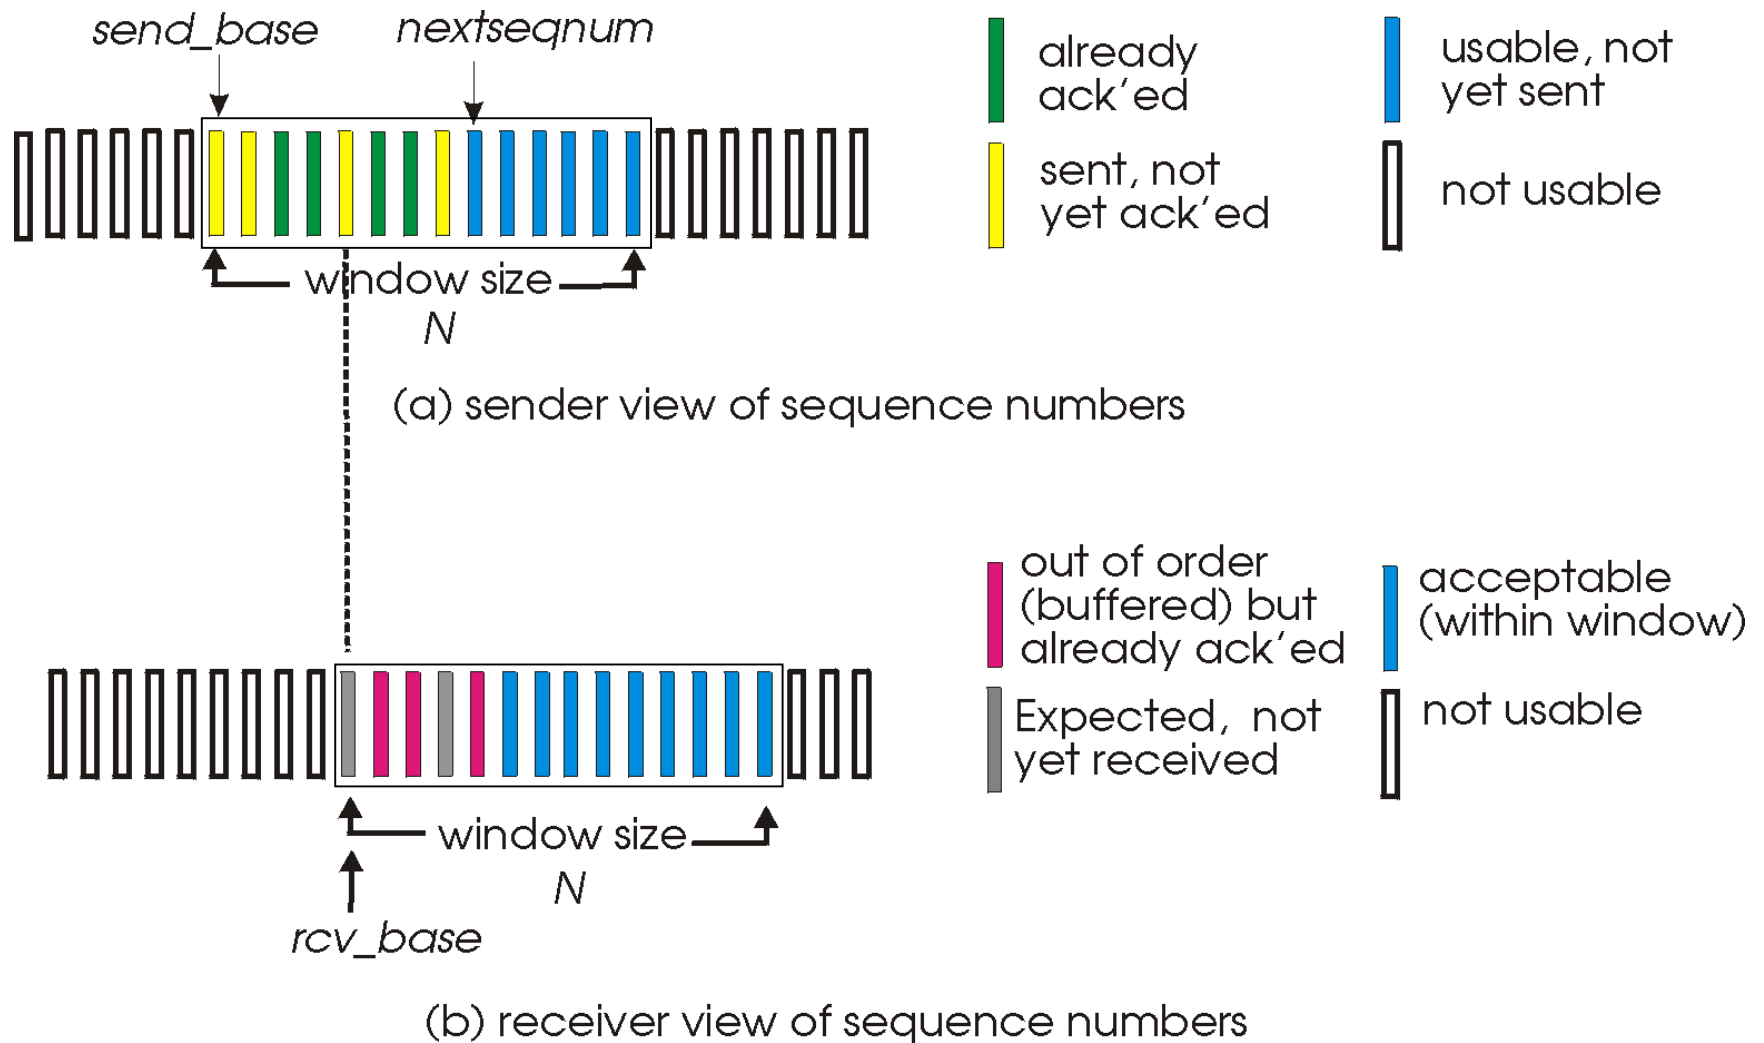
$d_{\text{prop}}$ : propagation delay:

- $d$ : length of physical link
- $s$ : propagation speed ( $\sim 2 \times 10^8$  m/sec)
- $d_{\text{prop}} = d/s$

←  $d_{\text{trans}}$  and  $d_{\text{prop}}$  →  
very different

- Check out the Java applet for an interactive animation on trans vs. prop delay [here](#)

# Selective repeat: sender, receiver windows



# Selective repeat

## sender

### data from above:

- if next available seq # in window, send pkt

### timeout(n):

- resend pkt n, restart timer

### ACK(n) in [sendbase, sendbase+N]:

- mark pkt n as received
- if n smallest unACKed pkt, advance window base to next unACKed seq #

## receiver

### pkt n in [rcvbase, rcvbase+N-1]

- send ACK(n)
- out-of-order: buffer
- in-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt

### pkt n in [rcvbase-N, rcvbase-1]

- ACK(n)

### otherwise:

- ignore

# TCP seq. numbers, ACKs

## sequence numbers:

- byte stream “number” of first byte in segment’s data

## acknowledgements:

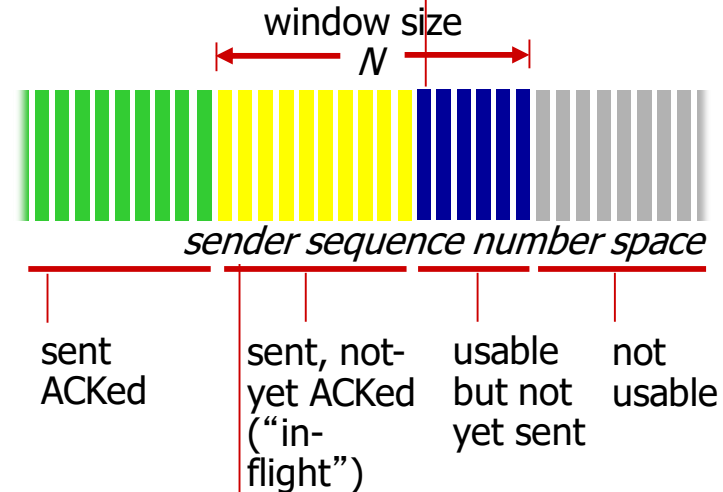
- seq # of next byte expected from other side
- cumulative ACK

**Q:** how receiver handles out-of-order segments

- **A:** TCP spec doesn’t say,  
- up to implementor

outgoing segment from sender

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer

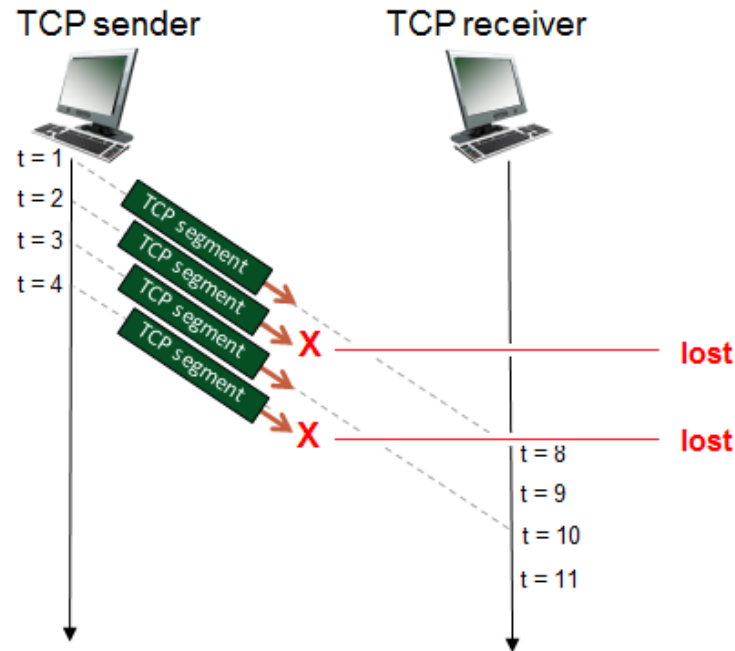


incoming segment to sender

source port #	dest port #
sequence number	
acknowledgement number	
	A
checksum	urg pointer

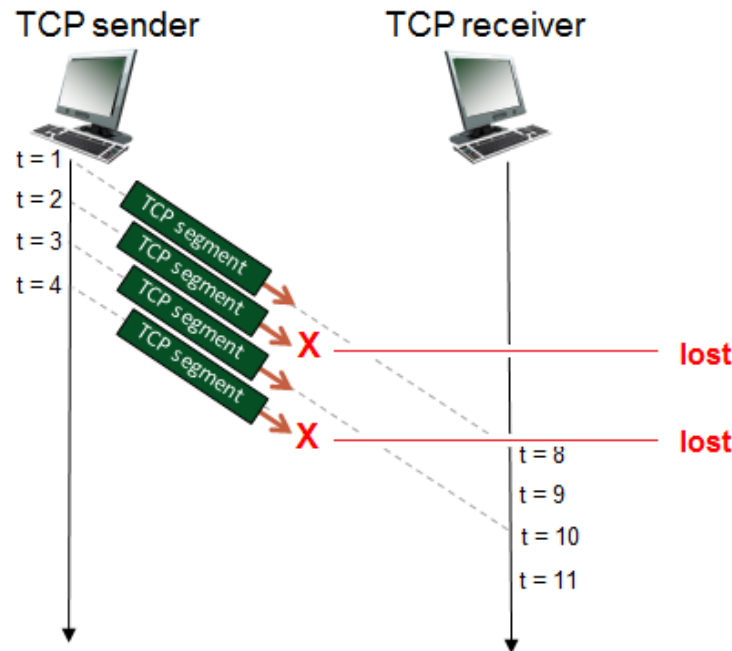
# Problem

- Consider the figure below in which TCP a sender and receiver communicate over a connection in which the sender-to-receiver segments may be lost.



- The TCP sender sends initial window of four segments at  $t=1, 2, 3, 4$ , respectively. Suppose the initial value of the sender-to-receiver sequence number is 123 and the first four segments *each* contain 575 bytes. The delay between the sender and the receiver is 7 time units, and so the first segment arrives at the receiver at  $t=8$ . As shown in the figure, two of the four segment(s) are lost between the sender and the receiver.

# Problem




- Consider the figure below in which TCP a sender and receiver communicate over a connection in which the sender-to-receiver segments may be lost.
- Answer the following questions:
  - Give the sequence numbers associated with each of the four segments sent by the sender
  - List the sequence of acknowledgements transmitted by the TCP receiver in response to the receipt of the segments actually received. In particular, give the value in the acknowledgement field of each receiver-to-sender acknowledgement, and give a brief explanation as to why that particular acknowledgement number value is being used



# TCP round trip time, timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- exponential weighted moving average
- influence of past sample decreases exponentially fast
- typical value:  $\alpha = 0.125$



# TCP round trip time, timeout

- **timeout interval:** `EstimatedRTT` plus “safety margin”
  - large variation in `EstimatedRTT` → larger safety margin
- estimate `SampleRTT` deviation from `EstimatedRTT`:

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



↑  
estimated RTT

↑  
“safety margin”

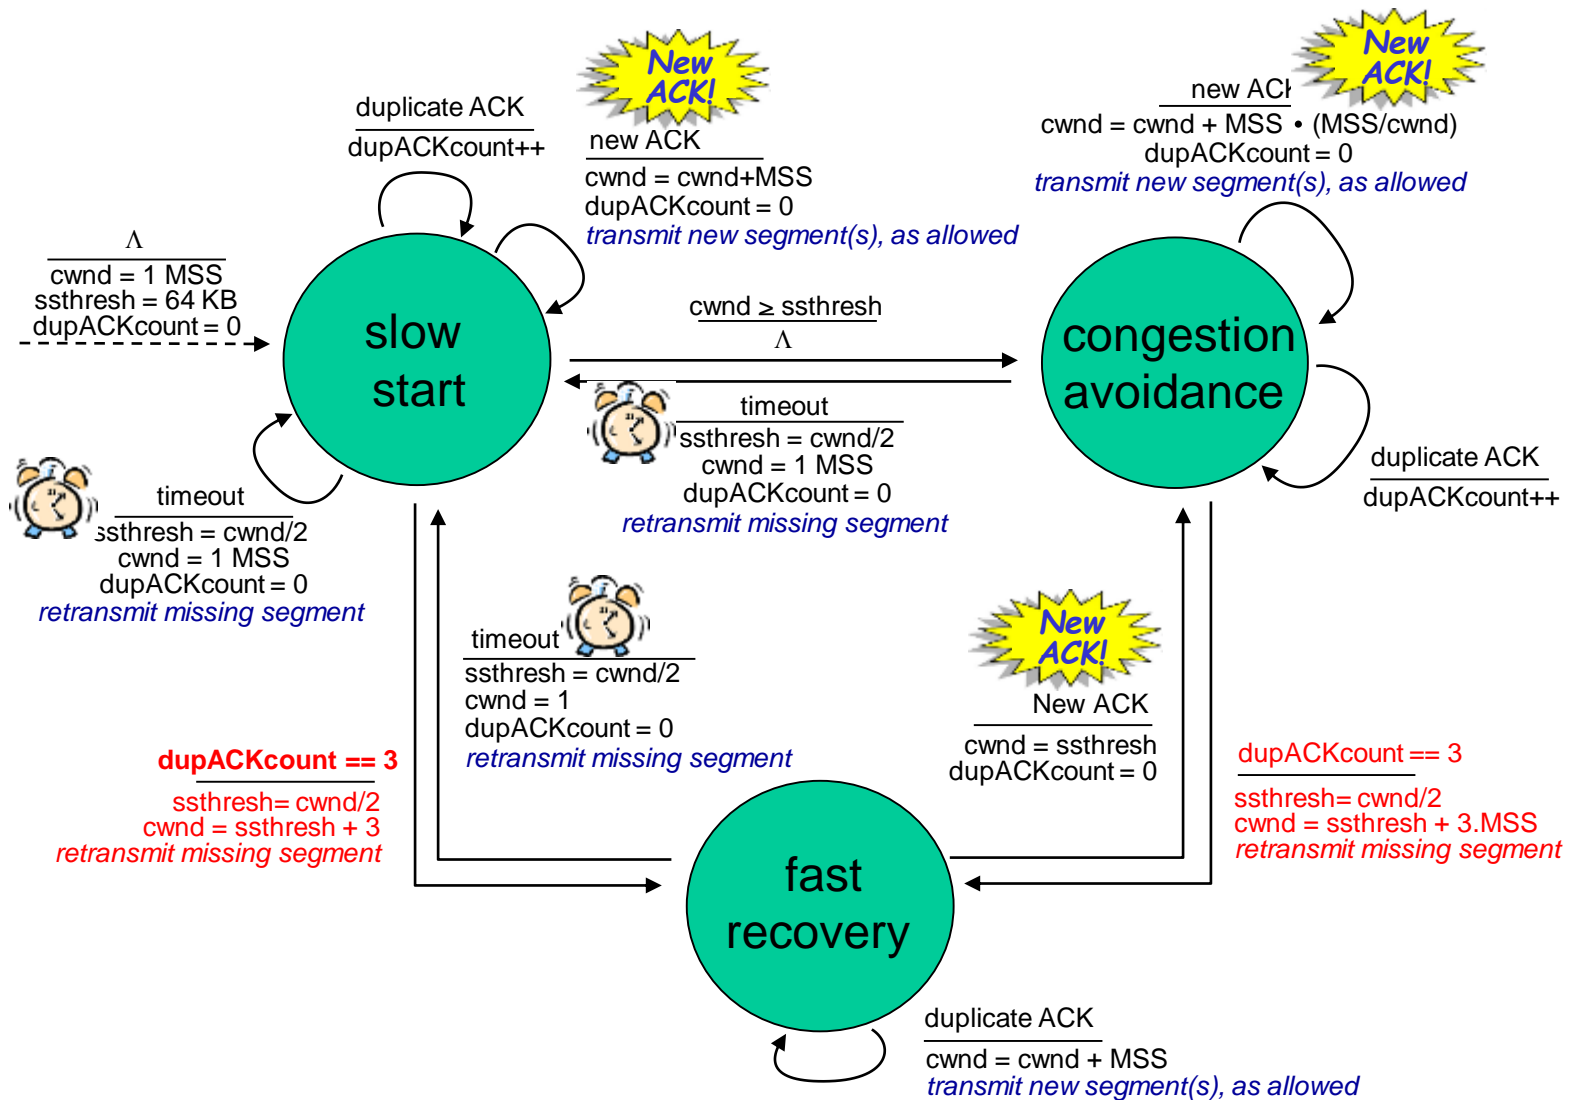
# TCP round trip time, timeout




Suppose that TCP's current estimated values for the round trip time ( $\text{estimatedRTT}$ ) and deviation in the RTT ( $\text{DevRTT}$ ) are 360 msec and 39 msec, respectively. Suppose that the next three measured values of the RTT are 260, 340, and 260 respectively.

Compute TCP's new value of  $\text{estimatedRTT}$ ,  $\text{DevRTT}$ , and the TCP timeout value after each of these three measured RTT values is obtained. Use the values of  $\alpha = 0.125$  and  $\beta = 0.25$ .

# Recap: TCP Congestion Control



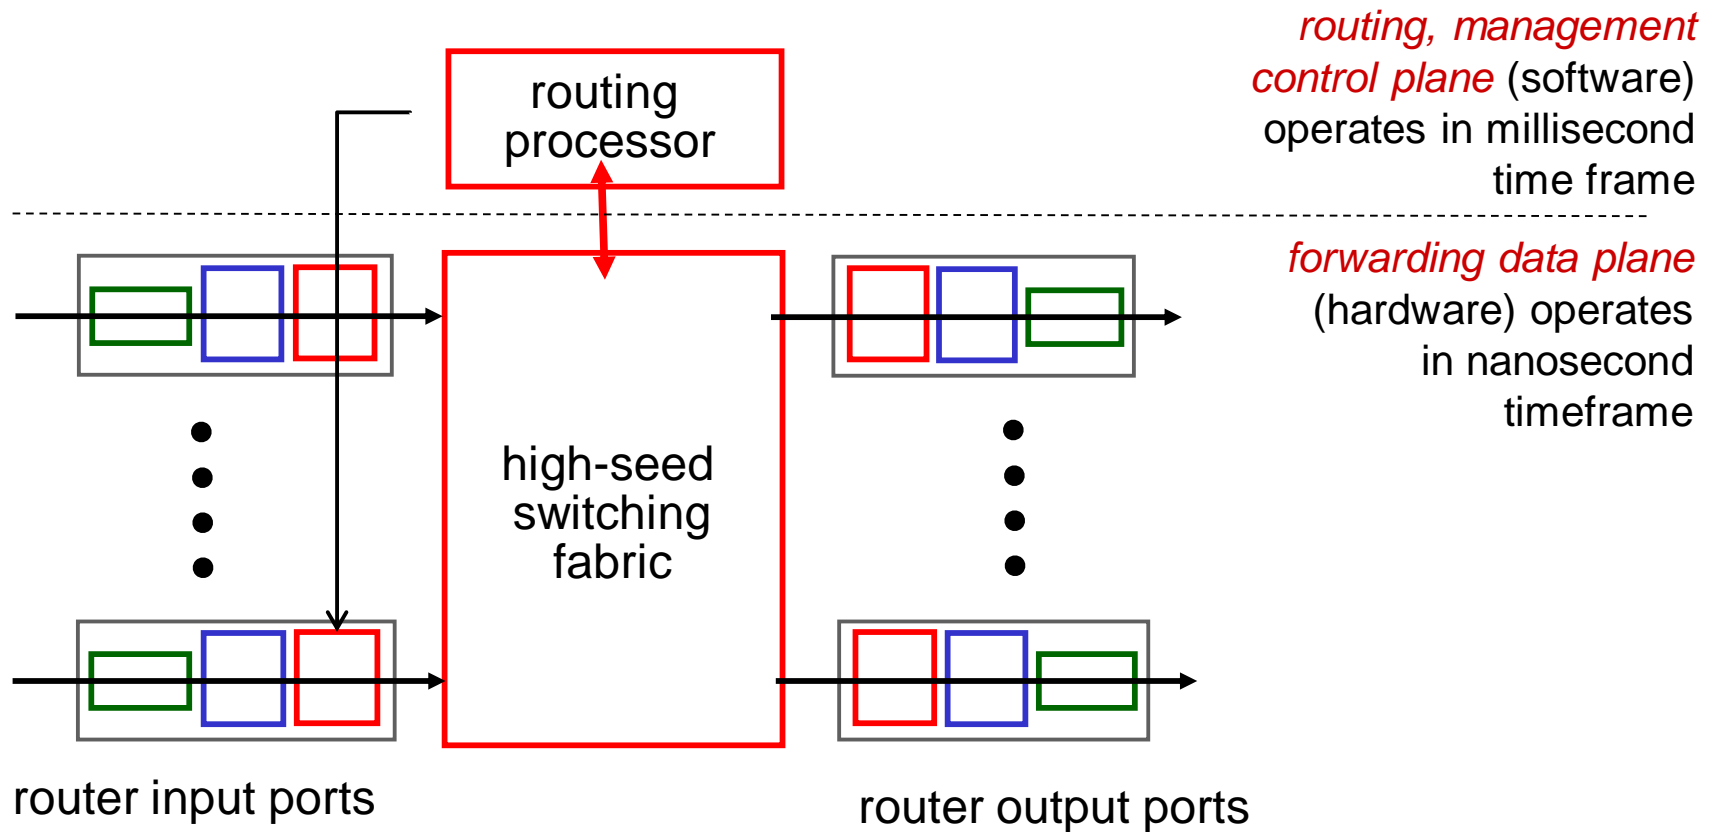


The result of sending that flight of packets is that either (i) all packets are ACKed at the end of the time unit, (ii) there is a timeout for the first packet, or (iii) there is a triple duplicate ACK for the first packet.

- Give the times at which TCP is in slow start, congestion avoidance and fast recovery at the start of a time slot, when the flight of packets is sent.
- Give the times at which the first packet in the sent flight of packets is lost, and indicate whether that packet loss is detected via timeout, or by triple duplicate ACKs

# Router architecture overview

- high-level view of generic router architecture:



# Longest prefix matching

## *longest prefix matching*

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

examples:

DA: 11001000 00010111 00010110 10100001


which interface?

DA: 11001000 00010111 00011000 10101010

which interface?

# IP fragmentation, reassembly

- network links have MTU (max.transfer size) - largest possible link-level frame
  - different link types, different MTUs
- large IP datagram divided (“fragmented”) within net
  - one datagram becomes several datagrams
  - **“reassembled” only at final destination**
  - IP header bits used to identify, order related fragments





# IP fragmentation, reassembly

## example:

- ❖ 4000 byte datagram
- ❖ MTU = 1500 bytes

	length	ID	fragflag	offset
	=4000	=x	=0	=0

*one large datagram becomes  
several smaller datagrams*

1480 bytes in  
data field


offset =  
 $1480/8$

	length	ID	fragflag	offset
	=1500	=x	=1	=0

	length	ID	fragflag	offset
	=1500	=x	=1	=185

	length	ID	fragflag	offset
	=1040	=x	=0	=370


# IP Address Classes



# IP addressing: CIDR

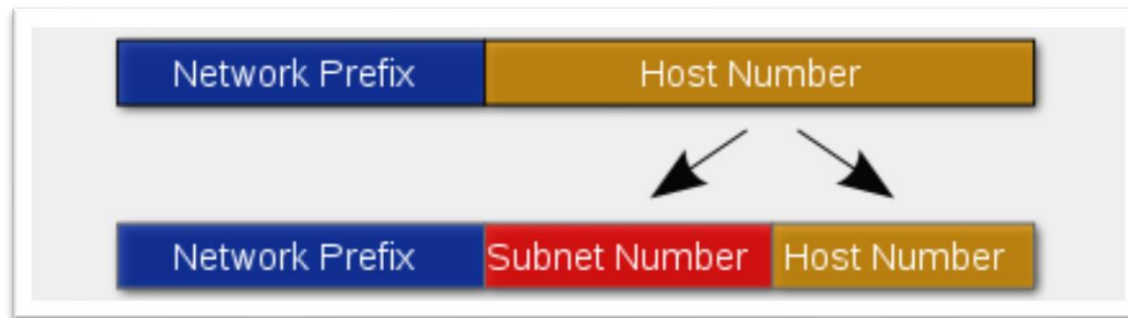
## CIDR: Classless InterDomain Routing

- subnet portion of address of arbitrary length
- address format: **a.b.c.d/x**, where x is # bits in subnet portion of address



# Subnetted Networks

- The network portion of the address is extended by splitting up the host number



- Borrowing 1 or more bits from the host bit portion

# Class C Subnetting

# of Subnets	# of Hosts/Subnet	NetMask	4 <sup>th</sup> Octet	CIDR Notation
2	126	255.255.255.128	10000000	/25
4	62	255.255.255.192	11000000	/26
8	30	255.255.255.224	11100000	/27
16	14	255.255.255.240	11110000	/28
32	6	255.255.255.248	11111000	/29
64	2	255.255.255.252	11111100	/30

# DHCP client-server scenario

DHCP server: 223.1.2.5

DHCP discover

arriving  
client



Broadcast: is there a  
DHCP server out there?

DHCP offer

Broadcast: I'm a DHCP  
server! Here's an IP  
address you can use  
-----


DHCP request

Broadcast: OK. I'll take  
that IP address!

DHCP ACK

Broadcast: OK. You've  
got that IP address!

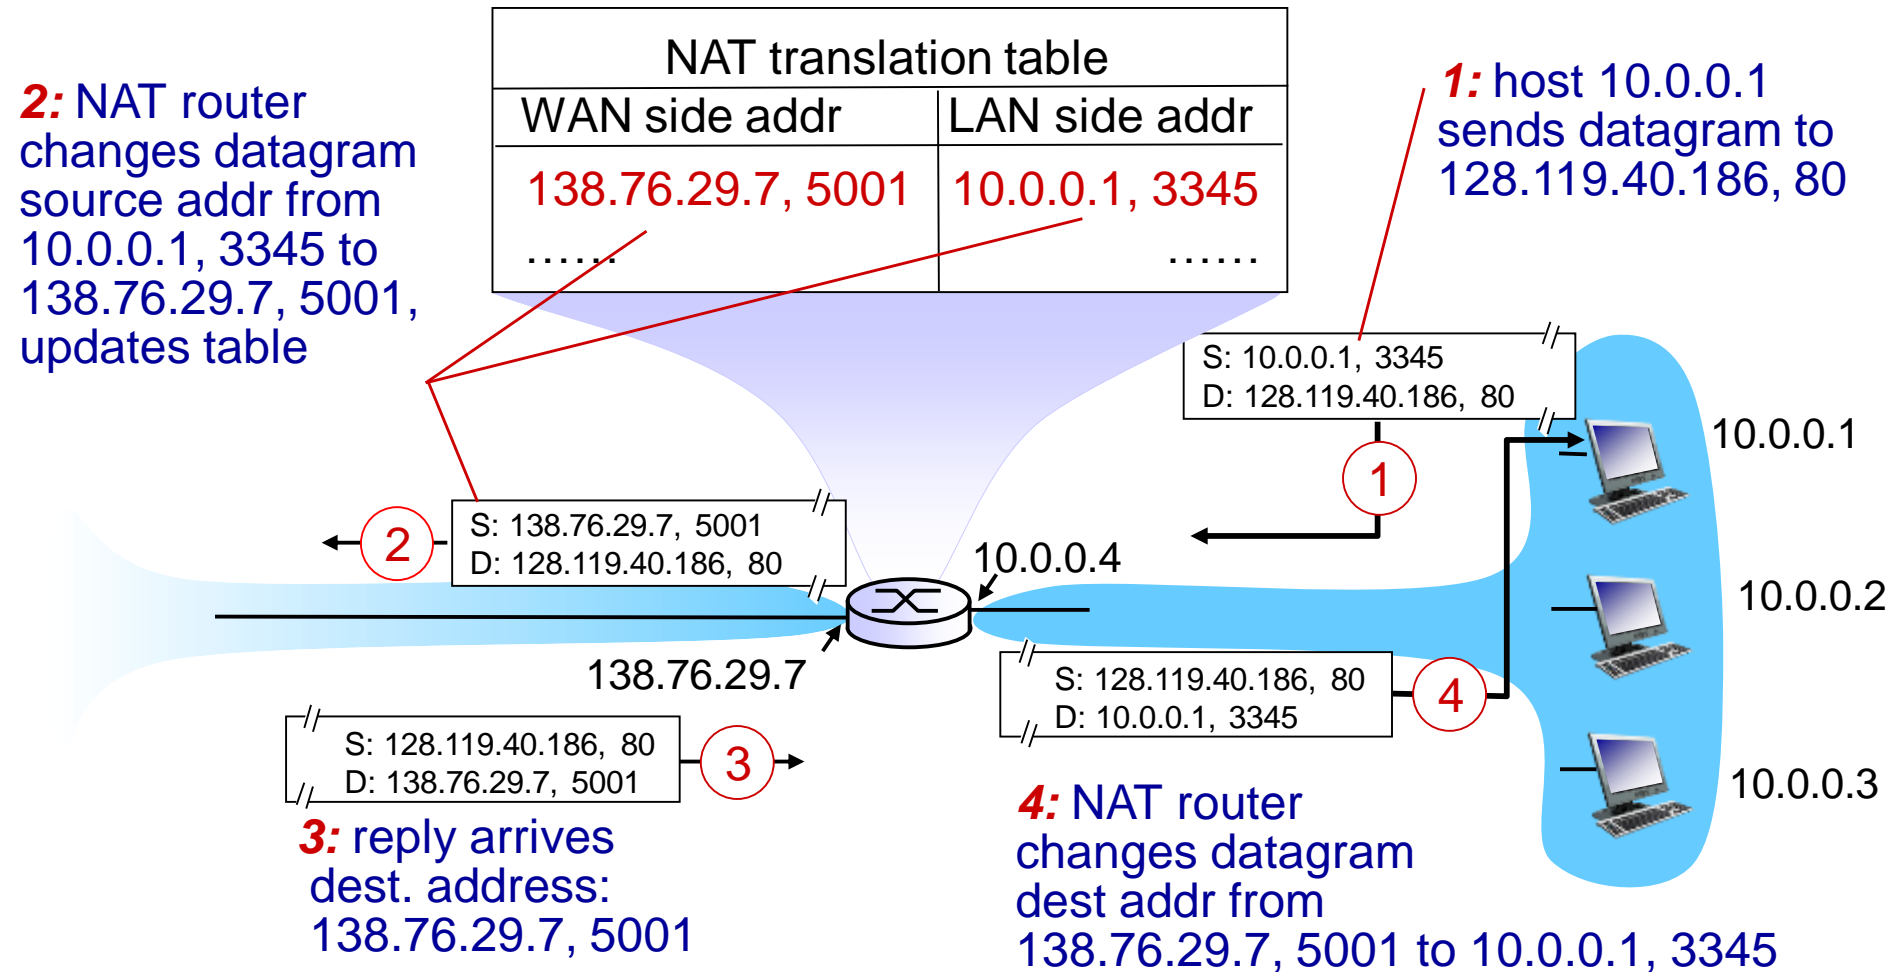
# NAT: network address translation



*all* datagrams *leaving* local network have *same* single source NAT IP address: 138.76.29.7, different source port numbers

datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)


# NAT: network address translation






# Tunneling

logical view:




physical view:





# Tunneling

logical view:




physical view:



# Address Resolution Protocol (ARP)

**Question:** how to determine interface's MAC address, knowing its IP address?



**ARP table:** each IP node (host, router) on LAN has table

- IP/MAC address mappings for some LAN nodes:  
< IP address; MAC address; TTL >
- TTL (Time To Live): time after which address mapping will be forgotten (typically 20 min)


# ARP protocol: same LAN

- A wants to send datagram to B
  - B's MAC address not in A's ARP table.
- A **broadcasts** ARP query packet, containing B's IP address
  - destination MAC address = FF-FF-FF-FF-FF-FF
  - all nodes on LAN receive ARP query
- B receives ARP packet, replies to A with its (B's) MAC address
  - frame sent to A's MAC address (unicast)
- A caches (saves) IP-to-MAC address pair in its ARP table until information becomes old (times out)
  - soft state: information that times out (goes away) unless refreshed
- ARP is “plug-and-play”:
  - nodes create their ARP tables *without intervention from net administrator*

# Addressing: routing to another LAN

walkthrough: **send datagram from A to B via R**

- focus on addressing – at IP (datagram) and MAC layer (frame)
- assume A knows B's IP address
- assume A knows IP address of first hop router, R (how?)
- assume A knows R's MAC address (how?)



# CSMA (carrier sense multiple access)


**CSMA:** listen before transmit:

if channel sensed idle: transmit entire frame

- if channel sensed busy, defer transmission
- human analogy: don't interrupt others!

# CSMA collisions

- collisions *can* still occur:  
propagation delay means  
two nodes may not hear  
each other's  
transmission
- collision: entire packet  
transmission time  
wasted
  - distance &  
propagation delay  
play role in in  
determining collision  
probability




# CSMA/CD (collision detection)

**CSMA/CD:** carrier sensing, deferral as in CSMA


- collisions *detected* within short time
- colliding transmissions aborted, reducing channel wastage
- collision detection:
  - easy in wired LANs: measure signal strengths, compare transmitted, received signals
  - difficult in wireless LANs: received signal strength overwhelmed by local transmission strength
- human analogy: the polite conversationalist



# CSMA/CD (collision detection)




# 802.11 LAN architecture




- wireless host communicates with base station
  - **base station = access point (AP)**
- **Basic Service Set (BSS)** (aka “cell”) in infrastructure mode contains:
  - wireless hosts
  - access point (AP): base station
  - ad hoc mode: hosts only

# 802.11: passive/active scanning



## passive scanning:

- (1) beacon frames sent from APs
- (2) association Request frame sent: H1 to selected AP
- (3) association Response frame sent from selected AP to H1



## active scanning:

- (1) Probe Request frame broadcast from H1
- (2) Probe Response frames sent from APs
- (3) Association Request frame sent: H1 to selected AP
- (4) Association Response frame sent from selected AP to H1


# IEEE 802.11 MAC Protocol: CSMA/CA

## 802.11 sender

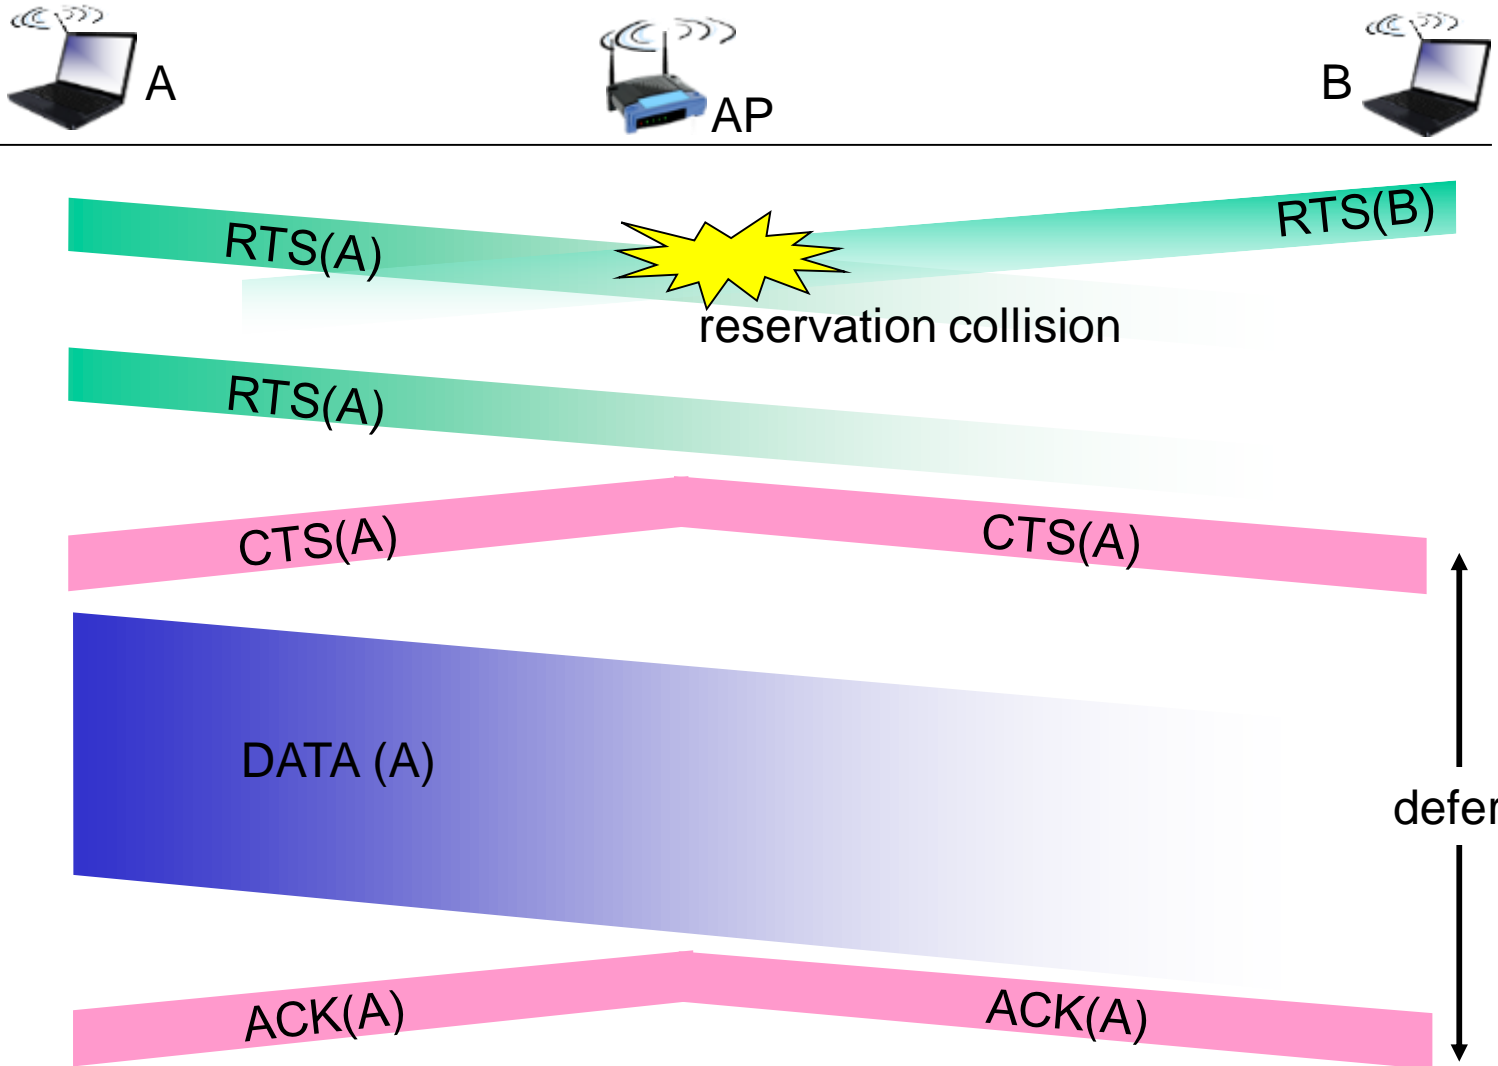
- 1 if sense channel idle for Distributed Inter-frame Space (**DIFS**) then  
transmit entire frame (no CD)
- 2 if sense channel busy then  
start random backoff time  
timer counts down while channel idle  
transmit when timer expires  
if no ACK, increase random backoff interval,  
repeat 2

## 802.11 receiver

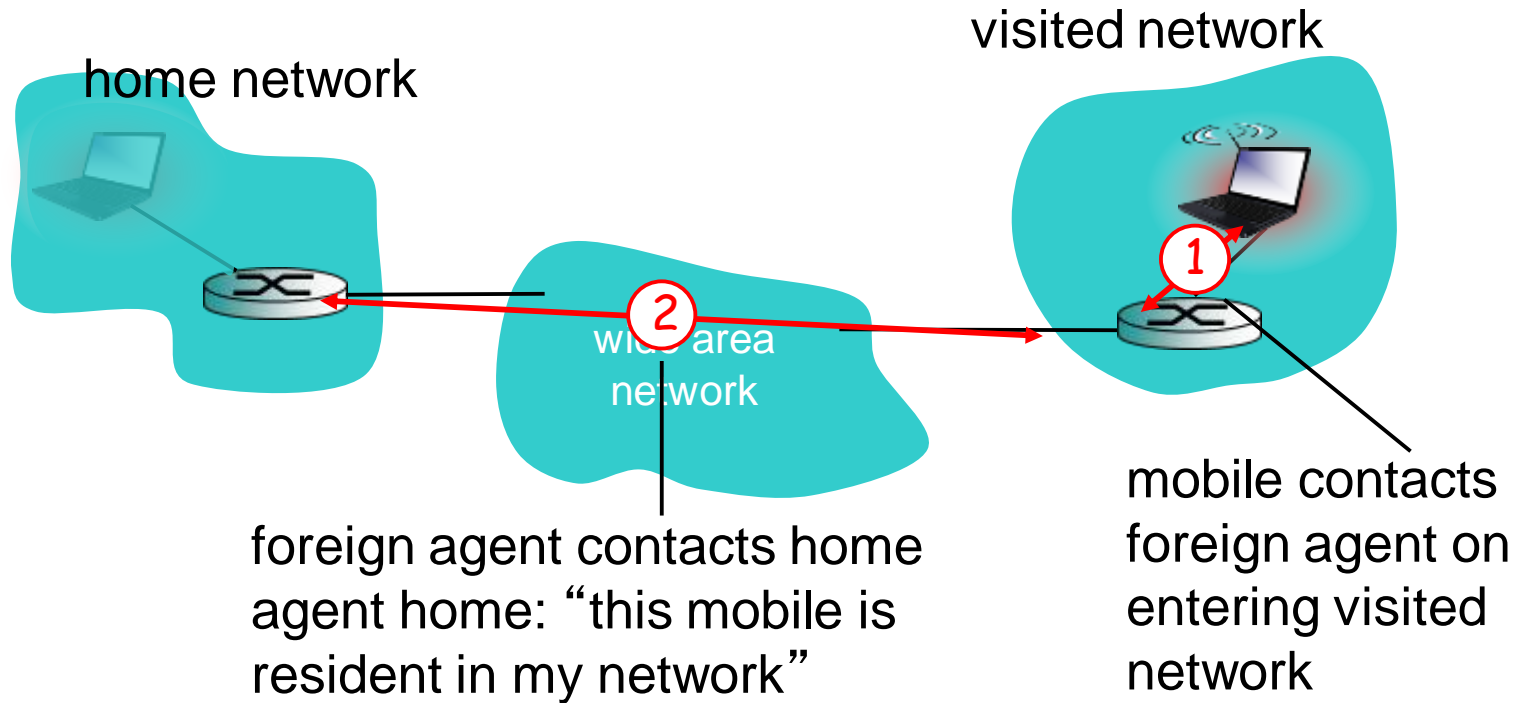
- if frame received OK  
return ACK after Short Inter-frame Spacing (**SIFS**) (ACK needed due to hidden terminal problem)



# Collision Avoidance: RTS-CTS exchange



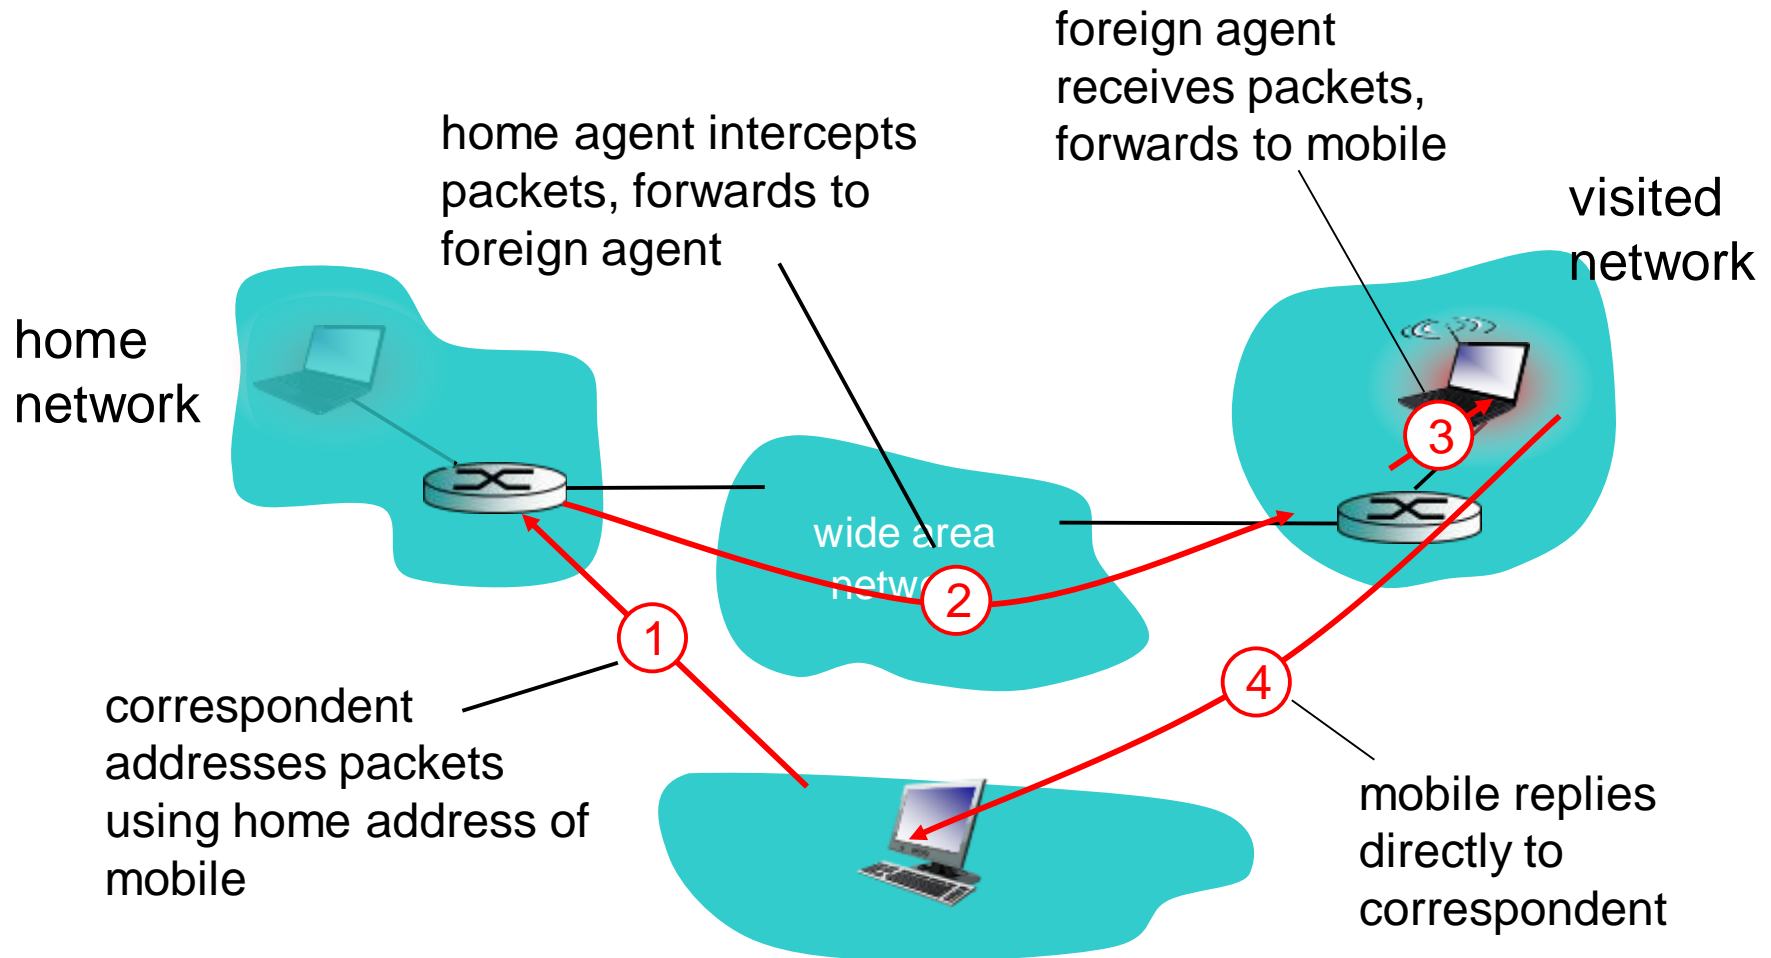
# Mobility: registration



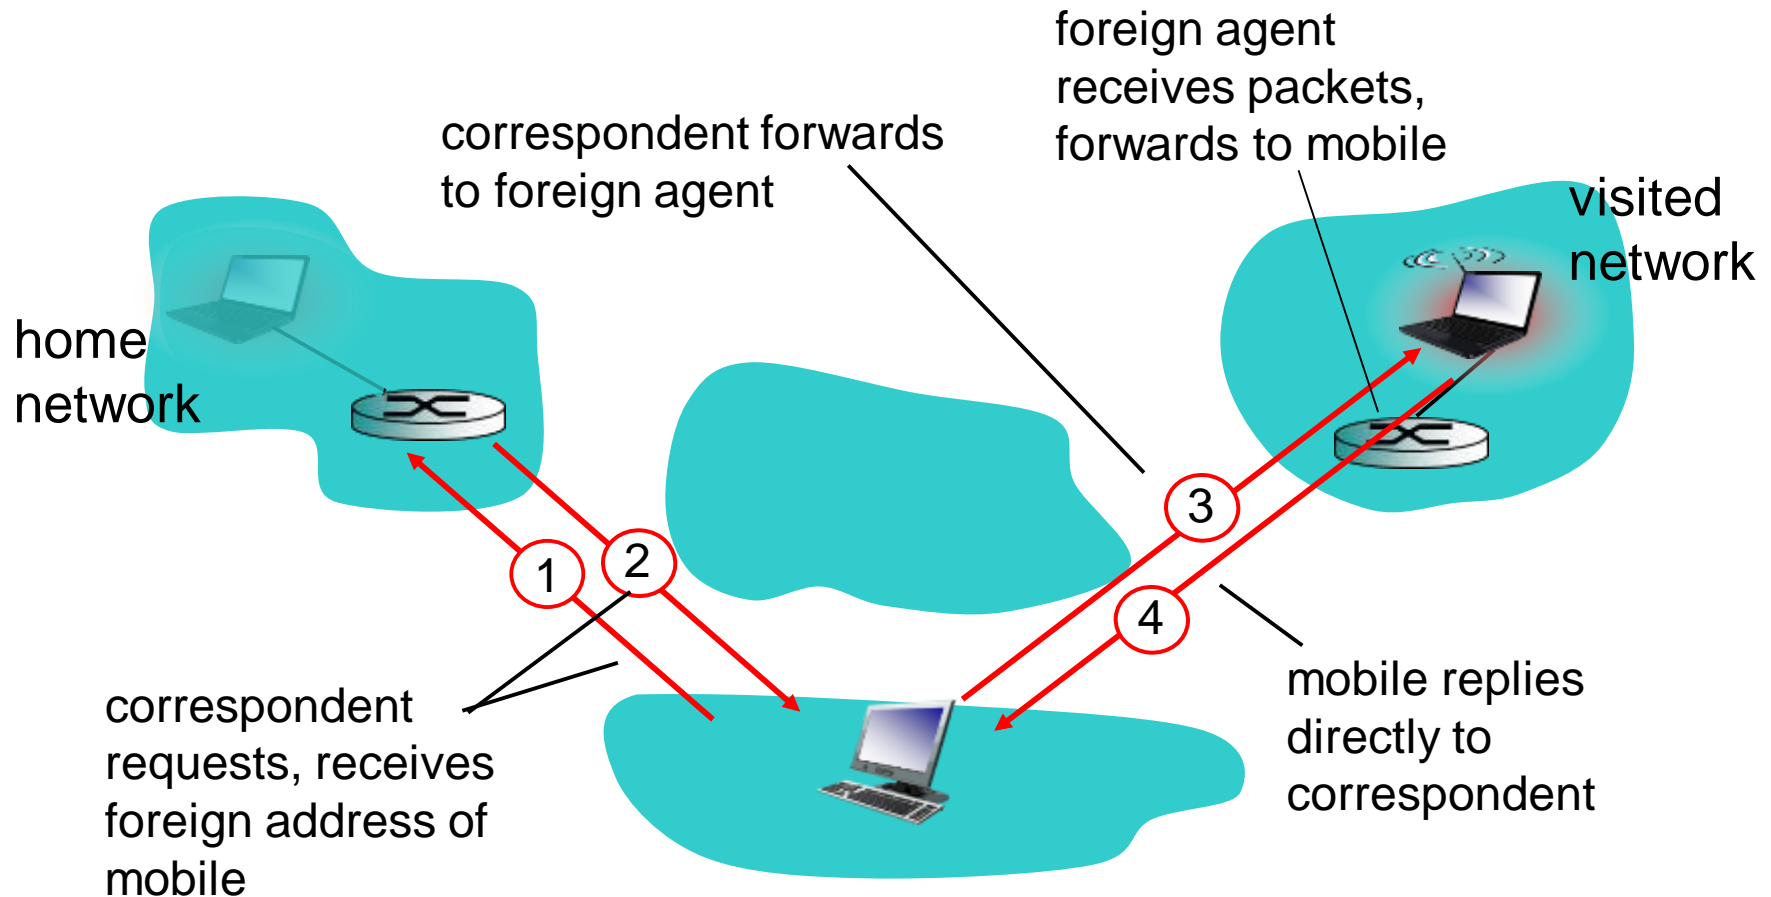
end result:

- foreign agent knows about mobile
- home agent knows location of mobile

# Mobility via indirect routing




# Mobility via direct routing






# Dijkstra's algorithm: another example




# Dijkstra's algorithm: another example

Step	N'	D(v),p(v)	D(w),p(w)	D(x),p(x)	D(y),p(y)	D(z),p(z)
0						
1						
2						
3						
4						
5						



# Dijkstra's algorithm: example (2)

resulting shortest-path tree from u:



resulting forwarding table in u:

destination	link
v	
x	
y	
w	
z	

# Distance vector algorithm

## *iterative, asynchronous:*

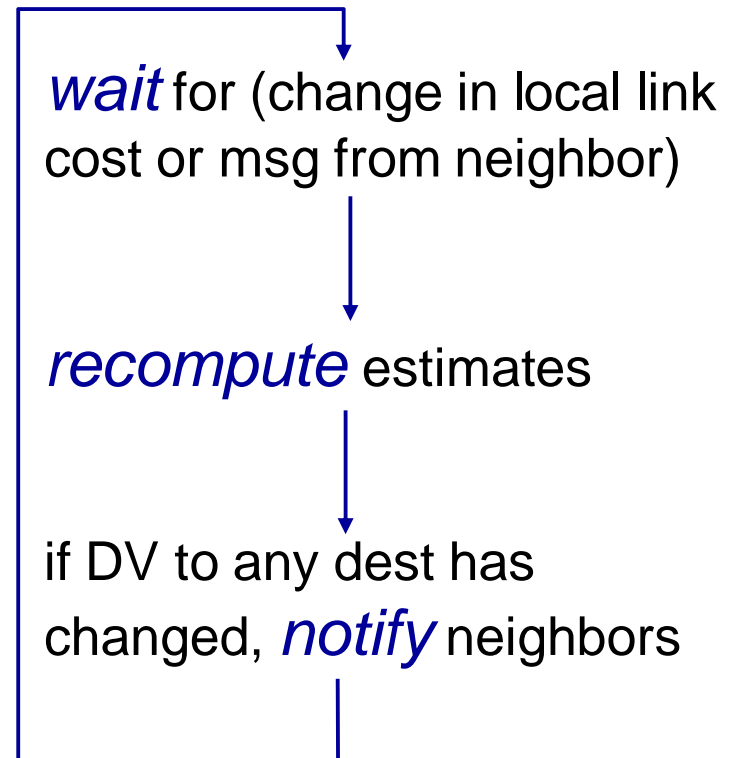
each local iteration  
caused by:

- local link cost change
- DV update message from neighbor

## *distributed:*

- each node notifies neighbors *only* when its DV changes
  - neighbors then notify their neighbors if necessary

## *each node:*



# Distance Vector algorithm

Let

$d_x(y) :=$  cost of least-cost path from  $x$  to  $y$

then

$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$


$v$   
|  
cost to neighbor  $v$


cost from neighbor  $v$  to destination  $y$

$\min$  taken over all neighbors  $v$  of  $x$

*Also called the Bellman-Ford equation (dynamic programming)*

# DVR Example





U	U	V	W	X	Y
U					
V					

U	U	V	W	X	Y
U					
V					

V	U	V	W	X	Y
U					
V					
W					
X					

V	U	V	W	X	Y
U					
V					
W					
X					

W	U	V	W	X	Y
V					
W					
X					


W	U	V	W	X	Y
V					
W					
X					

X	U	V	W	X	Y
V					
W					
X					
Y					

X	U	V	W	X	Y
V					
W					
X					
Y					

Y	U	V	W	X	Y
X					
Y					

Y	U	V	W	X	Y
X					
Y					



U	U	V	W	X	Y
U					
V					

U	U	V	W	X	Y
U					
V					

V	U	V	W	X	Y
U					
V					
W					
X					

V	U	V	W	X	Y
U					
V					
W					
X					

W	U	V	W	X	Y
V					
W					
X					

W	U	V	W	X	Y
V					
W					
X					

X	U	V	W	X	Y
V					
W					
X					
Y					

X	U	V	W	X	Y
V					
W					
X					
Y					

Y	U	V	W	X	Y
X					
Y					

Y	U	V	W	X	Y
X					
Y					



# Comparison of LS and DV algorithms

## *Communication costs*

- **LS:** broadcast link info to all
- **DV:** exchange between neighbors only
  - convergence time varies

*robustness:* what happens if router malfunctions?

## *LS:*

- node can advertise incorrect *link* cost
- each node computes only its own table

## *DV:*

- DV node can advertise incorrect *path* cost
- each node's table used by others
- error propagate thru network