

CS 3800: Computer Networks

Lecture 3: Application Layer

Instructor: John Korah

Acknowledgement

- The following slides include material from author resources for:
 - KR Text book
 - “Data and computer communications,” William Stallings, Tenth edition

Learning Goals


- Conceptual, implementation aspects of network application protocols
 - transport-layer service models
 - client-server paradigm
- learn about protocols by examining popular application-level protocols
 - HTTP
- Creating network applications
 - socket API

Topics

- **Application Layer**
 - Network Applications
- Web and HTTP
- Socket Programming: Introduction

Recall: Application Layer

- *Responsibilities*
 - Exchange information between hosts
 - Data unit is called **message**
- **Examples**
 - HTTP
 - DNS
 - SMTP
 - FTP
 - BitTorrent



Application Layer

- Application Layer protocol defines:
 - types of messages exchanged
 - e.g., request, response
 - message syntax
 - what fields in messages
 - how fields are delineated
 - message semantics
 - meaning of information in fields
 - rules for when and how processes send & respond to messages

open protocols:

- defined in RFCs
- allows for interoperability
- e.g., HTTP, SMTP

proprietary protocols:

- e.g., Skype

Common Network Apps

- e-mail
- web
- text messaging
- remote login
- P2P file sharing
- multi-user network games
- streaming stored video (YouTube, Hulu, Netflix)
- voice over IP (e.g., Skype)
- real-time video conferencing
- social networking
- search
- ...
- ...

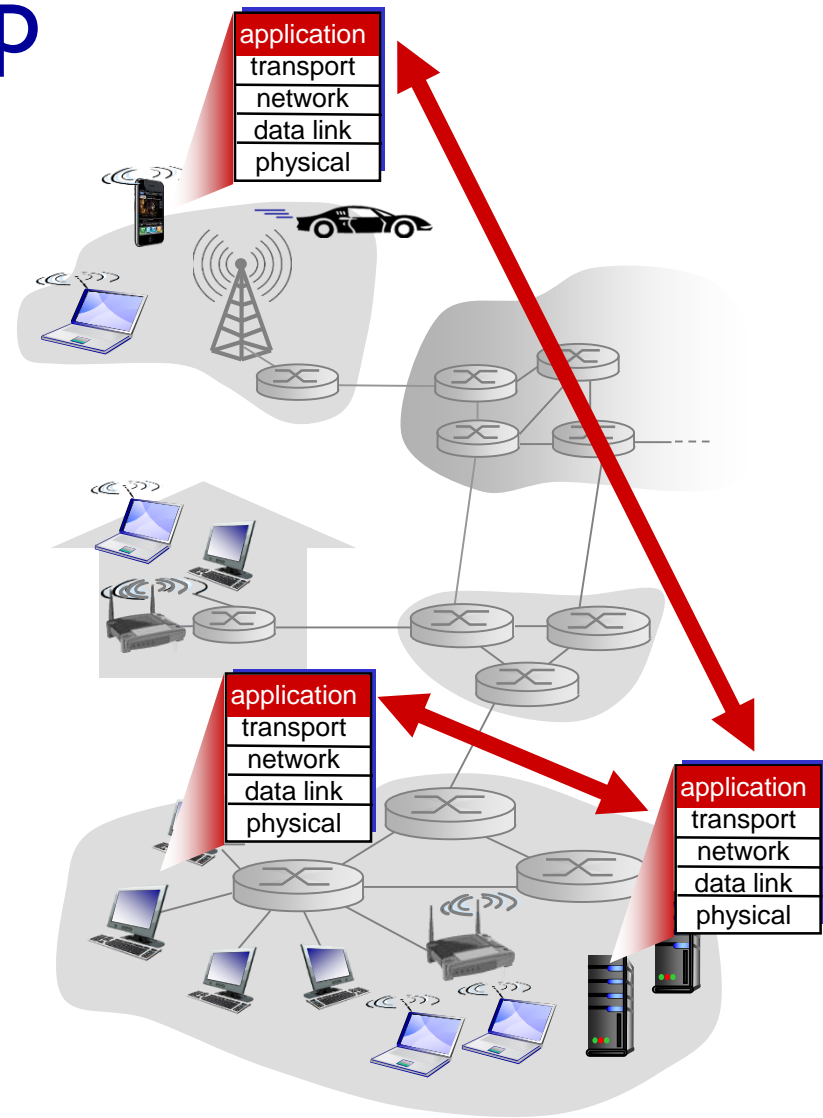
Creating a network app

write programs that:


- run on (different) *end systems*
- communicate over network
- e.g., web server software communicates with browser software

no need to write software
for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation



Recall: Client-server architecture



server:


- always-on host
- permanent IP address
- data centers for scaling

clients:

- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other

Recall: P2P architecture

- no always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
 - *self scalability* – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
 - complex management



Processes communicating

- **process:** program running within a host
 - within same host, two processes communicate using **inter-process communication** (defined by OS)
 - processes in different hosts communicate by exchanging **messages**

clients, servers

client process: process that initiates communication

server process: process that waits to be contacted

- Nodes in P2P architectures have both client processes & server processes.

Addressing processes

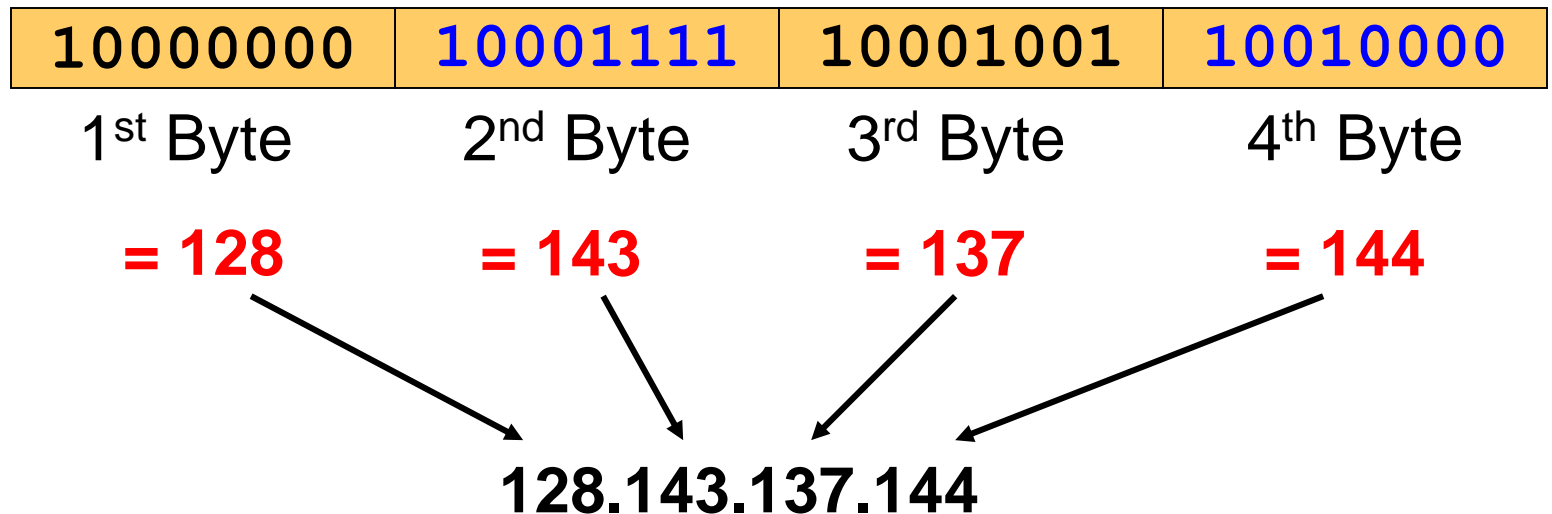
- to receive messages, process must have *identifier*
- host device has unique 32-bit IP address

What is an IP Address?

- An IP address is a unique global address for a network interface
- An IP address:
 - is a **32 bit long** identifier
 - encodes a network number
(**network prefix**)
and a **host number**

Dotted Decimal Notation

- IP addresses are written in **dotted decimal notation**
- Each byte is identified by a decimal number in the range [0..255]:
- **Example:**



Addressing processes

- Q: does IP address of host on which process runs suffice for identifying the process?

Addressing processes

- Q: does IP address of host on which process runs suffice for identifying the process?
- A: no, *many* processes can be running on same host
- *identifier* includes both **IP address** and **port numbers** associated with process on host.
- example port numbers:
 - HTTP server: 80
 - mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
 - **IP address**: 128.119.245.12
 - **port number**: 80

What transport service does an app need?

Data integrity

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

Delay

- some apps (e.g., Internet telephony, interactive games) are time sensitive
- Require low delay to be “effective”

Throughput

- some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- other apps (“elastic apps”) make use of whatever throughput they get

Security

- encryption, data integrity, ...

Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer			
e-mail			
Web documents			
real-time audio/video			
stored audio/video			
interactive games			
text messaging			

Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video			
stored audio/video			
interactive games			
text messaging			

Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	yes, 100s msec
stored audio/video			
interactive games			
text messaging			

Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	Yes: 100s msec
stored audio/video	loss-tolerant	same as above	
interactive games			
text messaging			

Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	Yes: 100s msec
stored audio/video	loss-tolerant	same as above	Yes: few secs
interactive games	loss-tolerant	few kbps up	Yes: 100s msec
text messaging			

Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	yes, 100s msec
stored audio/video	loss-tolerant	same as above	Yes: few secs
interactive games	loss-tolerant	few kbps – 10kbps	Yes: 100s msec
text messaging	no loss	elastic	Yes and no

Internet Transport Protocols

Transport Control Protocol (TCP) service:

- *Connection-oriented*
 - Setup required between client and server processes

Internet Transport Protocols

Transport Control Protocol (TCP) service:

- *Connection-oriented*
 - Setup required between client and server processes
- *Reliable transport* between sending and receiving process
 - Deal with packet drop

Internet Transport Protocols

Transport Control Protocol (TCP) service:

- *Connection-oriented*
 - Setup required between client and server processes
- *Reliable transport* between sending and receiving process
 - Deal with packet drop
- *Flow control*
 - Change rate of sending packets so that receiver is not overwhelmed

Internet Transport Protocols

Transport Control Protocol (TCP) service:

- *Connection-oriented*
 - Setup required between client and server processes
- *Reliable transport* between sending and receiving process
 - Deal with packet drop
- *Flow control*
 - Change rate of sending packets so that receiver is not overwhelmed
- *Does not provide security*

Internet Transport Protocols

Transport Control Protocol (TCP) service:

- *Connection-oriented*
 - Setup required between client and server processes
- *Reliable transport* between sending and receiving process
 - Deal with packet drop
- *Flow control*
 - Change rate of sending packets so that receiver is not overwhelmed
- *Does not provide security*

User Datagram Protocol (UDP) service:

- *unreliable data transfer* between sending and receiving process
- *does not provide:* reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

Q: why bother? Why is there a UDP?

Internet Apps: Application Layer, Transport Layer Protocols

	application	application layer protocol	underlying transport protocol
	e-mail	SMTP [RFC 2821]	TCP
remote terminal access		Telnet [RFC 854]	TCP
	Web	HTTP [RFC 2616]	?
	file transfer	FTP [RFC 959]	?
streaming multimedia		HTTP (e.g., YouTube),	?
Internet telephony		RTP [RFC 3550] SIP [RFC 3261], or proprietary (e.g., Skype)	?

Internet Apps: Application Layer, Transport Layer Protocols

	application	application layer protocol	underlying transport protocol
	e-mail	SMTP [RFC 2821]	TCP
remote terminal access		Telnet [RFC 854]	TCP
	Web	HTTP [RFC 2616]	TCP
	file transfer	FTP [RFC 959]	TCP
streaming multimedia		HTTP (e.g., YouTube),	TCP
Internet telephony		RTP [RFC 3550] SIP [RFC 3261], or proprietary (e.g., Skype)	UDP or TCP

Topics

- Application Layer
 - Network Applications
- Web and HTTP
- Socket Programming: Introduction

Web and HTTP

First, a review...

- *web page* consists of *objects*
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of *base HTML-file* which includes *several referenced objects*
- each object is addressable by a *URL*, e.g.,

`www.someschool.edu/someDept/pic.gif`

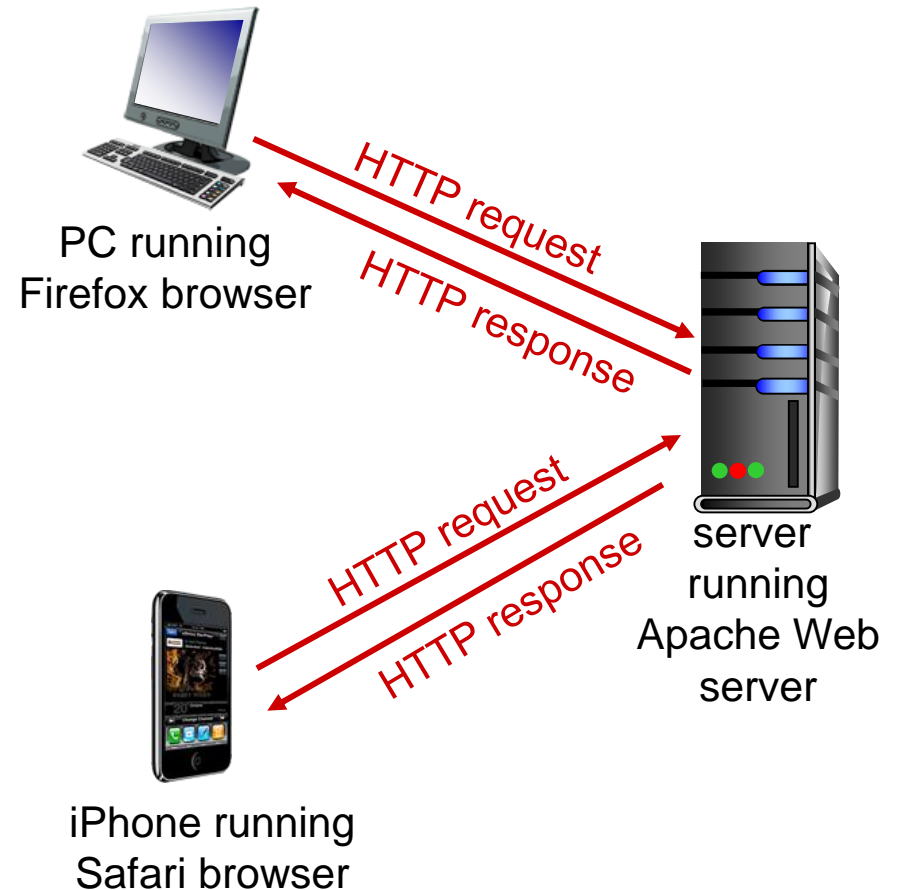
host name

path name

HTTP Overview

HTTP: Hypertext Transfer Protocol

- Web's application layer protocol
- client/server model
 - **client**: browser that requests, receives, (using HTTP protocol) and "displays" Web objects
 - **server**: Web server sends (using HTTP protocol) objects in response to requests



HTTP Overview (continued)

Uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

HTTP is “stateless”

- server maintains no information about past client requests

aside

protocols that maintain
“state” are complex!

- past history (state) must be maintained
- if server/client crashes, their views of “state” may be inconsistent, must be reconciled

HTTP connections

non-persistent HTTP

- at most one object sent over TCP connection
 - connection then closed
- downloading multiple objects required multiple connections

persistent HTTP

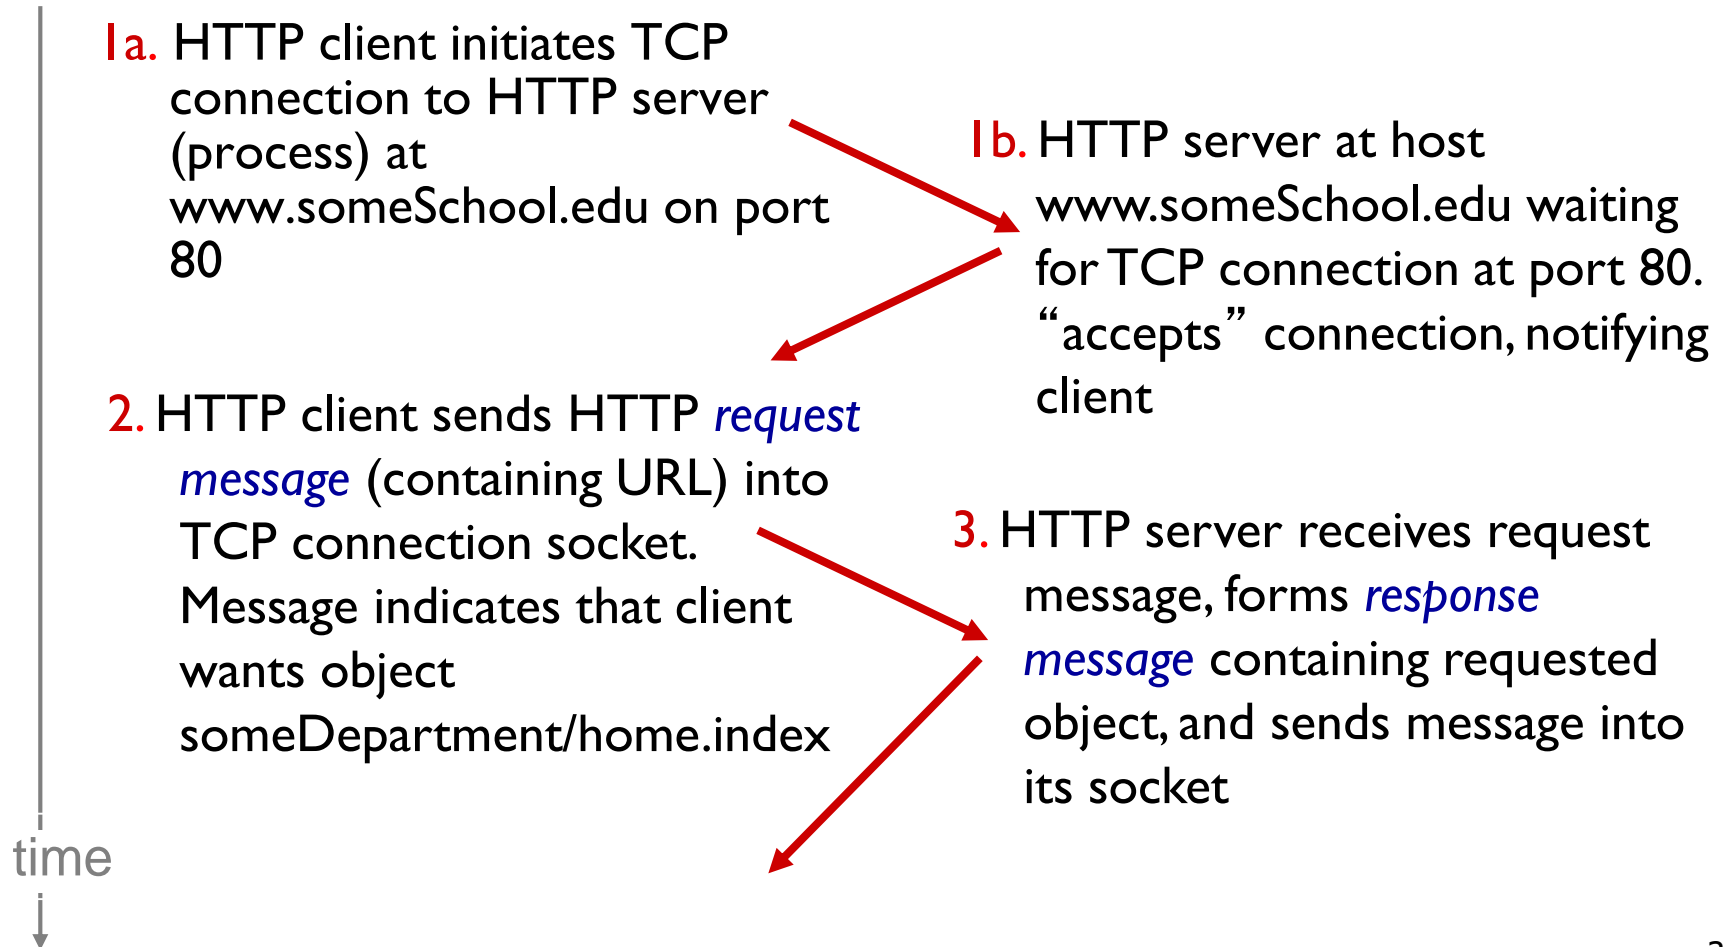
- multiple objects can be sent over single TCP connection between client, server

Non-persistent HTTP

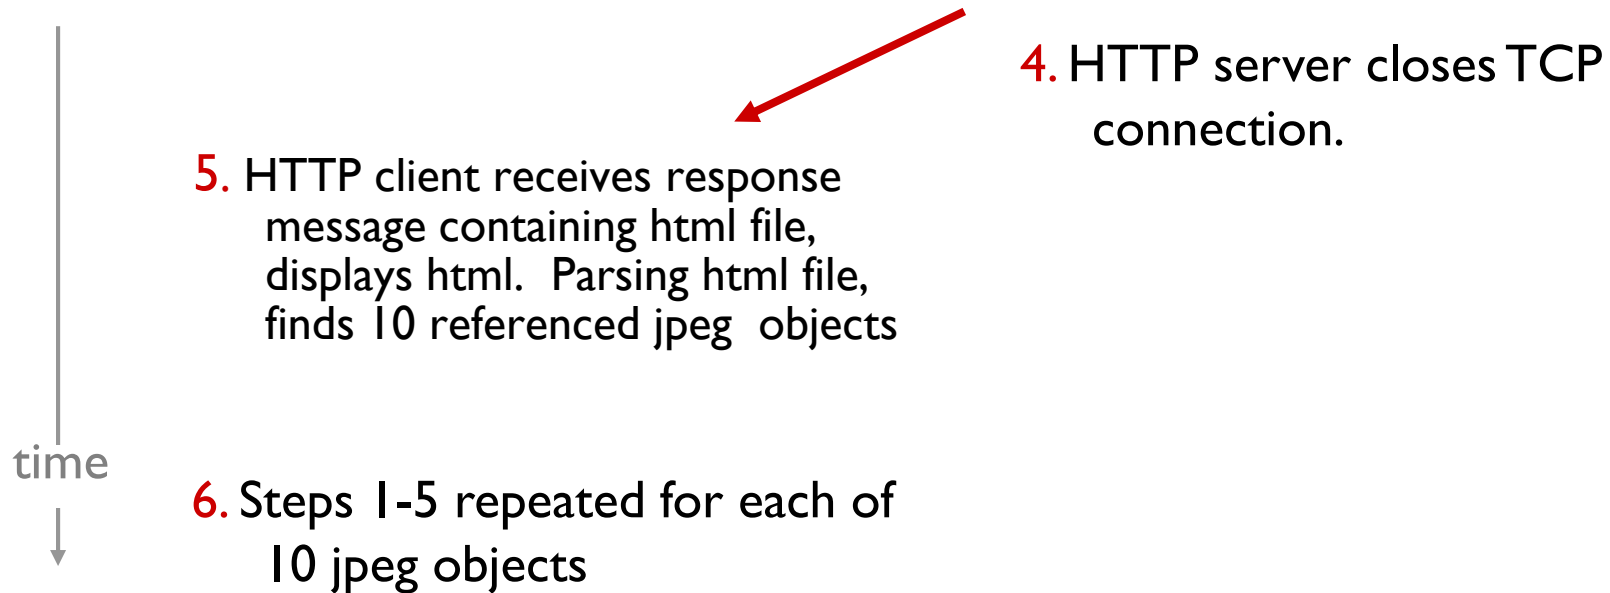
suppose user enters URL:

`www.someSchool.edu/someDepartment/home.index`

(contains text,
references to 10
jpeg images)



Non-persistent HTTP (cont.)




Non-persistent HTTP: response time

RTT (definition): time for a small packet to travel from client to server and back

HTTP response time:

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time
- non-persistent HTTP response time =
 $2\text{RTT} + \text{file transmission time}$



Persistent HTTP

non-persistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for *each* TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

persistent HTTP:

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

HTTP request message

- Two types of HTTP messages: *request, response*
- **HTTP request message:**
 - ASCII (human-readable format)

request line
(GET, POST,
HEAD commands)

header
lines

carriage return,
line feed at start
of line indicates
end of header lines

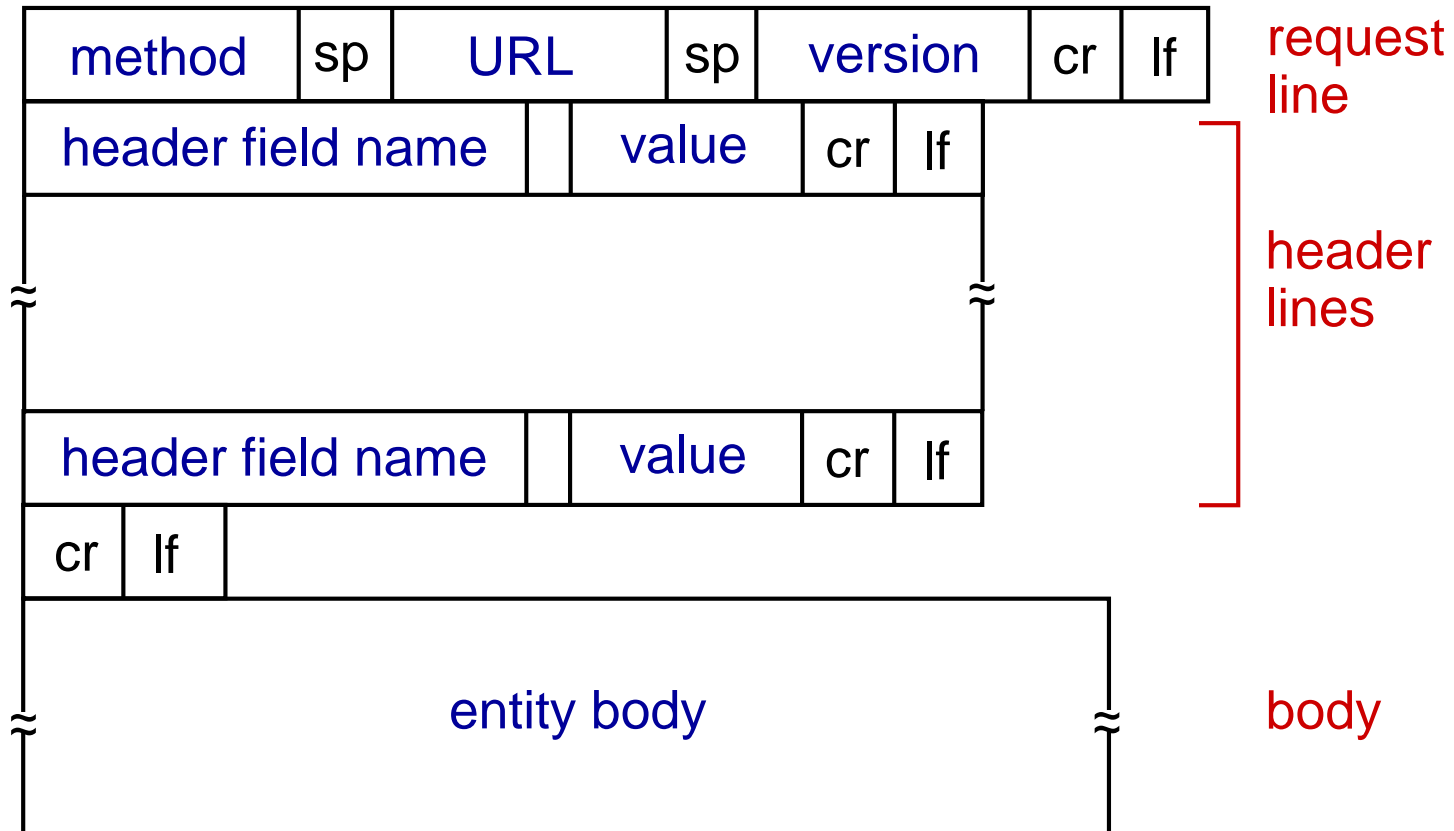
```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

carriage return character

line-feed character

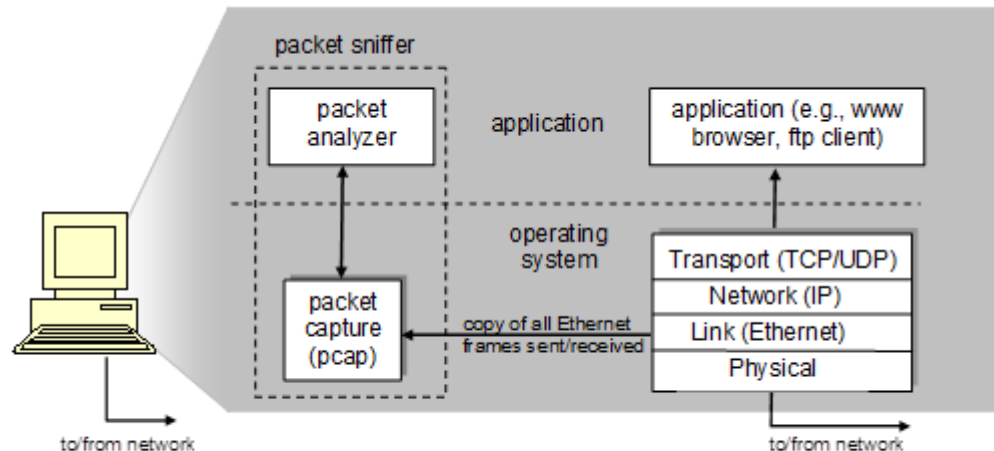
The diagram illustrates the structure of an HTTP request message. It shows a sequence of lines: a request line, followed by multiple header lines, and ending with a blank line. Blue arrows point from descriptive text labels to specific parts of the message. One arrow points from 'request line (GET, POST, HEAD commands)' to the first line. Another arrow points from 'header lines' to the block of lines between the request line and the final blank line. A third arrow points from 'carriage return, line feed at start of line indicates end of header lines' to the final blank line. Two additional arrows point from 'carriage return character' and 'line-feed character' to the backslash-r and backslash-n characters at the end of the first line.

HTTP request message: general format



Class Exercise

- Introducing Wireshark
 - Download from <https://www.wireshark.org/download.html>
- Protocol Analyzer/Package Sniffer
 - Captures messages being sent/received over a network
 - Packet capture library receives a copy of every **link-layer frame** that is sent from or received by your computer
 - Packet sniffer component displays the contents of all fields within a protocol message.



Class Exercise cont.

- Start up your favorite web browser, which will display your selected homepage.
- Start up the Wireshark software. Wireshark has not yet begun capturing packets.
- To begin packet capture, select the Capture pull down menu and select *Interfaces*.
- While Wireshark is running, enter the URL:
http://gaia.cs.umass.edu/kurose_ross/interactive/index.php
and have that page displayed in your browser. Click at other links on this webpage.
- Stop Wireshark packet capture by selecting stop in the Wireshark capture window.

Class Exercise cont.

- Type in “http” (without the quotes, and in lower case – all protocol names are in lower case in Wireshark) into the display filter specification window at the top of the main Wireshark window.
- Find the HTTP GET message that was sent from your computer to the gaia.cs.umass.edu HTTP server.

Class Exercise cont.

- Other questions
 - List 3 different protocols that appear in the protocol column in the unfiltered packet-listing window.
 - How long did it take from when the HTTP GET message was sent until the HTTP OK reply was received? (By default, the value of the Time column in the packet-listing window is the amount of time, in seconds, since Wireshark tracing began).
 - What is the Internet address of the `gaia.cs.umass.edu` (also known as `www-net.cs.umass.edu`)? What is the Internet address of your computer?

Class exercise 2

- Repeat steps from previous exercise:
 - Start up your favorite web browser, which will display your selected homepage.
 - Start up the Wireshark software. Wireshark has not yet begun capturing packets.
 - To begin packet capture, select the Capture pull down menu and select *Interfaces*.
 - While Wireshark is running, enter the URL:
http://gaia.cs.umass.edu/kurose_ross/interactive/index.php

Class exercise 2

- Is your browser running HTTP version 1.0 or 1.1? What version of HTTP is the server running?
- What languages (if any) does your browser indicate that it can accept to the server?
- What is the IP address of your computer? Of the `gaia.cs.umass.edu` server?
- What is the status code returned from the server to your browser?
- When was the HTML file that you are retrieving last modified at the server?
- How many bytes of content are being returned to your browser?

HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:

200 OK

- request succeeded, requested object later in this msg

301 Moved Permanently

- requested object moved, new location specified later in this msg (Location:)

400 Bad Request

- request msg not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

Uploading form input

POST method:

- web page often includes form input
- input is uploaded to server in entity body

URL method:

- uses GET method
- input is uploaded in URL field of request line:

`www.somesite.com/animalsearch?monkeys&banana`

Method types

HTTP/1.0:

- GET
- POST
- HEAD
 - asks server to leave requested object out of response

HTTP/1.1:

- GET, POST, HEAD
- PUT
 - uploads file in entity body to path specified in URL field
- DELETE
 - deletes file specified in the URL field

HTTP response message

status line
(protocol
status code
status phrase)

header
lines

data, e.g.,
requested
HTML file

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02
GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html; charset=ISO-8859-
1\r\n
\r\n
data data data data data ...
```

Trying out HTTP (client side) for yourself

1. Telnet to your favorite Web server:

`telnet gaia.cs.umass.edu 80` { opens TCP connection to port 80
(default HTTP server port)
at gaia.cs.umass.edu.
anything typed in will be sent
to port 80 at gaia.cs.umass.edu

2. type in a GET HTTP request:

`GET /kurose_ross/interactive/index.php HTTP/1.1`
`Host: gaia.cs.umass.edu` { by typing this in (hit carriage
return twice), you send
this minimal (but complete)
GET request to HTTP server

3. look at response message sent by HTTP server!
(or use Wireshark to look at captured HTTP request/response)

User-server state: cookies

many Web sites use cookies

four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

example:

- Susan always access Internet from PC
- visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
 - unique ID
 - entry in backend database for ID

Cookies: keeping “state” (cont.)

client



server



cookie file



ebay 8734
amazon 1678


usual http request msg

Amazon server
creates ID
1678 for user

usual http response
set-cookie: 1678

create
entry

backend
database



usual http request msg
cookie: 1678

cookie-
specific
action

access

usual http response msg

access

cookie-
specific
action

one week later:



ebay 8734
amazon 1678

usual http request msg
cookie: 1678

usual http response msg

Cookies (continued)

what cookies can be used for:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

aside

cookies and privacy:

- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites

how to keep “state”:

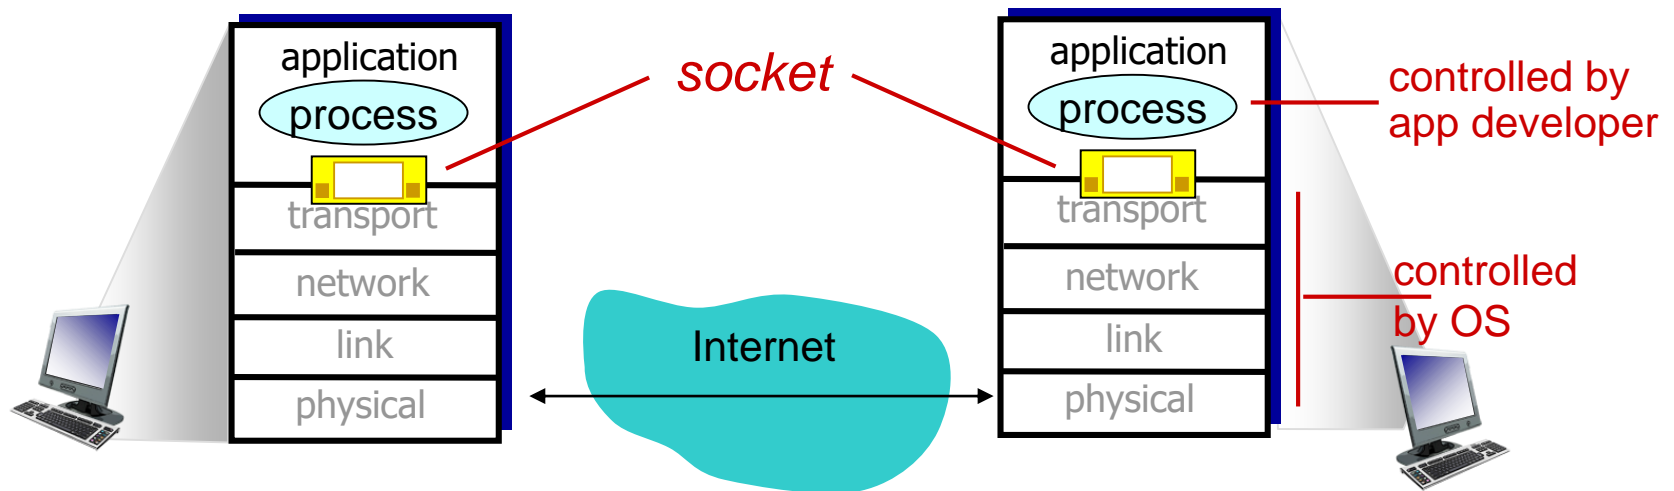
- protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: http messages carry state

Topics

- Application Layer
 - Network Applications
- Web and HTTP
- Socket Programming: Introduction

Sockets

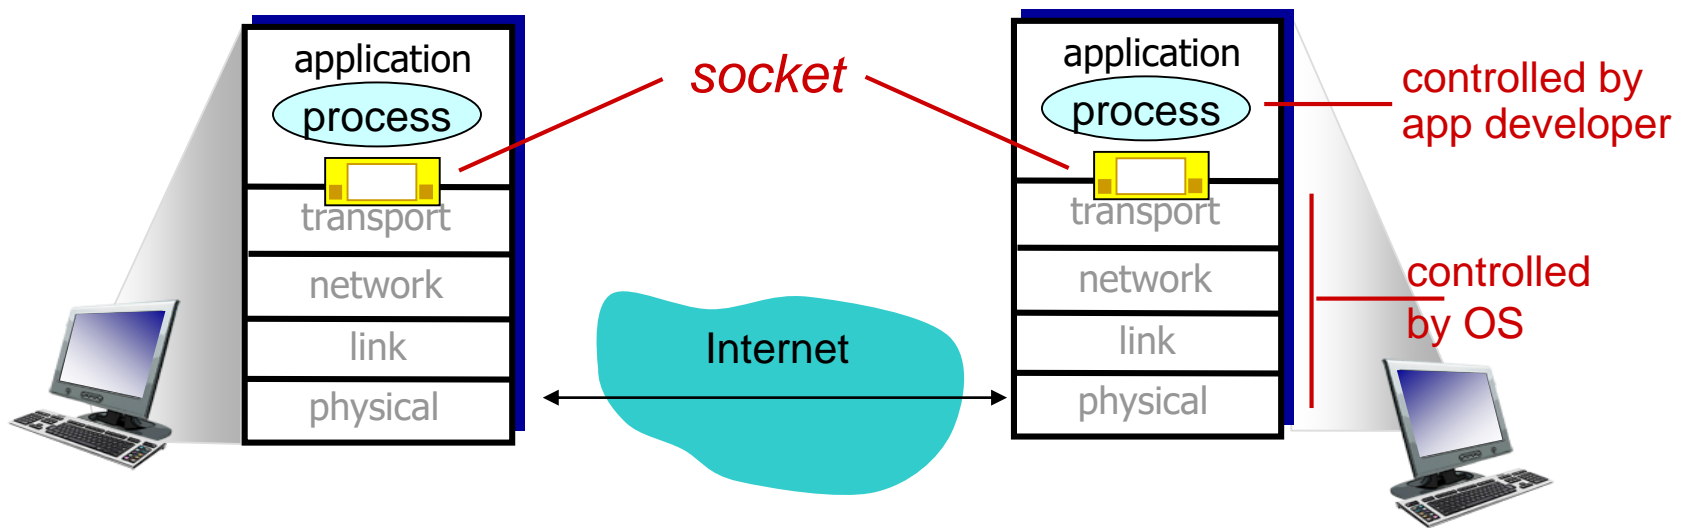
- Process sends/receives messages to/from its **socket**
- Socket analogous to door
 - sending process shoves message out door
 - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process



Socket programming

goal: learn how to build client/server applications that communicate using sockets

socket: door between application process and end-end-transport protocol



Sockets Programming

- Concept was developed in the 1980s in the UNIX environment as the Berkeley Sockets Interface
 - De facto standard application programming interface (API)
 - Basis for Window Sockets (WinSock)
- Enables communication between a client and server process
- May be connection oriented or connectionless

The Socket

- Formed by the concatenation of a port value and an IP address
 - Unique throughout the Internet
- Used to define an API
 - Generic communication interface for writing programs that use TCP or UDP
- Stream sockets
 - All blocks of data sent between a pair of sockets are guaranteed for delivery and arrive in the order that they were sent
- Datagram sockets
 - Delivery is not guaranteed, nor is order necessarily preserved
- Raw sockets
 - Allow direct access to lower-layer protocols

Socket programming

Two socket types for two transport services:

- **UDP:** Datagram Sockets
- **TCP:** Stream Sockets

Application Example:

1. client reads a line of characters (data) from its keyboard and sends data to server
2. server receives the data and converts characters to uppercase
3. server sends modified data to client
4. client receives modified data and displays line on its screen

Java Sockets Programming

- The package `java.net` provides support for sockets programming (and more).
- Typically you import everything defined in this package with:

```
import java.net.*;
```

Classes

InetAddress

Socket

ServerSocket

DatagramSocket

DatagramPacket

InetAddress class

- static methods you can use to create new InetAddress objects.
 - `getByName(String host)`
 - `getAllByName(String host)`
 - `getLocalHost()`

```
InetAddress x = InetAddress.getByName (  
                                "cse.unr.edu" ) ;
```

❖ **Throws UnknownHostException**

Sample Code: Lookup.java

```
try {  
  
    InetAddress a = InetAddress.getByName(hostname);  
  
    System.out.println(hostname + ":" +  
                        a.getHostAddress());  
  
} catch (UnknownHostException e) {  
  
    System.out.println("No address found for " +  
                      hostname);  
  
}
```

Sample Code: Lookup.java cont.

- Uses InetAddress class to lookup hostnames found on command line.

```
> java Lookup cse.unr.edu www.yahoo.com  
cse.unr.edu:134.197.40.9  
www.yahoo.com:209.131.36.158
```

Socket class

- Corresponds to active TCP sockets only!
 - client sockets
 - socket returned by `accept()`;
- Passive sockets are supported by a different class:
 - `ServerSocket`
- UDP sockets are supported by
 - `DatagramSocket`

JAVA TCP Sockets

- `java.net.Socket`
 - Implements client sockets (also called just “sockets”).
 - An endpoint for communication between two machines.
 - Constructor and Methods
 - `Socket(String host, int port)`: Creates a stream socket and connects it to the specified port number on the named host.
 - `InputStream getInputStream()`
 - `OutputStream getOutputStream()`
 - `close()`
- `java.net.ServerSocket`
 - Implements server sockets.
 - Waits for requests to come in over the network.
 - Performs some operation based on the request.
 - Constructor and Methods
 - `ServerSocket(int port)`
 - `Socket Accept()`: Listens for a connection to be made to this socket and accepts it. This method blocks until a connection is made.

Socket Constructors

- Constructor creates a TCP connection to a named TCP server.
 - There are a number of constructors:

```
Socket(InetAddress server, int port);
```

```
Socket(InetAddress server, int port,  
        InetAddress local, int localport);
```

```
Socket(String hostname, int port);
```

Socket Methods

```
void close() ;
```

```
InetAddress getAddress() ;
```

```
InetAddress getLocalAddress() ;
```

```
InputStream getInputStream() ;
```

```
OutputStream getOutputStream() ;
```

- Lots more (setting/getting socket options, partial close, etc.)

ServerSocket Class

(TCP Passive Socket)

- Constructors:

```
ServerSocket(int port) ;
```

```
ServerSocket(int port, int backlog) ;
```

```
ServerSocket(int port, int backlog,  
             InetAddress bindAddr) ;
```

ServerSocket Methods

`Socket accept() ;`

`void close() ;`

`InetAddress getInetAddress() ;`

`int getLocalPort() ;`

`throw IOException, SecurityException`

Socket functional calls

- `socket ()`: Create a socket
- `bind()`: bind a socket to a local IP address and port #
- `listen()`: passively waiting for connections
- `connect()`: initiating connection to another socket
- `accept()`: accept a new connection
- `Write()`: write data to a socket
- `Read()`: read data from a socket
- `sendto()`: send a datagram to another UDP socket
- `recvfrom()`: read a datagram from a UDP socket
- `close()`: close a socket (tear down the connection)

Socket programming *with TCP*

client must contact server

- server process must first be running
- server must have created socket (door) that welcomes client's contact

client contacts server by:

- Creating TCP socket, specifying IP address, port number of server process
- *when client creates socket:* client TCP establishes connection to server TCP


- when contacted by client, *server TCP creates new socket* for server process to communicate with that particular client
 - allows server to talk with multiple clients
 - source port numbers used to distinguish clients (more in Chap 3)

application viewpoint:

TCP provides reliable, in-order byte-stream transfer (“pipe”) between client and server

Stream jargon

- A **stream** is a sequence of characters that flow into or out of a process.
- An **input stream** is attached to some input source for the process, eg, keyboard or socket.
- An **output stream** is attached to an output source, eg, monitor or socket.



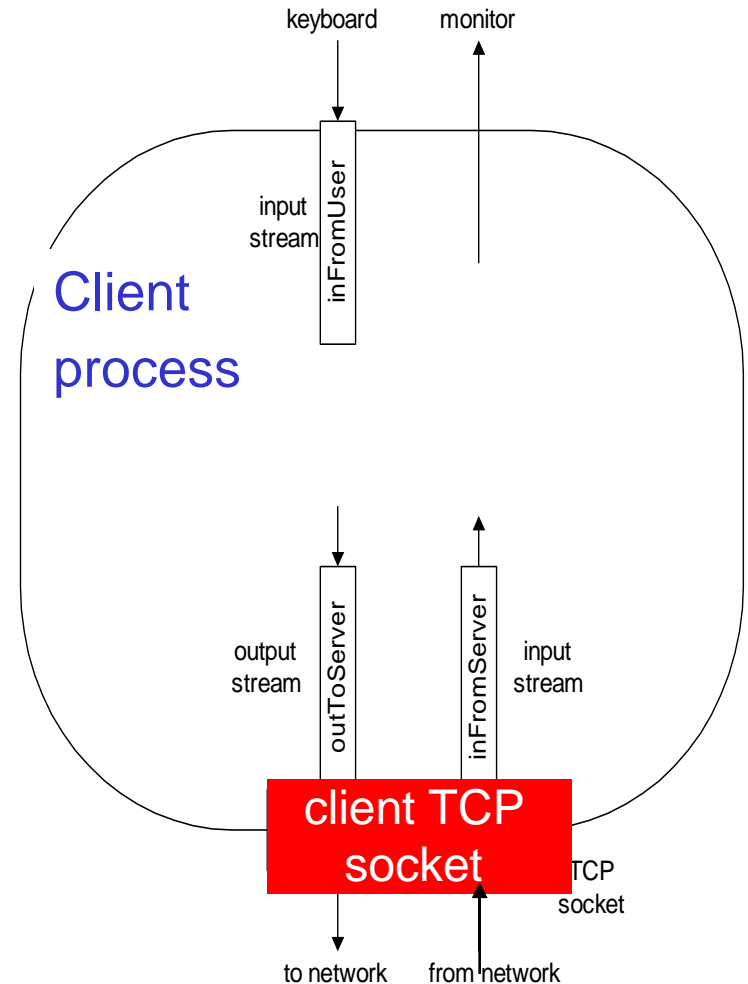
Socket I/O

- Socket I/O is based on the Java I/O support
 - in the package `java.io`
- `InputStream` and `OutputStream` are abstract classes
 - common operations defined for all kinds of `InputStreams`, `OutputStreams`...

Socket programming with TCP

Example client-server app:

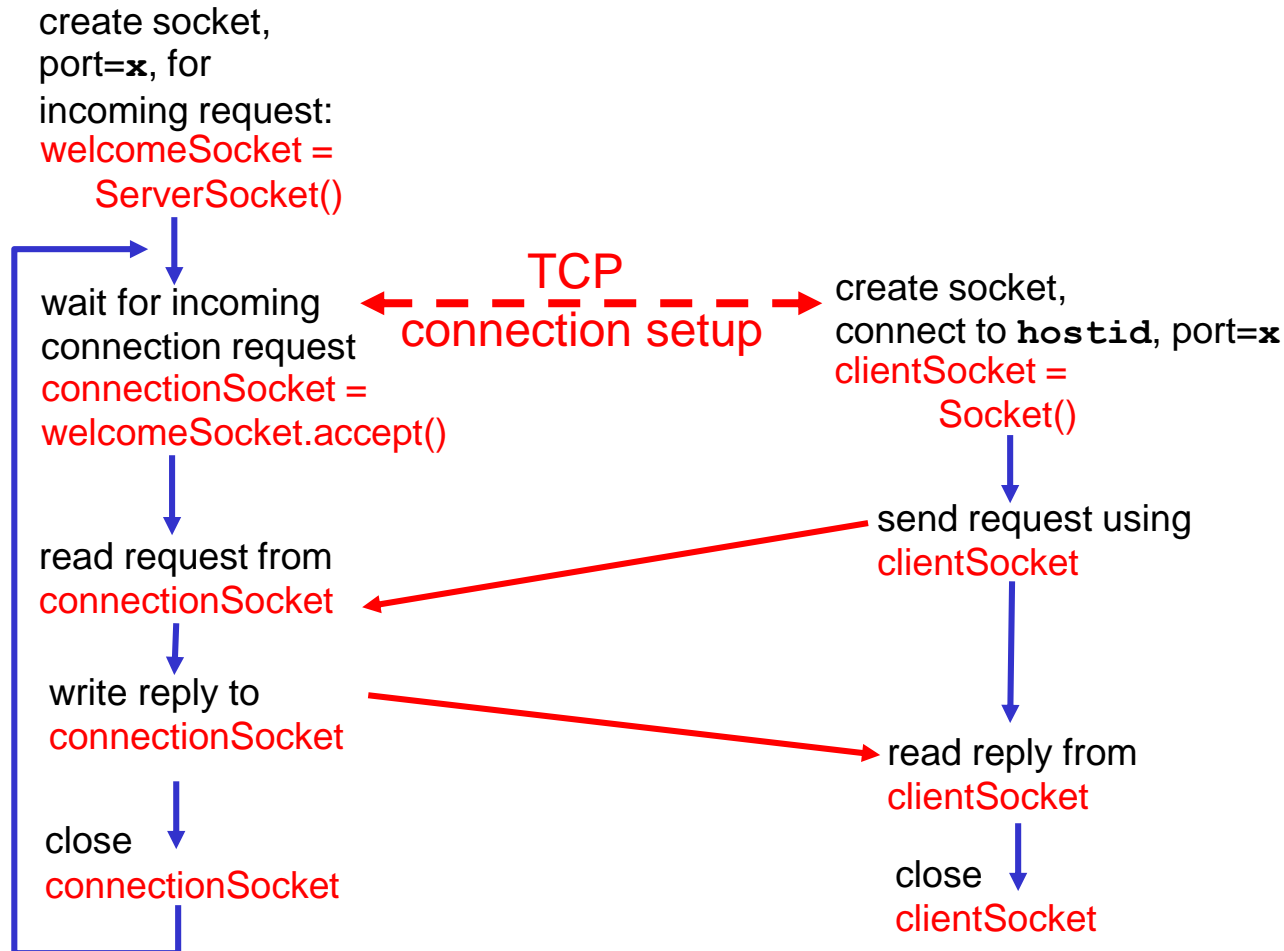
- 1) client reads line from standard input (**inFromUser** stream), sends to server via socket (**outToServer** stream)
- 2) server reads line from socket
- 3) server converts line to uppercase, sends back to client
- 4) client reads, prints modified line from socket (**inFromServer** stream)



Client/server socket interaction: TCP

Server (running on `hostid`)

Client



Example: Java client (TCP)

```
import java.io.*;
import java.net.*;
class TCPCClient {
```

```
    public static void main(String argv[]) throws Exception
    {
```

```
        String sentence;
        String modifiedSentence;
```

Create
input stream



```
        BufferedReader inFromUser =
            new BufferedReader(new InputStreamReader(System.in));
```

Create
client socket,
connect to server



```
        Socket clientSocket = new Socket("hostname", 6789);
```

Create
output stream
attached to socket



```
        DataOutputStream outToServer =
            new DataOutputStream(clientSocket.getOutputStream());
```

Example: Java client (TCP), cont.

Create
input stream
attached to socket

Send line
to server

Read line
from server

```
BufferedReader inFromServer =  
    new BufferedReader(new  
        InputStreamReader(clientSocket.getInputStream()));  
  
sentence = inFromUser.readLine();  
  
outToServer.writeBytes(sentence + '\n');  
  
modifiedSentence = inFromServer.readLine();  
  
System.out.println("FROM SERVER: " + modifiedSentence);  
  
clientSocket.close();  
  
}  
}
```


Example: Java server (TCP)

```
import java.io.*;  
import java.net.*;
```

```
class TCPServer {
```

```
    public static void main(String argv[]) throws Exception  
    {
```

```
        String clientSentence;  
        String capitalizedSentence;
```

Create
welcoming socket
at port 6789

```
        ServerSocket welcomeSocket = new ServerSocket(6789);
```

Wait, on welcoming
socket for contact
by client

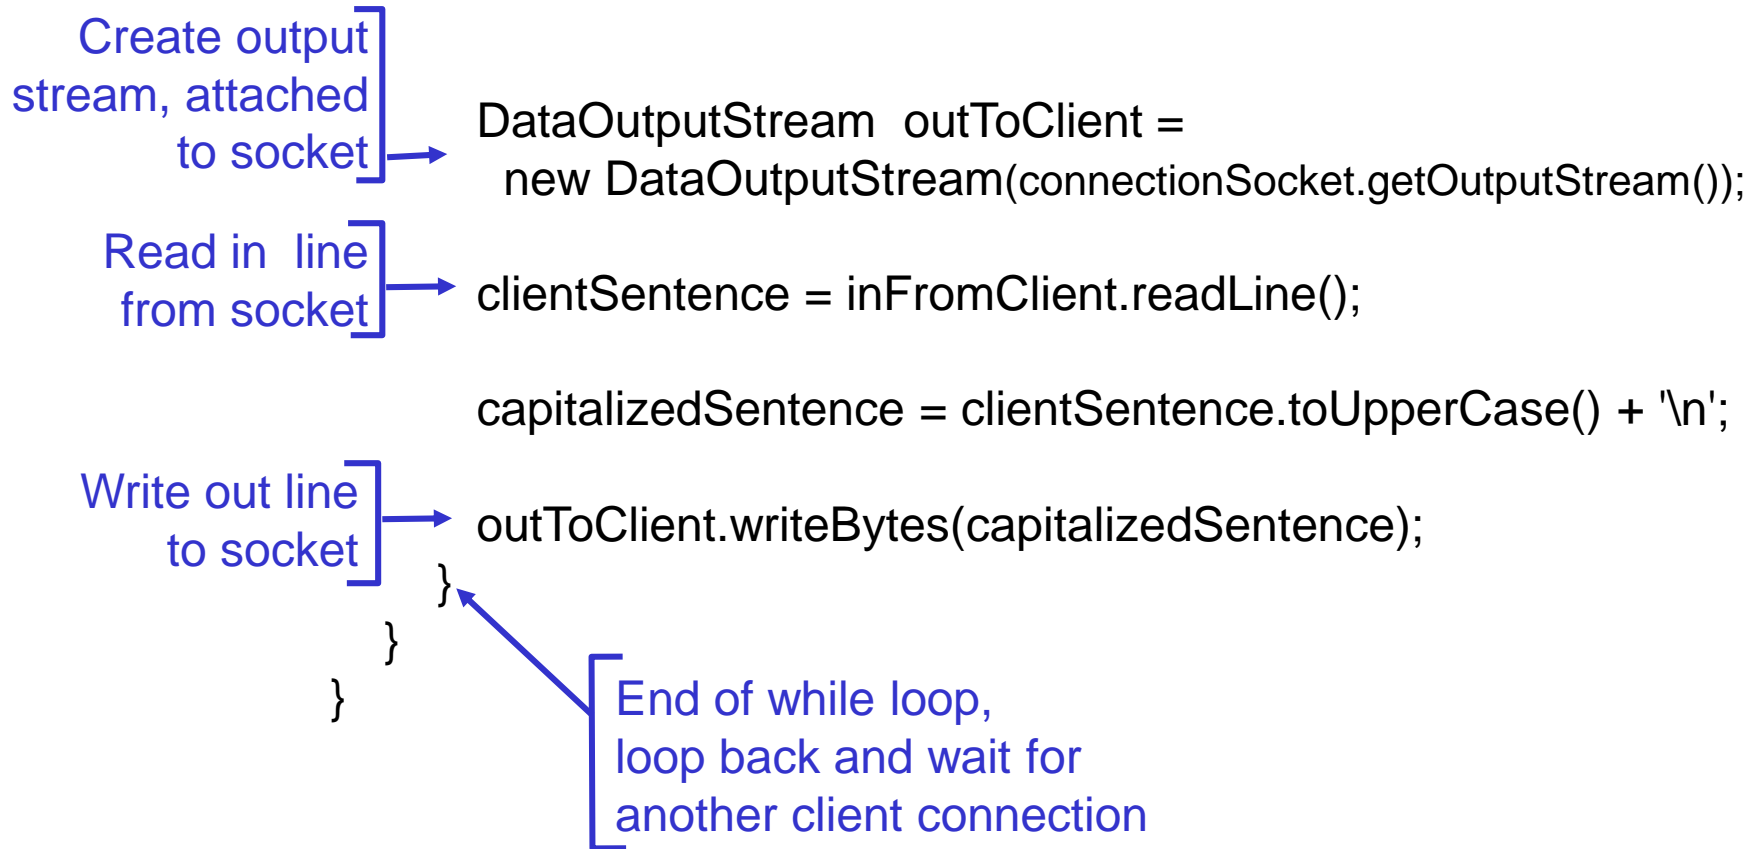
```
        while(true) {
```

```
            Socket connectionSocket = welcomeSocket.accept();
```

Create input
stream, attached
to socket

```
            BufferedReader inFromClient =  
                new BufferedReader(new  
                    InputStreamReader(connectionSocket.getInputStream()));
```

Example: Java server (TCP), cont



Class Exercise

- Make the server and client code from previous exercise work on two different nodes.

Socket programming *with UDP*

UDP: no “connection” between client & server

- no handshaking before sending data
- sender explicitly attaches IP destination address and port # to each packet
- receiver extracts sender IP address and port# from received packet

UDP: transmitted data may be lost or received out-of-order

Application viewpoint:

- UDP provides *unreliable* transfer of groups of bytes (“datagrams”) between client and server

Client/server socket interaction: UDP

server (running on *serverIP*)

create socket, port= x:
`serverSocket =
socket(AF_INET,SOCK_DGRAM)`

↓
read datagram from
`serverSocket`

↓
write reply to
`serverSocket`
specifying
client address,
port number

client

create socket:
`clientSocket =
socket(AF_INET,SOCK_DGRAM)`

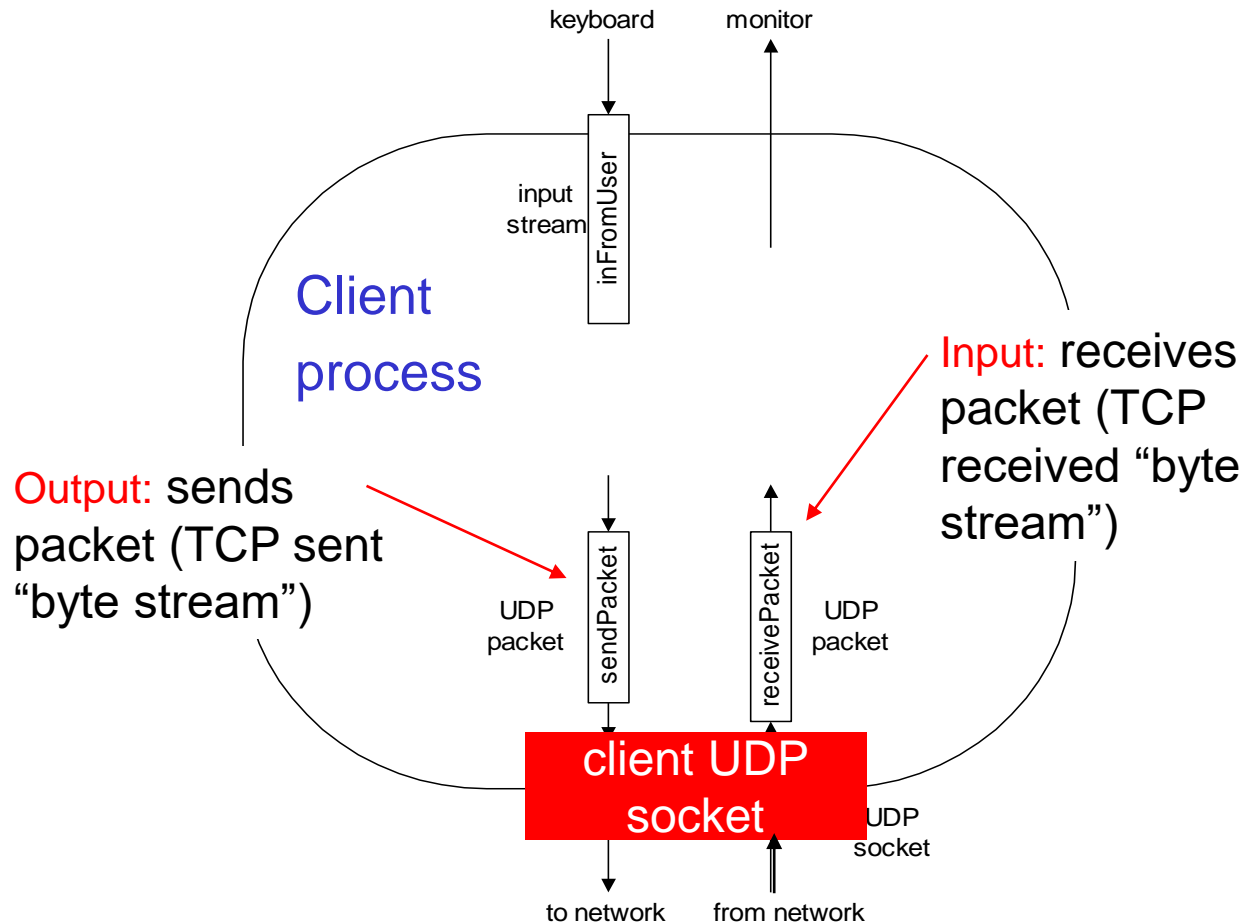
↓
Create datagram with server IP and
port=x; send datagram via
`clientSocket`

↓
read datagram from
`clientSocket`

↓
close
`clientSocket`



Example: Java client (UDP)



Example: Java client (UDP)

```
import java.io.*;
import java.net.*;
```

```
class UDPClient {
    public static void main(String args[]) throws Exception
    {
```

Create
input stream

```
        BufferedReader inFromUser =
```

Create
client socket

```
        new BufferedReader(new InputStreamReader(System.in));
```

```
        DatagramSocket clientSocket = new DatagramSocket();
```

Translate
hostname to IP
address using DNS

```
        InetAddress IPAddress = InetAddress.getByName("hostname");
```

```
        byte[] sendData = new byte[1024];
```

```
        byte[] receiveData = new byte[1024];
```

```
        String sentence = inFromUser.readLine();
```

```
        sendData = sentence.getBytes();
```

Example: Java client (UDP), cont.

Create datagram with
data-to-send,
length, IP addr, port

Send datagram
to server

Read datagram
from server

```
DatagramPacket sendPacket =  
    new DatagramPacket(sendData, sendData.length, IPAddress, 9876);  
  
clientSocket.send(sendPacket);  
  
DatagramPacket receivePacket =  
    new DatagramPacket(receiveData, receiveData.length);  
  
clientSocket.receive(receivePacket);  
  
String modifiedSentence =  
    new String(receivePacket.getData());  
  
System.out.println("FROM SERVER:" + modifiedSentence);  
clientSocket.close();  
}  
}
```


Example: Java server (UDP)

```
import java.io.*;  
import java.net.*;
```

```
class UDPServer {  
    public static void main(String args[]) throws Exception  
    {
```

Create
datagram socket
at port 9876

```
        DatagramSocket serverSocket = new DatagramSocket(9876);
```

```
        byte[] receiveData = new byte[1024];  
        byte[] sendData = new byte[1024];
```

```
        while(true)  
        {
```

Create space for
received datagram

```
            DatagramPacket receivePacket =  
                new DatagramPacket(receiveData, receiveData.length);
```

Receive
datagram

```
            serverSocket.receive(receivePacket);
```

Example: Java server (UDP), cont

```
String sentence = new String(receivePacket.getData());
```

Get IP addr
port #, of
sender

```
InetAddress IPAddress = receivePacket.getAddress();
```

```
int port = receivePacket.getPort();
```

```
String capitalizedSentence = sentence.toUpperCase();
```

```
sendData = capitalizedSentence.getBytes();
```

Create datagram
to send to client

```
DatagramPacket sendPacket =  
    new DatagramPacket(sendData, sendData.length, IPAddress,  
                        port);
```

Write out
datagram
to socket

```
serverSocket.send(sendPacket);
```


```
}  
}  
}
```

End of while loop,
loop back and wait for
another datagram

Web caches (proxy server)

goal: satisfy client request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
 - object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client



More about Web caching

- cache acts as both client and server
 - server for original requesting client
 - client to origin server
- typically cache is installed by ISP (university, company, residential ISP)

why Web caching?

- reduce response time for client request
- reduce traffic on an institution's access link
- Internet dense with caches: enables “poor” content providers to effectively deliver content (so too does P2P file sharing)


Caching example:

assumptions:

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

consequences:

- LAN utilization: 15%
- access link utilization = **99%** *problem!*
- total delay = Internet delay + access delay + LAN delay
= 2 sec + minutes + usecs




Caching example: fatter access link

assumptions:

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: ~~1.54 Mbps~~ → 154 Mbps

consequences:

- LAN utilization: 15%
- access link utilization = ~~99%~~ → 9.9%
- total delay = Internet delay + access delay + LAN delay
= 2 sec + ~~minutes~~ → msec



Cost: increased access link speed (not cheap!)

Caching example: install local cache

assumptions:


- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

consequences:

- LAN utilization: 15%
- access link utilization = ?
- total delay = ?

How to compute link utilization, delay?


Cost: web cache (cheap!)



Caching example: install local cache

Calculating access link utilization, delay with cache:

- suppose cache hit rate is 0.4
 - 40% requests satisfied at cache, 60% requests satisfied at origin
- access link utilization:
 - 60% of requests use access link
- data rate to browsers over access link
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
 - utilization $= 0.9 / 1.54 = .58$
- total delay
 - $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
 - $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$
 - less than with 154 Mbps link (and cheaper too!)



Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version
 - no object transmission delay
 - lower link utilization
- **cache:** specify date of cached copy in HTTP request
If-modified-since: <date>
- **server:** response contains no object if cached copy is up-to-date:
HTTP/1.0 304 Not Modified

