

# **CS 3800: Computer Networks**

## **Lecture 7: Network Layer**

Instructor: John Korah

# Acknowledgement

- The following slides include material from author resources for:
  - KR Text book
  - “Data and computer communications,” William Stallings, Tenth edition

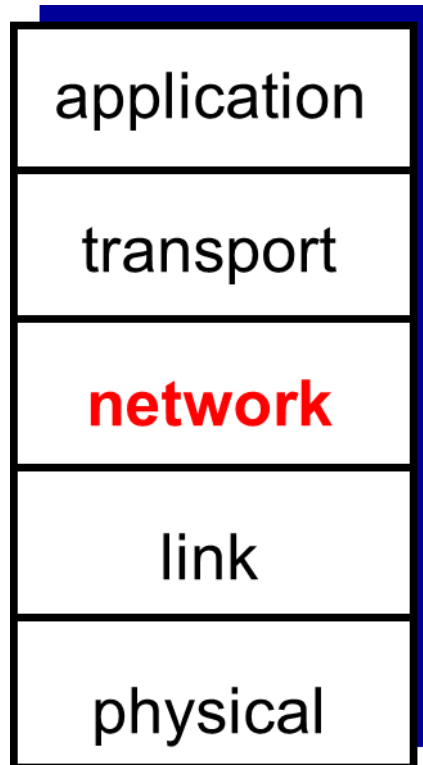
# Learning Goals

- Understand principles behind network layer services, focusing on data plane:
  - network layer service models
  - forwarding versus routing
  - how a router works
  - generalized forwarding
- instantiation, implementation in the Internet

# Topics

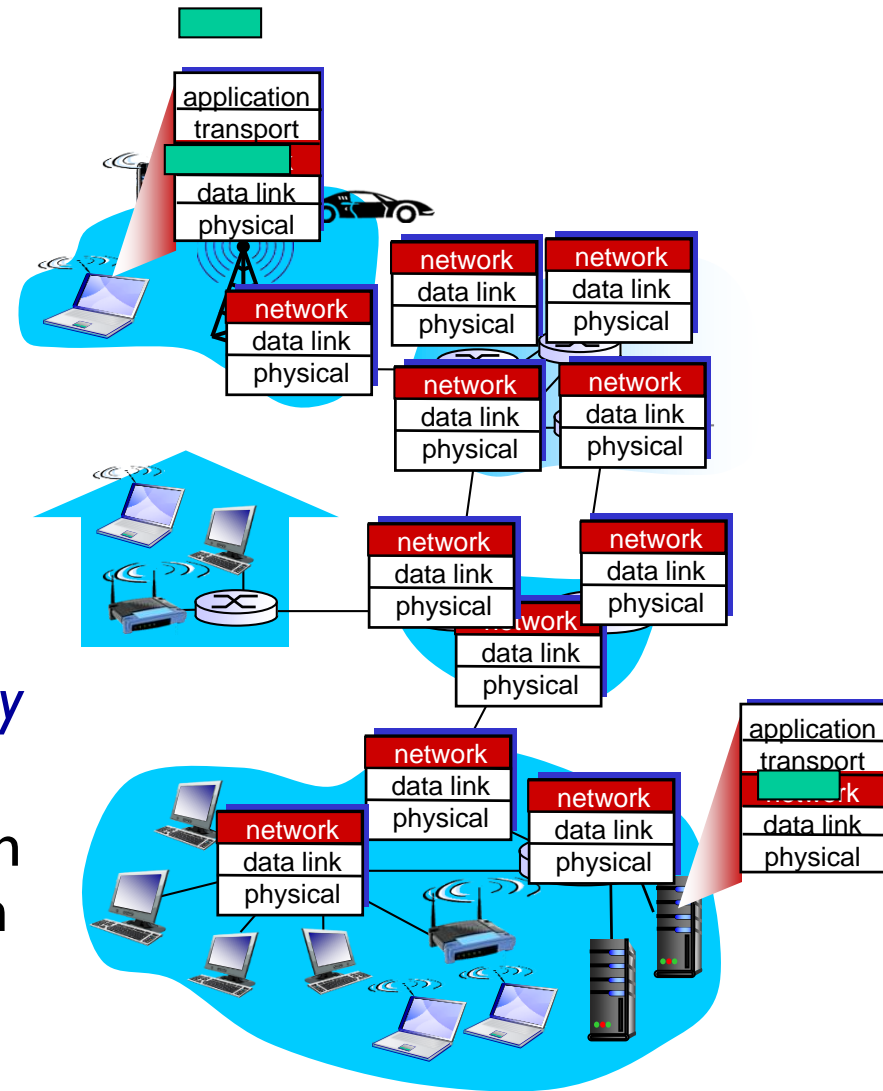
- Overview of Network layer
  - data plane
  - control plane
- What's inside a router
- IP: Internet Protocol
  - datagram format
  - fragmentation
  - IPv4 addressing
  - network address translation
  - IPv6

# Recall



# Network layer

- Transport segment from sending to receiving host
- Encapsulation: On sending side encapsulates segments into datagrams
- De-encapsulation: on receiving side, delivers segments to transport layer
- network layer protocols in *every* host, router
- router examines header fields in all IP datagrams passing through it



# Two key network-layer functions

*Network-layer functions:*

- *forwarding: ?*
- *routing: ?*

# Two key network-layer functions

## *Network-layer functions:*

- *Forwarding*: Move packets from router's input to appropriate router output
- *Routing*: Determine route taken by packets from source to destination
  - *routing algorithms*



# Two key network-layer functions

## *Network-layer functions:*

- *forwarding*: move packets from router's input to appropriate router output
- *routing*: determine route taken by packets from source to destination
  - *routing algorithms*

## *Analogy: taking a trip*

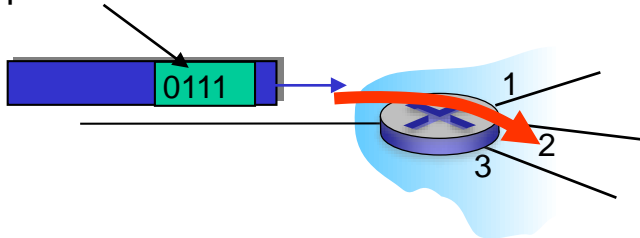
- *forwarding*: process of getting through single interchange
- *routing*: process of planning trip from source to destination

# Network layer: Data plane, Control plane

## *Data plane (forwarding)*

- local, per-router function
- determines how datagram arriving on router input port is forwarded to router output port
- forwarding function

values in arriving packet header

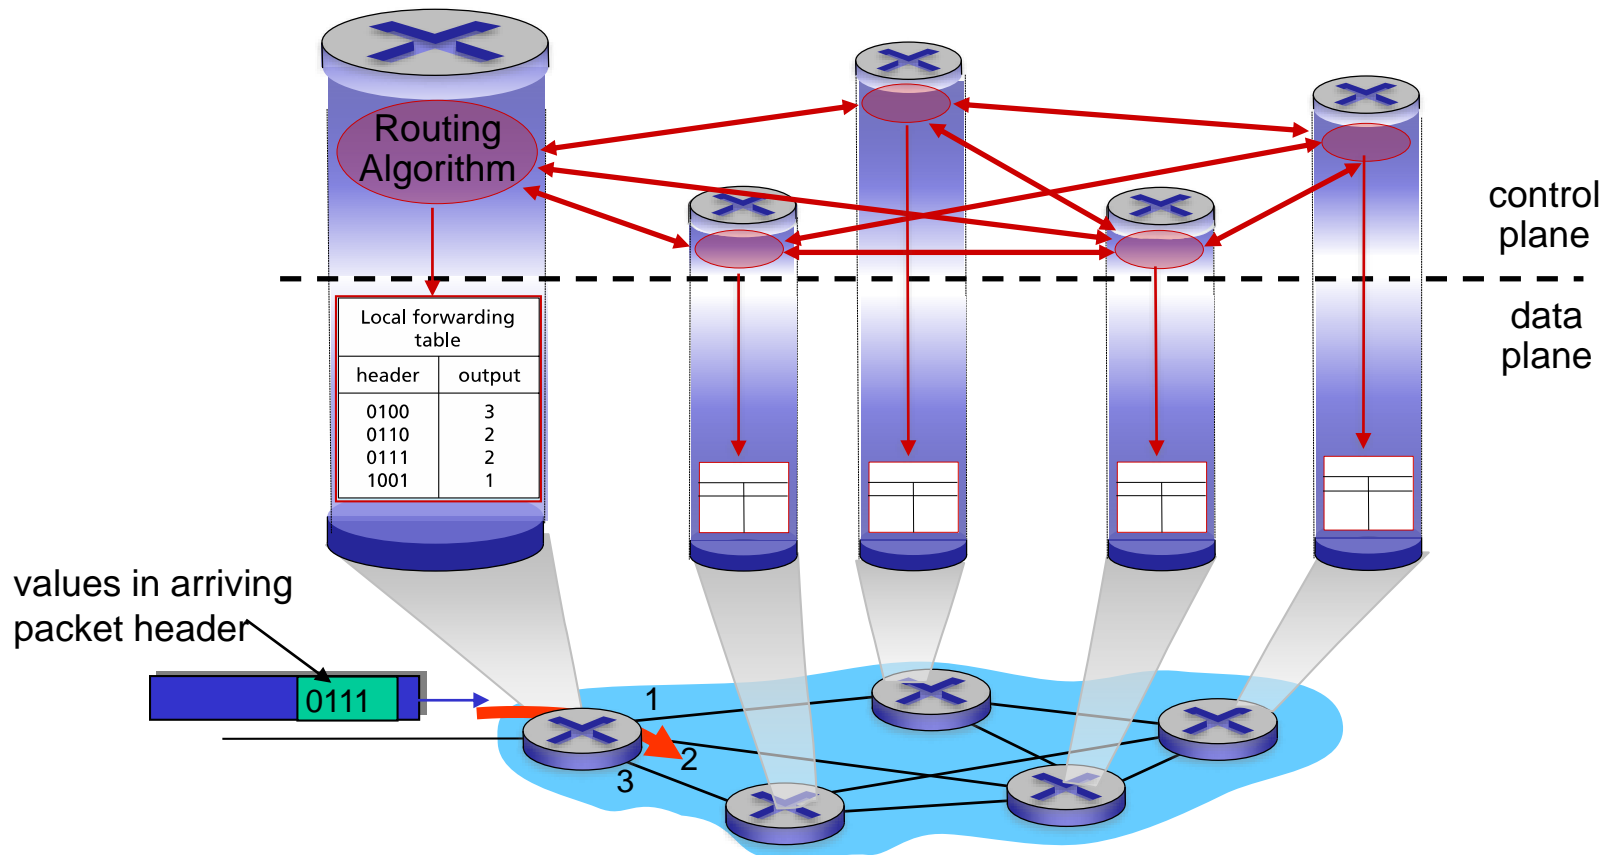


## *Control plane (routing)*

- network-wide logic
- determines how datagram is routed among routers along end-end path from source host to destination host
- two control-plane approaches:
  - *traditional routing algorithms*: implemented in routers
  - *software-defined networking (SDN)*: implemented in (remote) servers

# Per-router control plane

Individual routing algorithm components *in each and every router* interact in the control plane

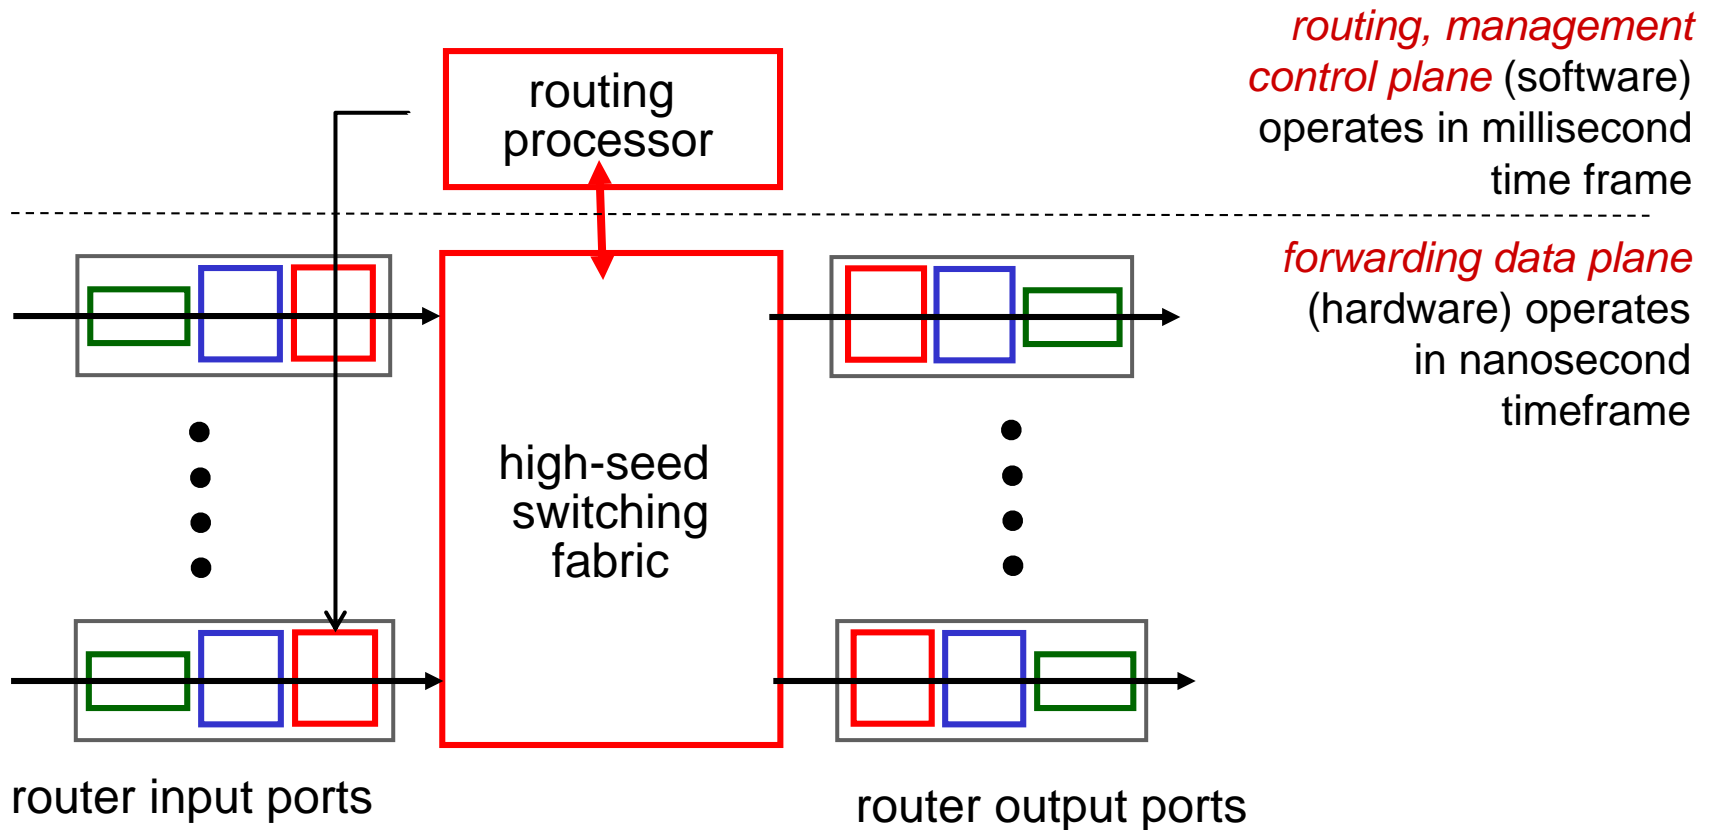


# Topics

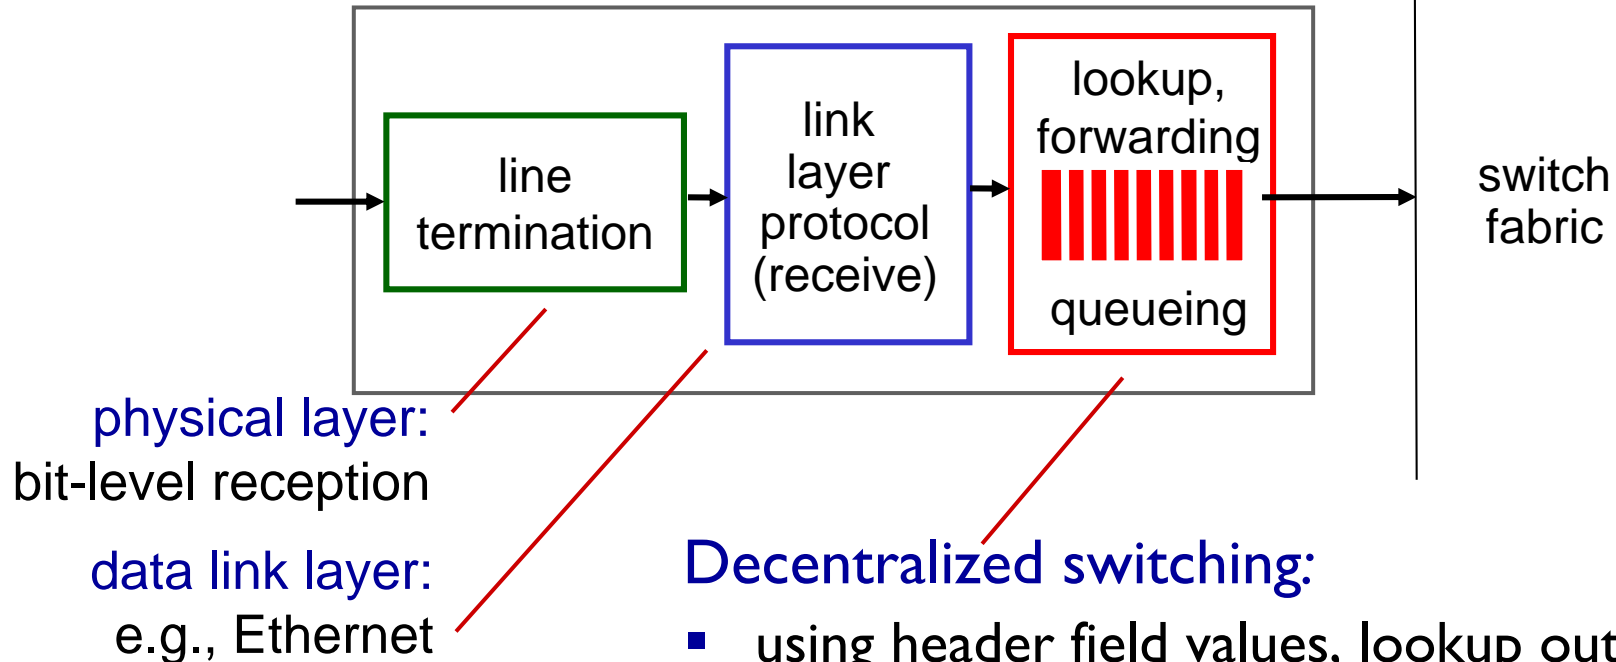
- Overview of Network layer
  - data plane
  - control plane
- What's inside a router
- IP: Internet Protocol
  - datagram format
  - fragmentation
  - IPv4 addressing
  - network address translation
  - IPv6

# Router architecture overview

- high-level view of generic router architecture:



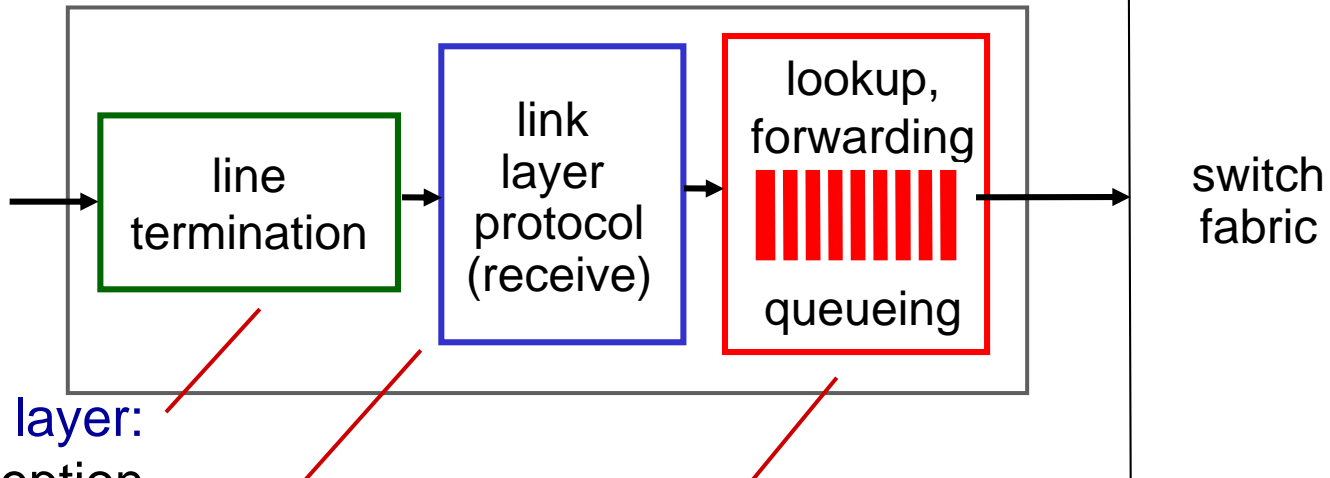
# Input port functions



## Decentralized switching:

- using header field values, lookup output port using forwarding table in input port memory (“*match plus action*”)
- goal: complete input port processing at ‘line speed’
- queuing: if datagrams arrive faster than forwarding rate into switch fabric

# Input port functions



physical layer:  
bit-level reception

data link layer:  
e.g., Ethernet

## Decentralized switching:

- Using header field values, lookup output port using forwarding table in input port memory (“*match plus action*”)
- **Destination-based forwarding:** forward based only on destination IP address (traditional)
- **Generalized forwarding:** forward based on any set of header field values

# Destination-based forwarding

*forwarding table*

Destination Address Range	Link Interface
11001000 00010111 00010000 00000000 through 11001000 00010111 00010111 11111111	0
11001000 00010111 00011000 00000000 through 11001000 00010111 00011000 11111111	1
11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	2
otherwise	3



# Longest prefix matching

## *longest prefix matching*

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

examples:

DA: 11001000 00010111 00010110 10100001

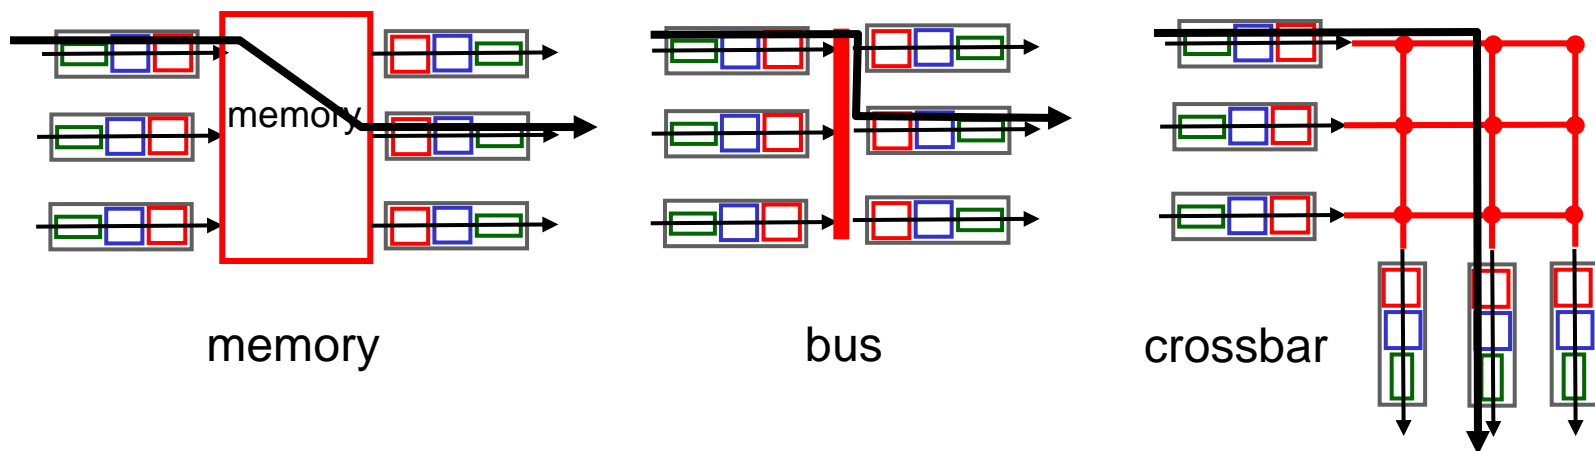
which interface?

DA: 11001000 00010111 00011000 10101010

which interface?

# Switching fabrics

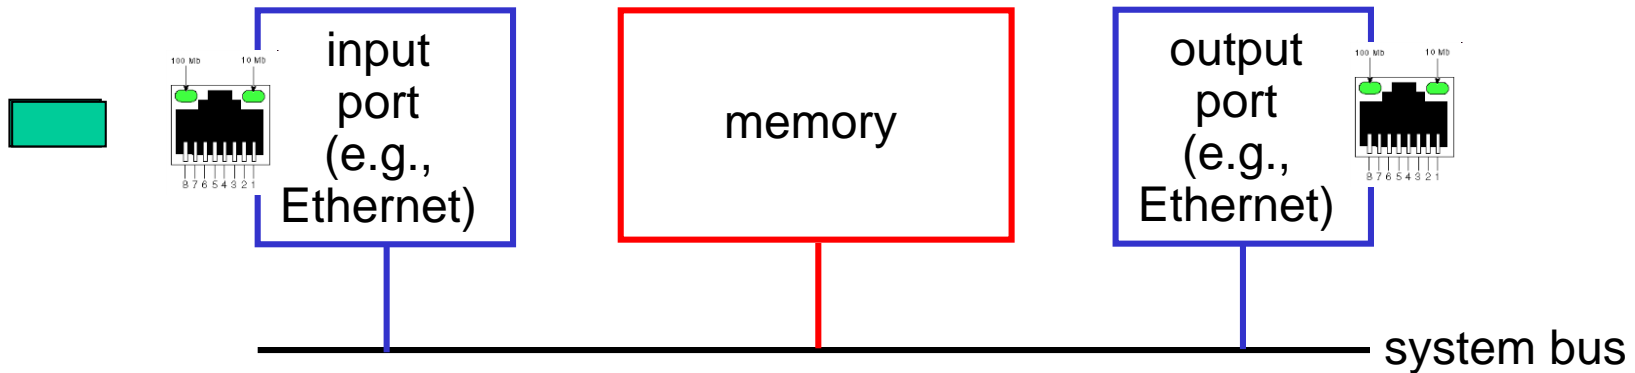
- Transfer packet from input buffer to appropriate output buffer
- Switching rate: rate at which packets can be transfer from inputs to outputs
  - often measured as multiple of input/output line rate
  - N inputs: switching rate N times line rate desirable
- Three types of switching fabrics



# Switching via memory

## *First generation routers:*

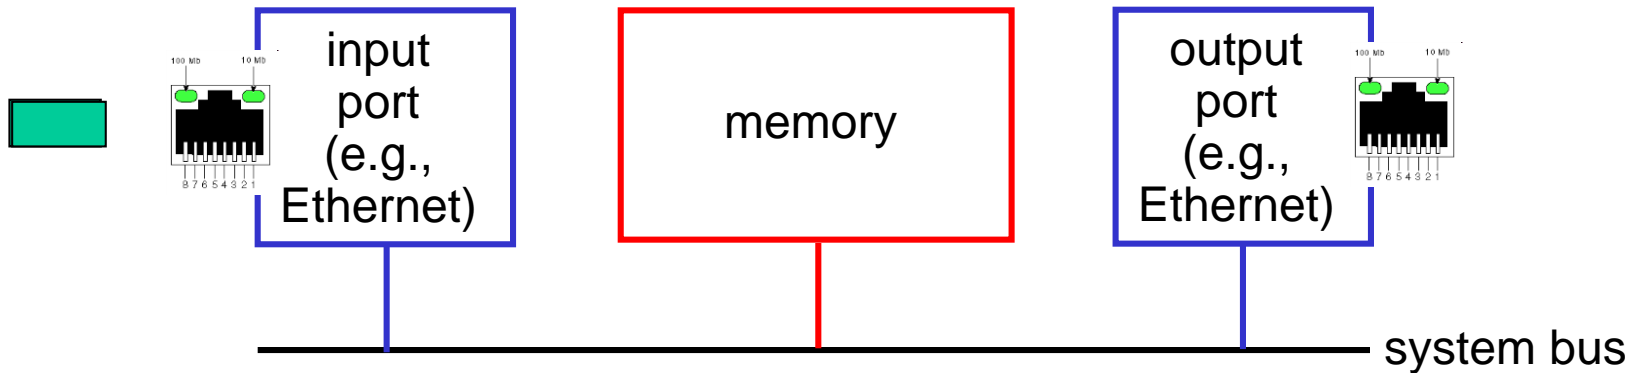
- Traditional computers with switching under direct control of CPU
- Packet copied to system's memory
- Speed limited by memory bandwidth (2 bus crossings per datagram)



# Switching via memory

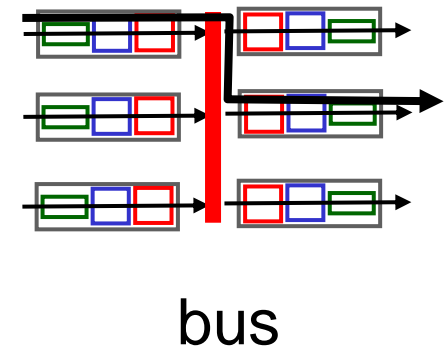
Advantages: ?

Disadvantages: ?



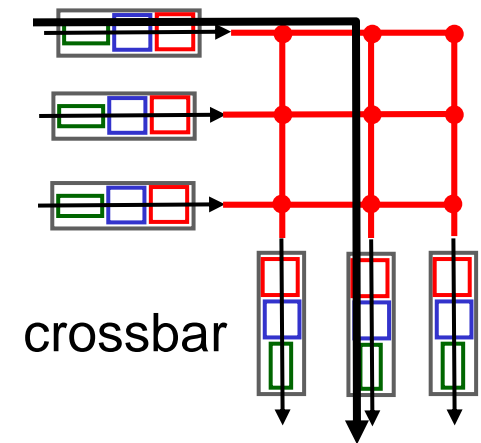
# Switching via a bus

- Datagram from input port memory to output port memory via a shared bus
- *Bus contention*: switching speed limited by bus bandwidth
- Advantages:
- Disadvantages:



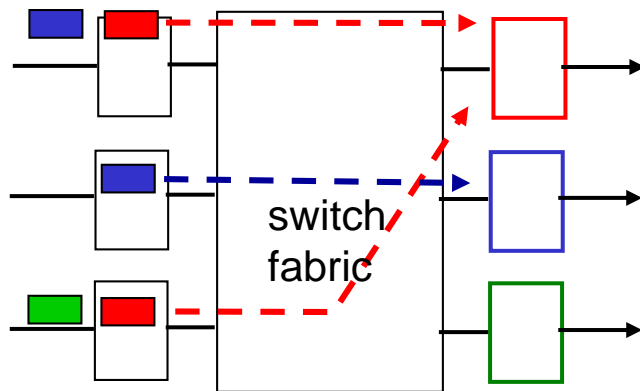
# Switching via interconnection network

- Overcome bus bandwidth limitations
- Banyan networks, crossbar, other interconnection nets initially developed to connect processors in multiprocessors
- Advanced design: fragmenting datagram into fixed length cells, switch cells through the fabric.

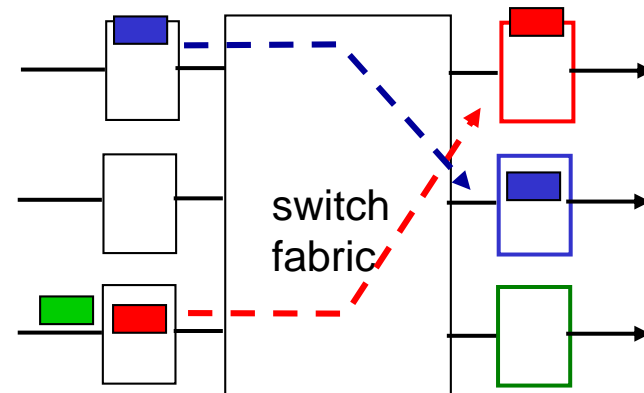


# Input port queuing

- Fabric slower than input ports combined → queueing may occur at input queues
  - *queueing delay and loss due to input buffer overflow!*
- **Head-of-the-Line (HOL) blocking:** queued datagram at front of queue prevents others in queue from moving forward

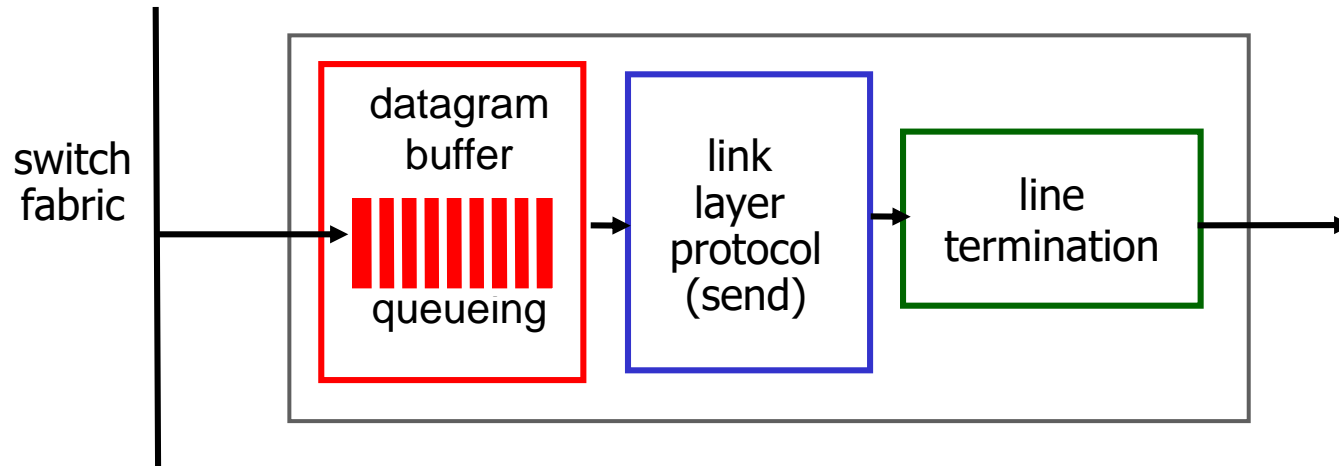


output port contention:  
only one red datagram can be  
transferred.  
*lower red packet is blocked*



one packet time later:  
green packet  
experiences HOL  
blocking

# Output ports



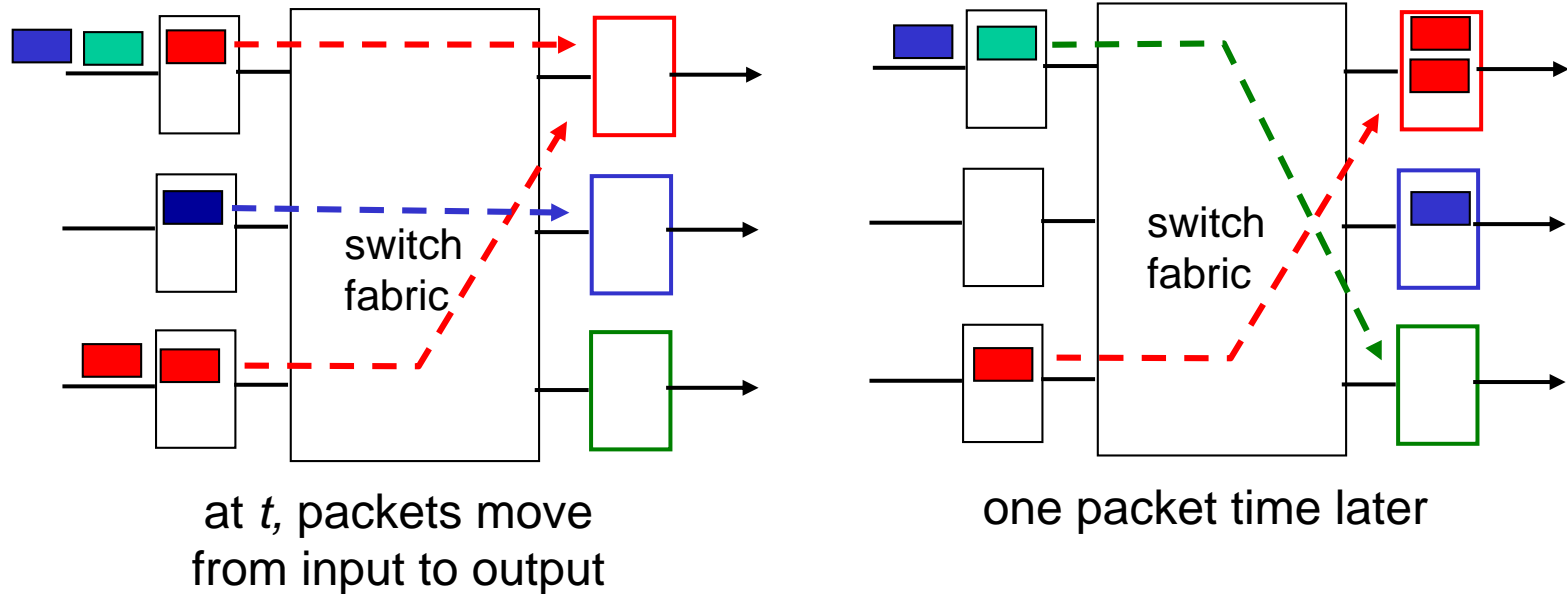
- *buffering* required when switch fabric faster than the line
- *scheduling discipline* chooses among queued datagrams for transmission

Datagram (packets) can be lost due to congestion, lack of buffers

Priority scheduling – who gets best performance, network neutrality



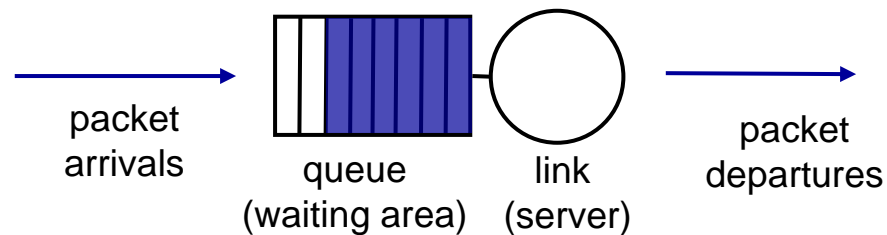
# Output port queueing



- buffering when arrival rate via switch exceeds output line speed
- *queueing (delay) and loss due to output port buffer overflow!*

# Scheduling mechanisms

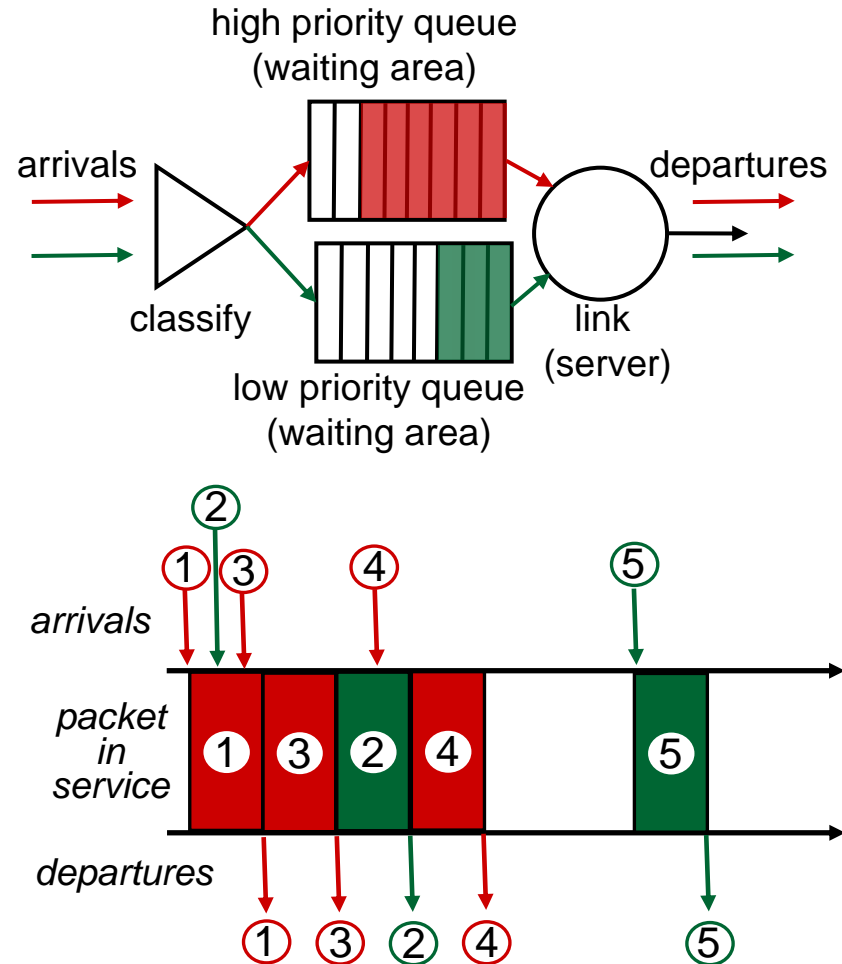
- *Scheduling*: choose next packet to send on link
- *FIFO (first in first out) scheduling*: send in order of arrival to queue
  - real-world example?
  - *discard policy*: if packet arrives to full queue: who to discard?
    - *tail drop*: drop arriving packet
    - *priority*: drop/remove on priority basis
    - *random*: drop/remove randomly



# Scheduling policies: priority

*priority scheduling*: send highest priority queued packet

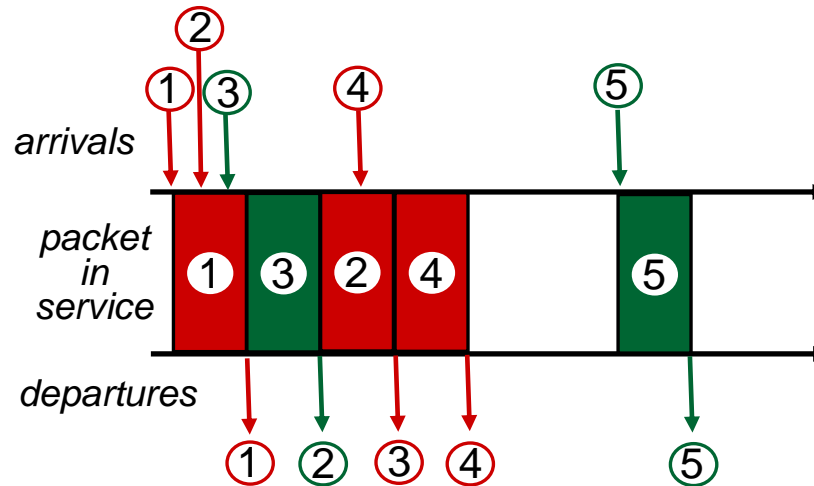
- multiple *classes*, with different priorities
  - class may depend on marking or other header info, e.g. IP source/dest, port numbers, etc.
  - real world example?



# Scheduling policies: still more

## *Round Robin (RR) scheduling:*

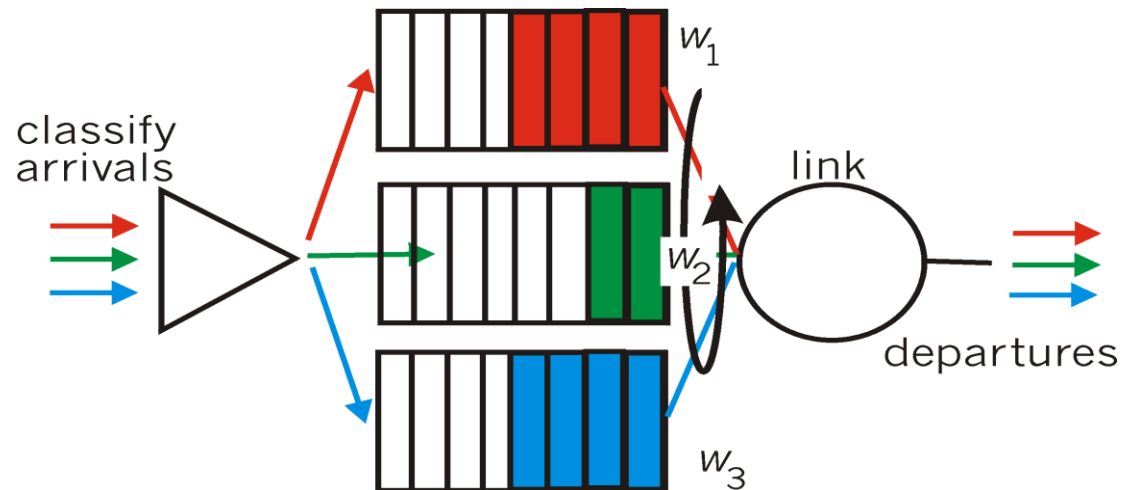
- multiple classes
- cyclically scan class queues, sending one complete packet from each class (if available)
- real world example?



# Scheduling policies: still more

## *Weighted Fair Queuing (WFQ):*

- generalized Round Robin
- each class gets weighted amount of service in each cycle

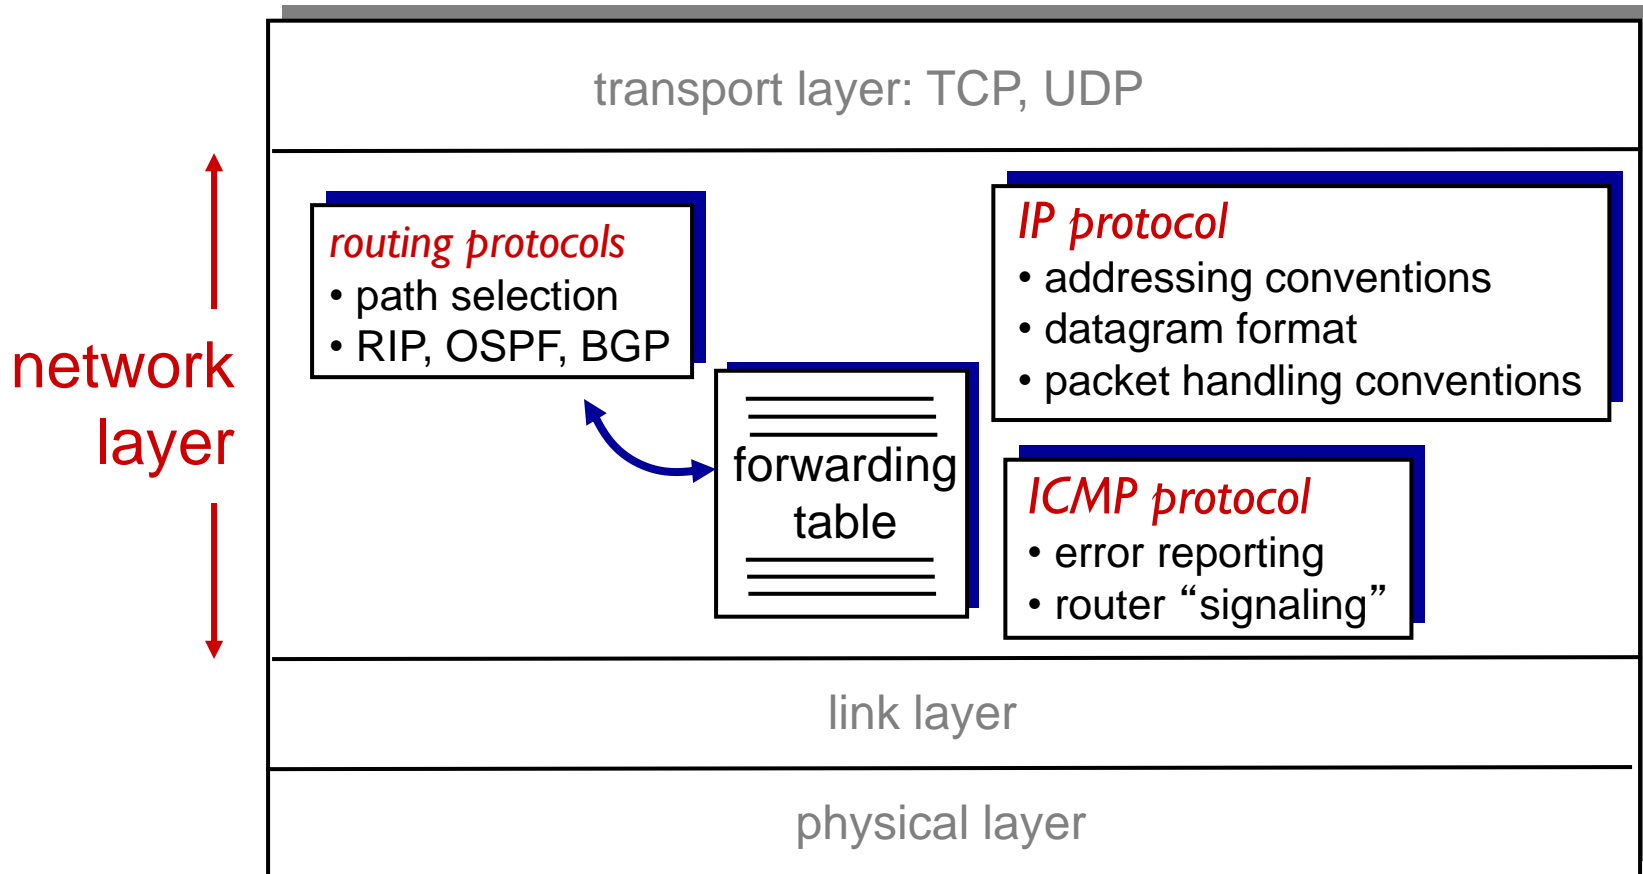


# Topics

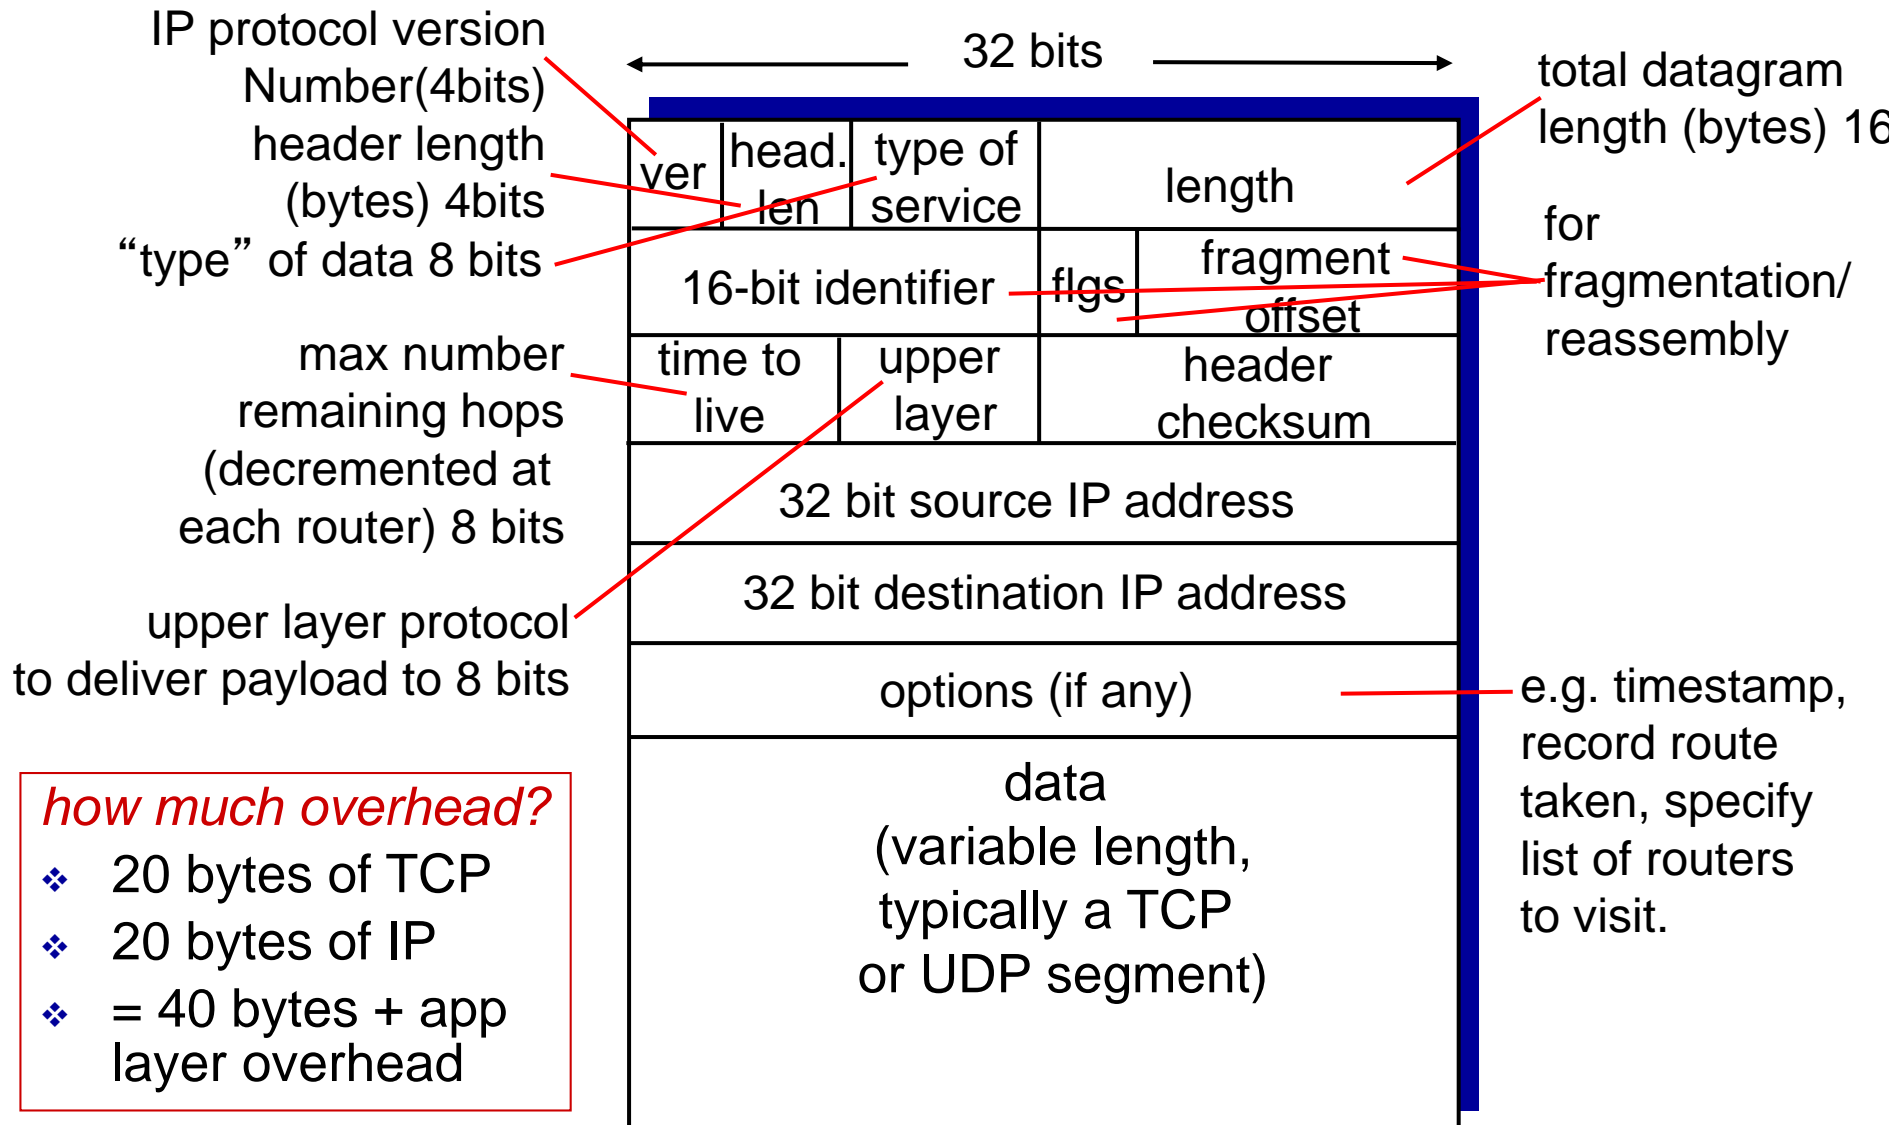
- Overview of Network layer
  - data plane
  - control plane
- What's inside a router
- IP: Internet Protocol
  - datagram format
  - fragmentation
  - IPv4 addressing
  - network address translation
  - IPv6

# The Internet network layer

host, router network layer functions:



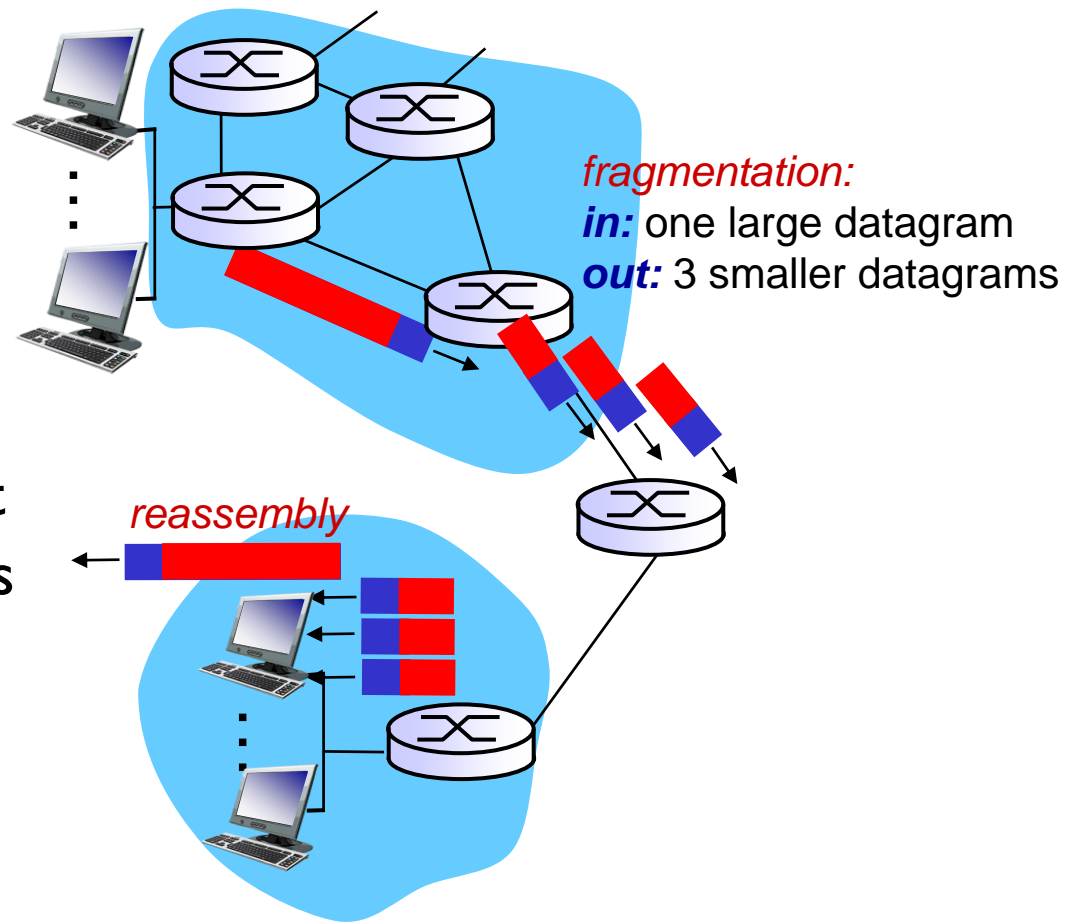
# IP datagram format





# IP fragmentation, reassembly

- network links have MTU (max.transfer size) - largest possible link-level frame
  - different link types, different MTUs
- large IP datagram divided (“fragmented”) within net
  - one datagram becomes several datagrams
  - **“reassembled” only at final destination**
  - IP header bits used to identify, order related fragments



# IP fragmentation, reassembly

## example:

- ❖ 4000 byte datagram
- ❖ MTU = 1500 bytes

	length	ID	fragflag	offset	
	=4000	=x	=0	=0	

*one large datagram becomes  
several smaller datagrams*

1480 bytes in  
data field

offset =  
 $1480/8$

	length	ID	fragflag	offset	
	=1500	=x	=1	=0	

	length	ID	fragflag	offset	
	=1500	=x	=1	=185	

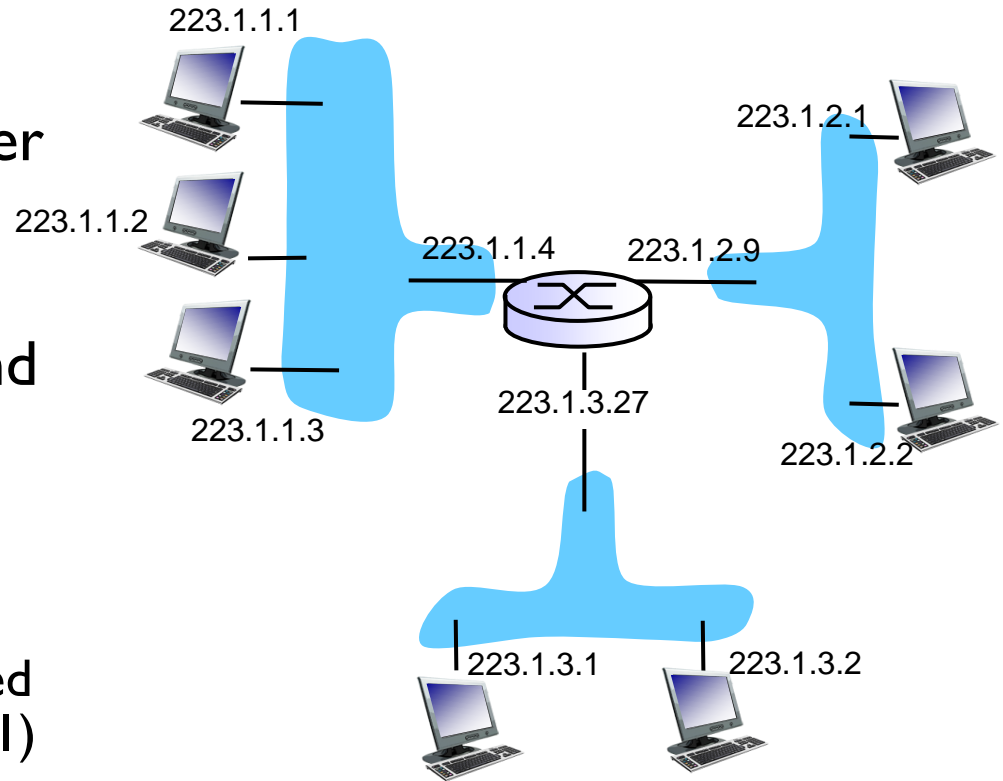
	length	ID	fragflag	offset	
	=1040	=x	=0	=370	

# Topics

- Overview of Network layer
  - data plane
  - control plane
- What's inside a router
- IP: Internet Protocol
  - datagram format
  - fragmentation
  - IPv4 addressing
  - network address translation
  - IPv6

# Recap: IP addressing

- **IP address:** 32-bit identifier for host, router interface
- **interface:** connection between host/router and physical link
  - router's typically have multiple interfaces
  - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)
- **IP addresses associated with each interface**

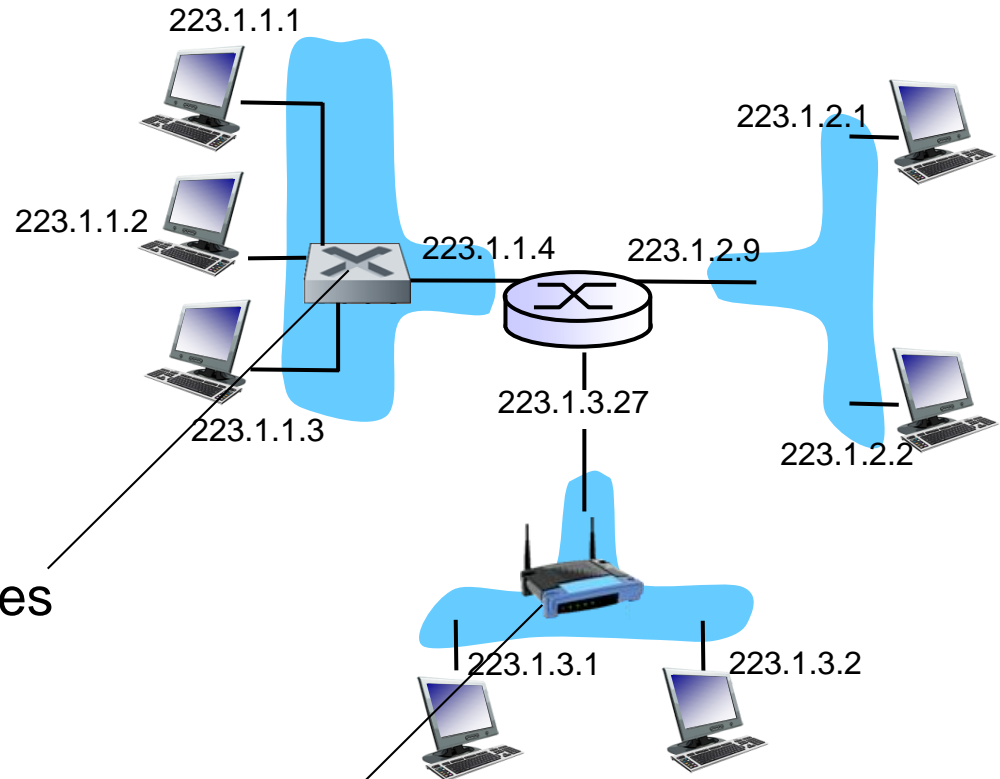


$$223.1.1.1 = \underbrace{11011111}_{223} \underbrace{00000001}_1 \underbrace{00000001}_1 \underbrace{00000001}_1$$

# IP addressing: Introduction

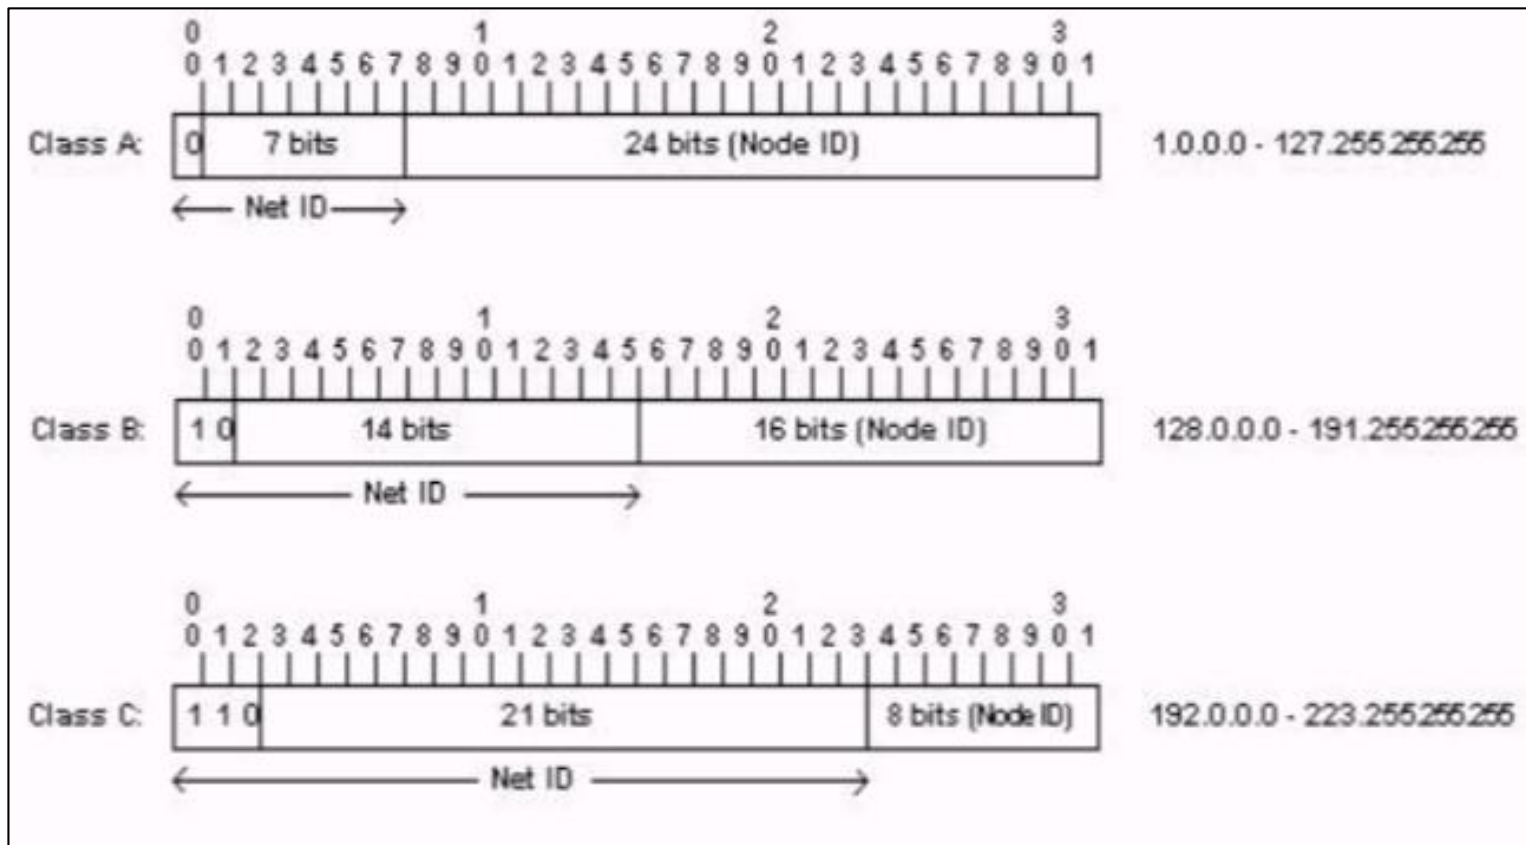
*Q: how are interfaces actually connected?*

*A:* wired Ethernet interfaces connected by Ethernet switches



*A:* wireless WiFi interfaces connected by WiFi base station

# IP Address Classes



# IP Address Classes

- Class A: The first octet is the network portion.
  - Octets 2, 3, and 4 are for subnets/hosts
- Class B: The first two octets are the network portion.
  - Octets 3 and 4 are for subnets/hosts
- Class C: The first three octets are the network portion.
  - Octet 4 is for subnets/hosts

# Private Address Range

Address Class	Reserved Address Space
Class A	10.0.0.0 - 10.255.255.255
Class B	172.16.0.0 - 172.31.255.255
Class C	192.168.0.0 - 192.168.255.255



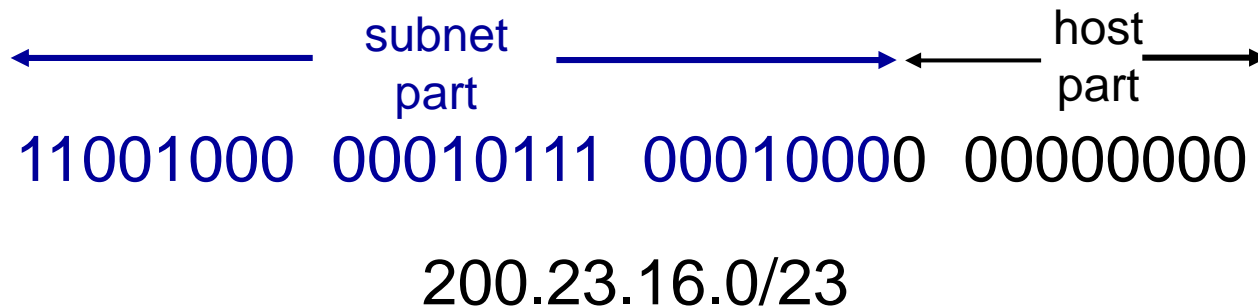
# Problem with IP Address Classes

- Inefficient usage of IP addresses!
- Solution: Subnets
  - Creates multiple logical networks that exist within a single Class A, B, or C network.
  - If you do not subnet, you will only be able to use one network from your Class A, B, or C network, which is unrealistic
  - Each LAN must have a unique network ID, with every node on that link being a member of the same network

# IP addressing: CIDR

## CIDR: Classless InterDomain Routing

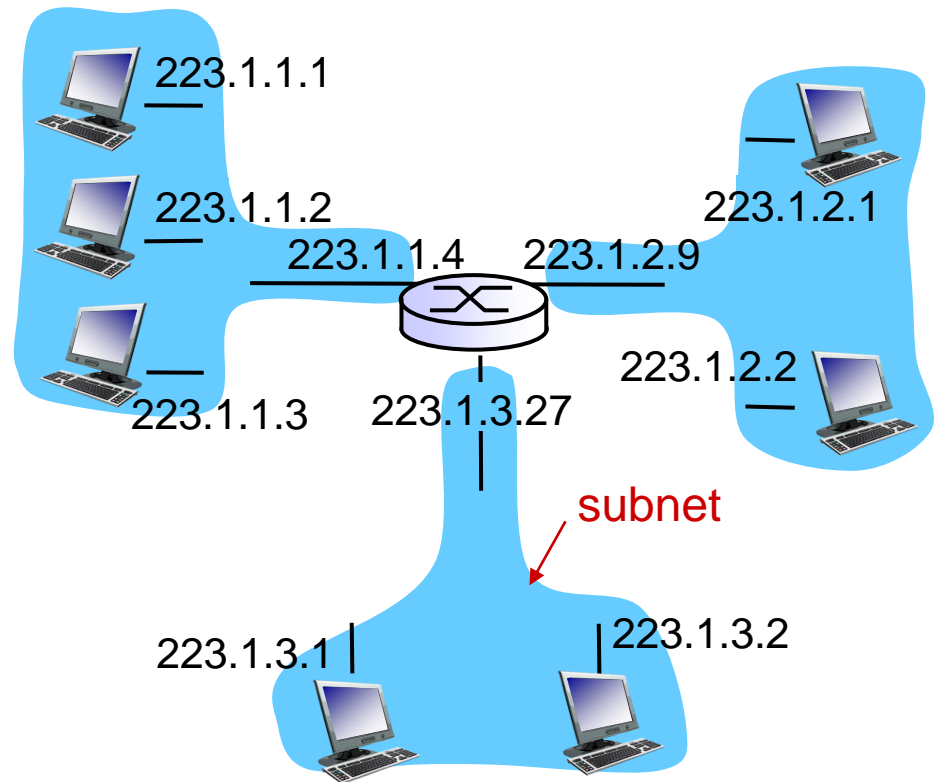
- subnet portion of address of arbitrary length
- address format: **a.b.c.d/x**, where x is # bits in subnet portion of address



# Subnets

## ■ What's a subnet ?

- device interfaces with same subnet part of IP address
- can physically reach each other *without intervening router*



network consisting of 3 subnets

# Subnetting Your Network

- Every IP address has a Subnet Mask
  - 172.16.25.2 255.255.0.0
    - subnet part - high order bits
    - host part - low order bits
- Classless Interdomain Routing(CIDR)
  - 172.16.25.2 /16

# Subnet Mask

- Determines the way an IP address is split into network and hosts portions
- **Class A** - 0nnnnnnnn.hhhhhhhh.hhhhhhhh.hhhhhhhh  
Subnet Mask = 255.0.0.0

# Subnet Mask

- Determines the way an IP address is split into network and hosts portions
- **Class A** - 0nnnnnnnn.hhhhhhhh.hhhhhhhh.hhhhhhhh  
Subnet Mask = 255.0.0.0 → IP Address /8

# Subnet Mask

- Determines the way an IP address is split into network and hosts portions
- **Class A** - 0nnnnnnnn.hhhhhhhh.hhhhhhhh.hhhhhhhh  
Subnet Mask = 255.0.0.0 → IP Address /8
- **Class B** - 10nnnnnnn.nnnnnnnnn.hhhhhhhh.hhhhhhhh  
Subnet Mask = 255.255.0.0 → IP Address /?

# Subnet Mask

- Determines the way an IP address is split into network and hosts portions
- **Class A** - 0nnnnnnnn.hhhhhhhh.hhhhhhhh.hhhhhhhh  
Subnet Mask = 255.0.0.0 → IP Address /8
- **Class B** - 10nnnnnnn.nnnnnnnn.hhhhhhhh.hhhhhhhh  
Subnet Mask = 255.255.0.0 → IP Address /16



# Subnet Mask

- Determines the way an IP address is split into network and hosts portions

- **Class A** - 0nnnnnnnn.hhhhhhhh.hhhhhhhh.hhhhhhhh

Subnet Mask = 255.0.0.0 → IP Address /8

- **Class B** - 10nnnnnnn.nnnnnnnnn.hhhhhhhh.hhhhhhhh

Subnet Mask = 255.255.0.0 → IP Address /16

- **Class C** - 100nnnnnn.nnnnnnnnn.nnnnnnnnn.hhhhhhhh

Subnet Mask = 255.255.255.0    IP Address /?

# Subnet Mask

- Determines the way an IP address is split into network and hosts portions

- **Class A** - 0nnnnnnnn.hhhhhhhh.hhhhhhhh.hhhhhhhh

Subnet Mask = 255.0.0.0 → IP Address /8

- **Class B** - 10nnnnnnn.nnnnnnnnn.hhhhhhhh.hhhhhhhh

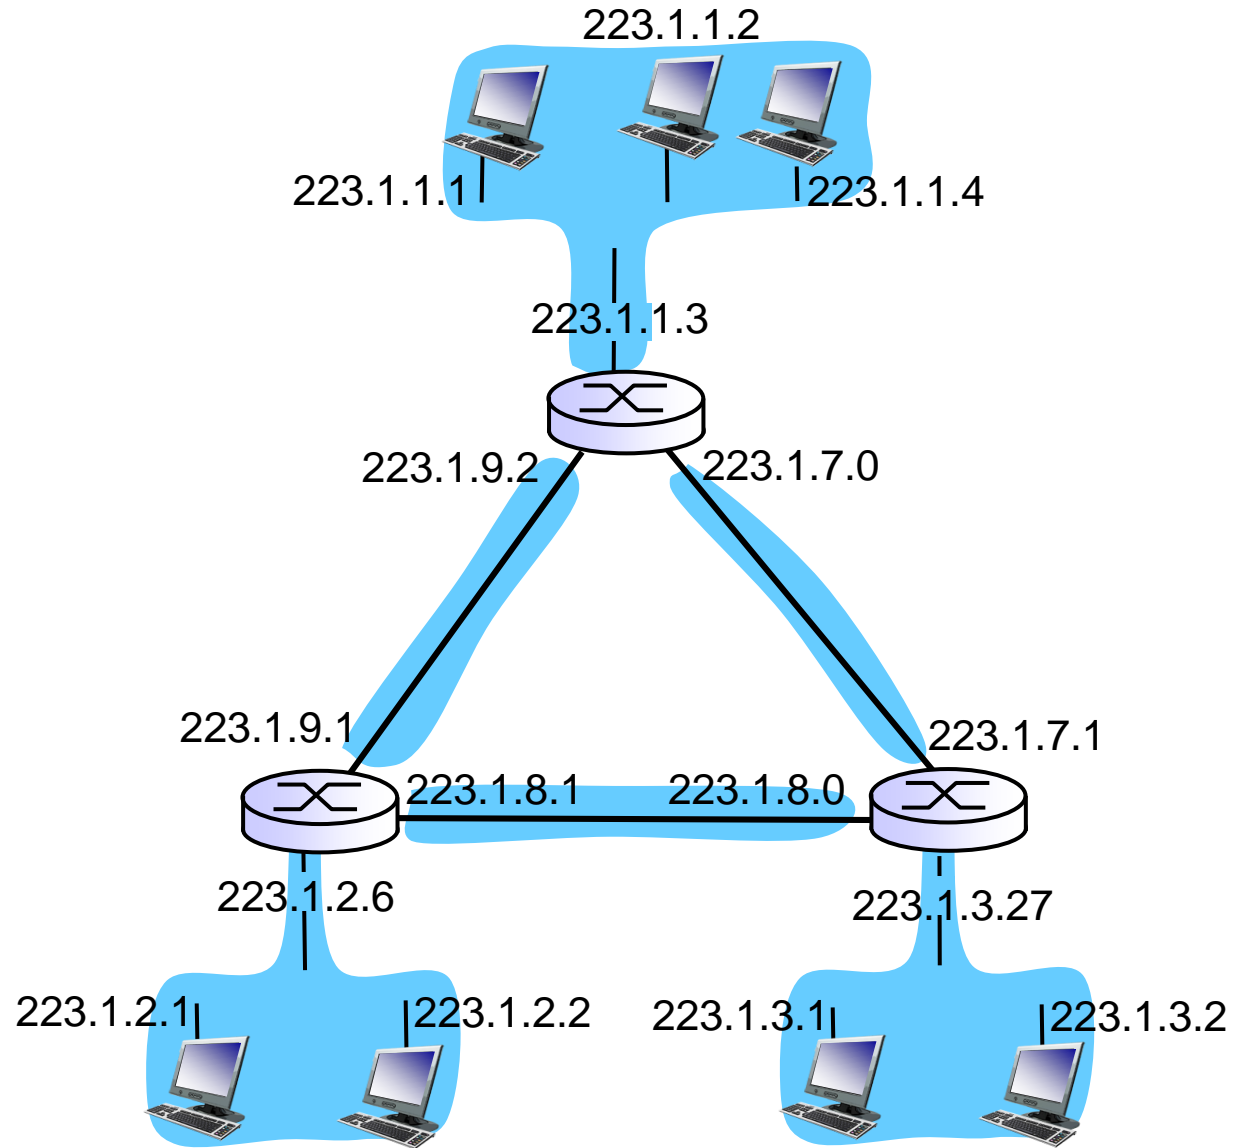
Subnet Mask = 255.255.0.0 → IP Address /16

- **Class C** - 100nnnnnn.nnnnnnnnn.nnnnnnnnn.hhhhhhhh

Subnet Mask = 255.255.255.0    IP Address /24

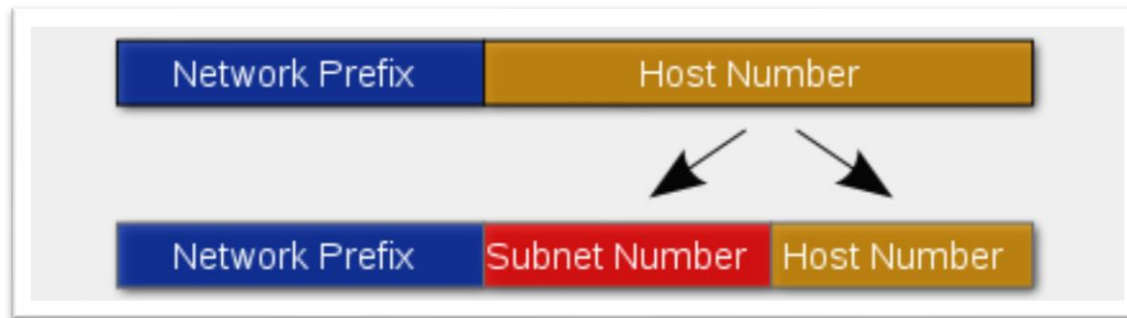
# Subnets

how many?



# Subnetted Networks

- The network portion of the address is extended by splitting up the host number




- Borrowing 1 or more bits from the host bit portion

# Example:

Dividing a network into 2 subnets requires to borrow 1 bit

Class C:

11111111.11111111.11111111.10000000(255.255.255.128)  
CIDR  IP address /25

This would allow 126 hosts per subnet

All 1's are reserved for broadcast ID

All 0's are reserved for network ID

# Class C Subnetting

# of Subnets	# of Hosts/Subnet	NetMask	4 <sup>th</sup> Octet	CIDR Notation
2				
4				
8				
16				
32				
64				

# Class C Subnetting

# of Subnets	# of Hosts/Subnet	NetMask	4 <sup>th</sup> Octet	CIDR Notation
2	126	255.255.255.128	10000000	/25
4				
8				
16				
32				
64				

# Class C Subnetting

# of Subnets	# of Hosts/Subnet	NetMask	4 <sup>th</sup> Octet	CIDR Notation
2	126	255.255.255.128	10000000	/25
4	62	255.255.255.192	11000000	/26
8				
16				
32				
64				



# Class C Subnetting

# of Subnets	# of Hosts/Subnet	NetMask	4 <sup>th</sup> Octet	CIDR Notation
2	126	255.255.255.128	10000000	/25
4	62	255.255.255.192	11000000	/26
8	30	255.255.255.224	11100000	/27
16				
32				
64				

# Class C Subnetting

# of Subnets	# of Hosts/Subnet	NetMask	4 <sup>th</sup> Octet	CIDR Notation
2	126	255.255.255.128	10000000	/25
4	62	255.255.255.192	11000000	/26
8	30	255.255.255.224	11100000	/27
16	14	255.255.255.240	11110000	/28
32				
64				

# Class C Subnetting

# of Subnets	# of Hosts/Subnet	NetMask	4 <sup>th</sup> Octet	CIDR Notation
2	126	255.255.255.128	10000000	/25
4	62	255.255.255.192	11000000	/26
8	30	255.255.255.224	11100000	/27
16	14	255.255.255.240	11110000	/28
32	6	255.255.255.248	11111000	/29
64				

# Class C Subnetting

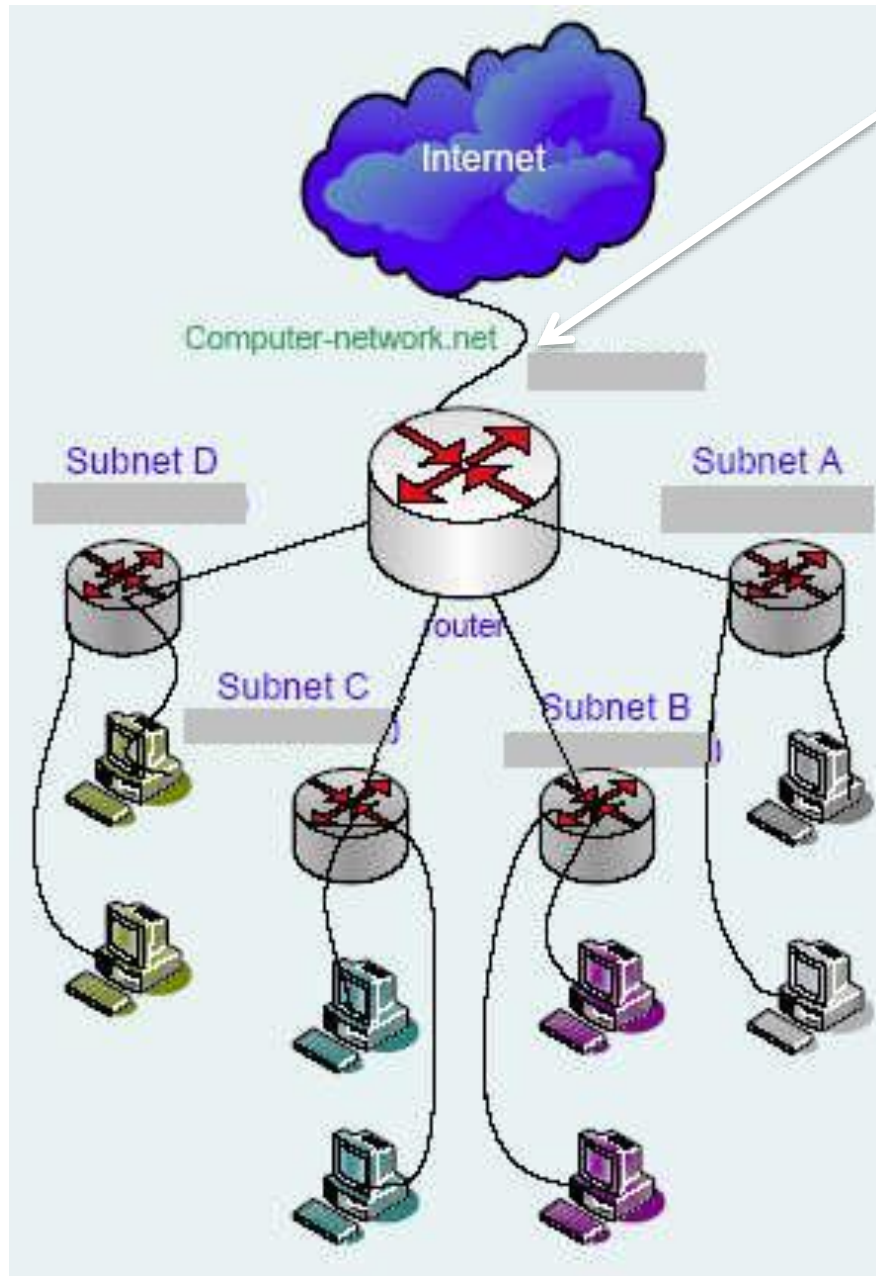
# of Subnets	# of Hosts/Subnet	NetMask	4 <sup>th</sup> Octet	CIDR Notation
2	126	255.255.255.128	10000000	/25
4	62	255.255.255.192	11000000	/26
8	30	255.255.255.224	11100000	/27
16	14	255.255.255.240	11110000	/28
32	6	255.255.255.248	11111000	/29
64	2	255.255.255.252	11111100	/30

192.168.5.130 /24

Subnet mask = 255.255.255.0

192.168.5.0 = Network ID

Create 4 Subnets. How?

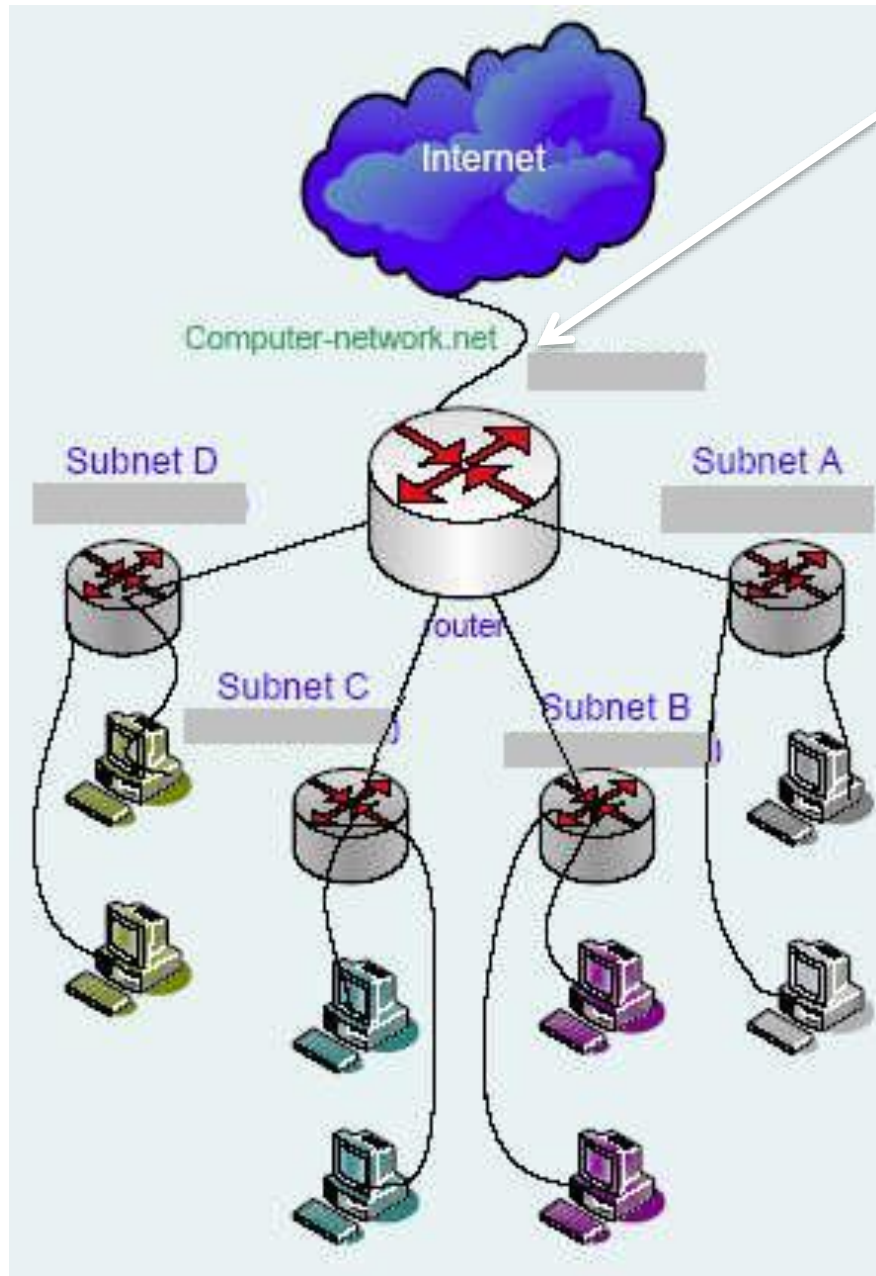


192.168.5.130 /24

Subnet mask = 255.255.255.0

192.168.5.0 = Network ID

Create 4 Subnets. How?



192.168.5.0 =>

11000000.10101000.00000101.00000000

Borrow 2 bits from host byte

11000000.10101000.00000101.**11**000000

Subnet Mask = 255.255.255.192

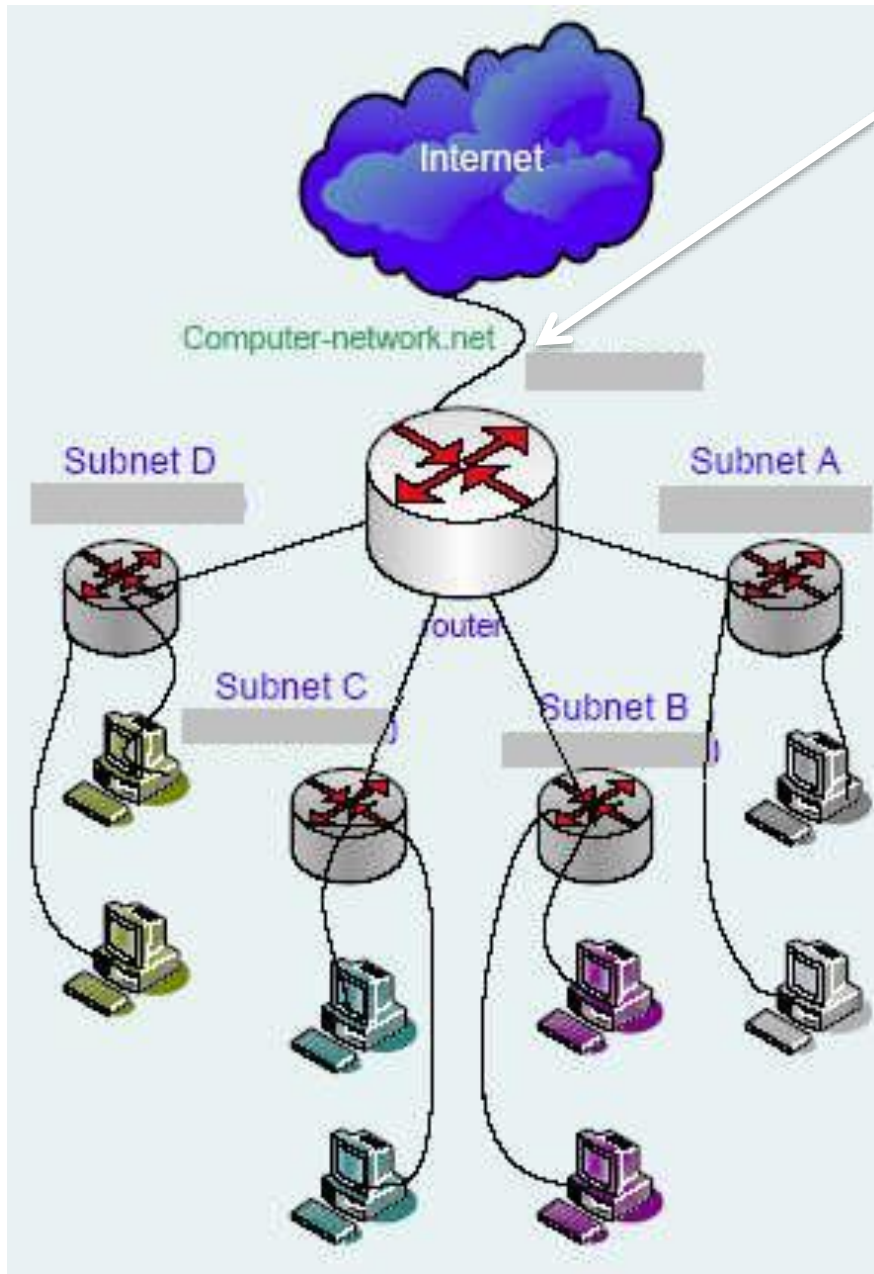
= /26

192.168.5.130 /24

Subnet mask = 255.255.255.0

192.168.5.0 = Network ID

Create 4 Subnets. How?



192.168.5.0 =>

11000000.10101000.00000101.00000000

Borrow 2 bits from host byte

11000000.10101000.00000101.11000000

Subnet Mask = 255.255.255.192

= /26

What are my new subnets?

192.168.5.130 /24

Subnet mask = 255.255.255.0

192.168.5.0 = Network ID

Create 4 Subnets. How?

192.168.5.0 =>

11000000.10101000.00000101.00000000

Borrow 2 bits from host byte

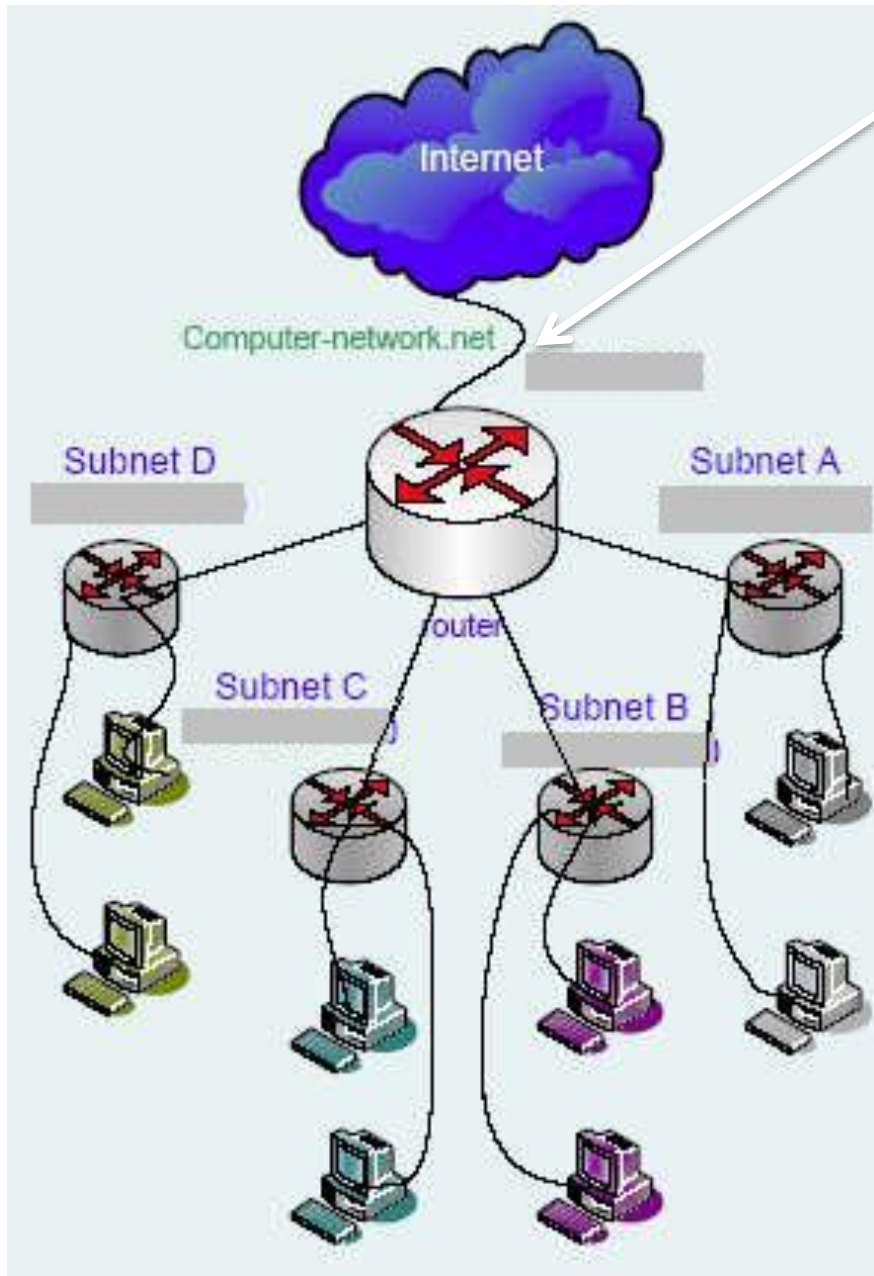
11000000.10101000.00000101.11000000

Subnet Mask = 255.255.255.192

= /26

Subnet A -> 192.168.5.1/26

to 192.168.5.62/26





192.168.5.130 /24

Subnet mask = 255.255.255.0

192.168.5.0 = Network ID

Create 4 Subnets. How?

192.168.5.0 =>

11000000.10101000.00000101.00000000

Borrow 2 bits from host byte

11000000.10101000.00000101.11000000

Subnet Mask = 255.255.255.192

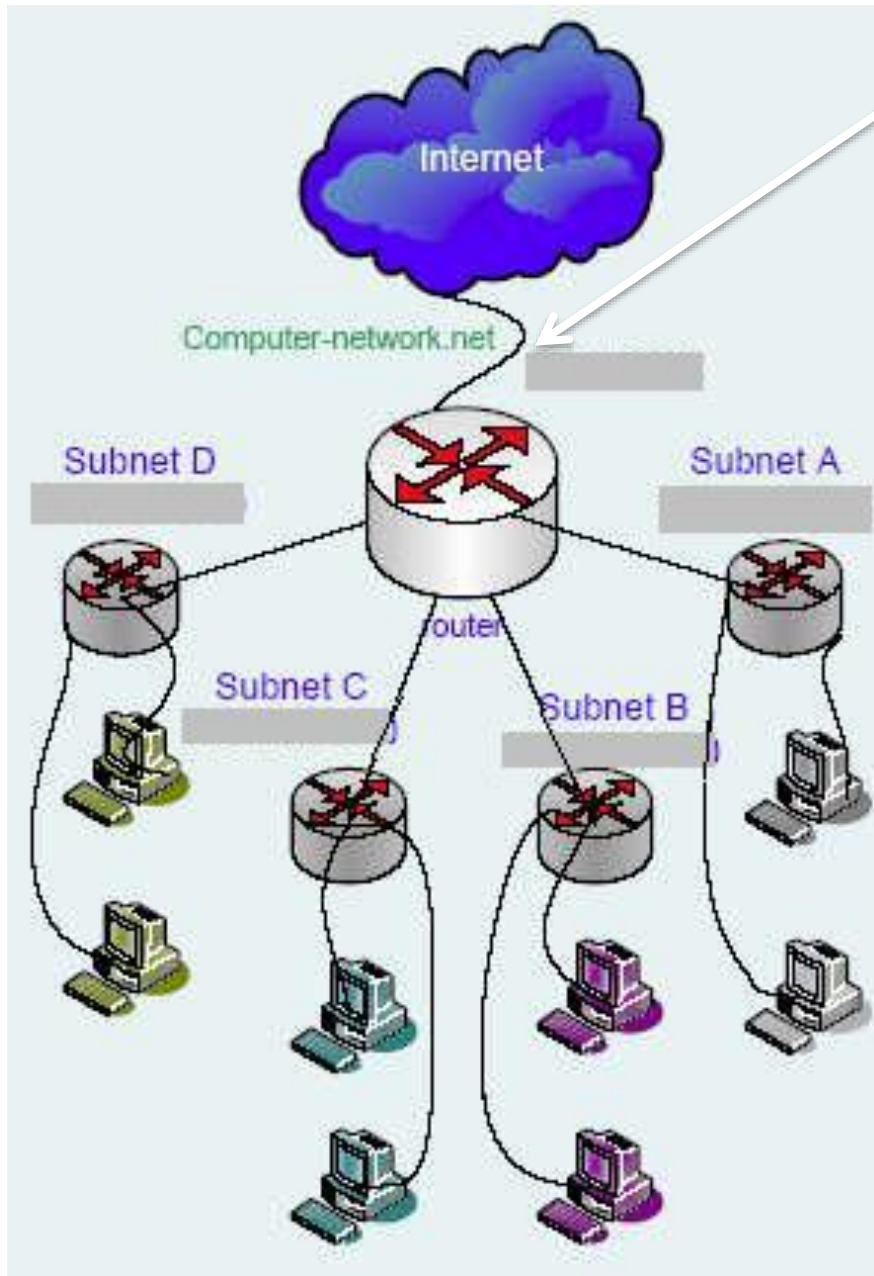
= /26

Subnet A -> 192.168.5.1/26

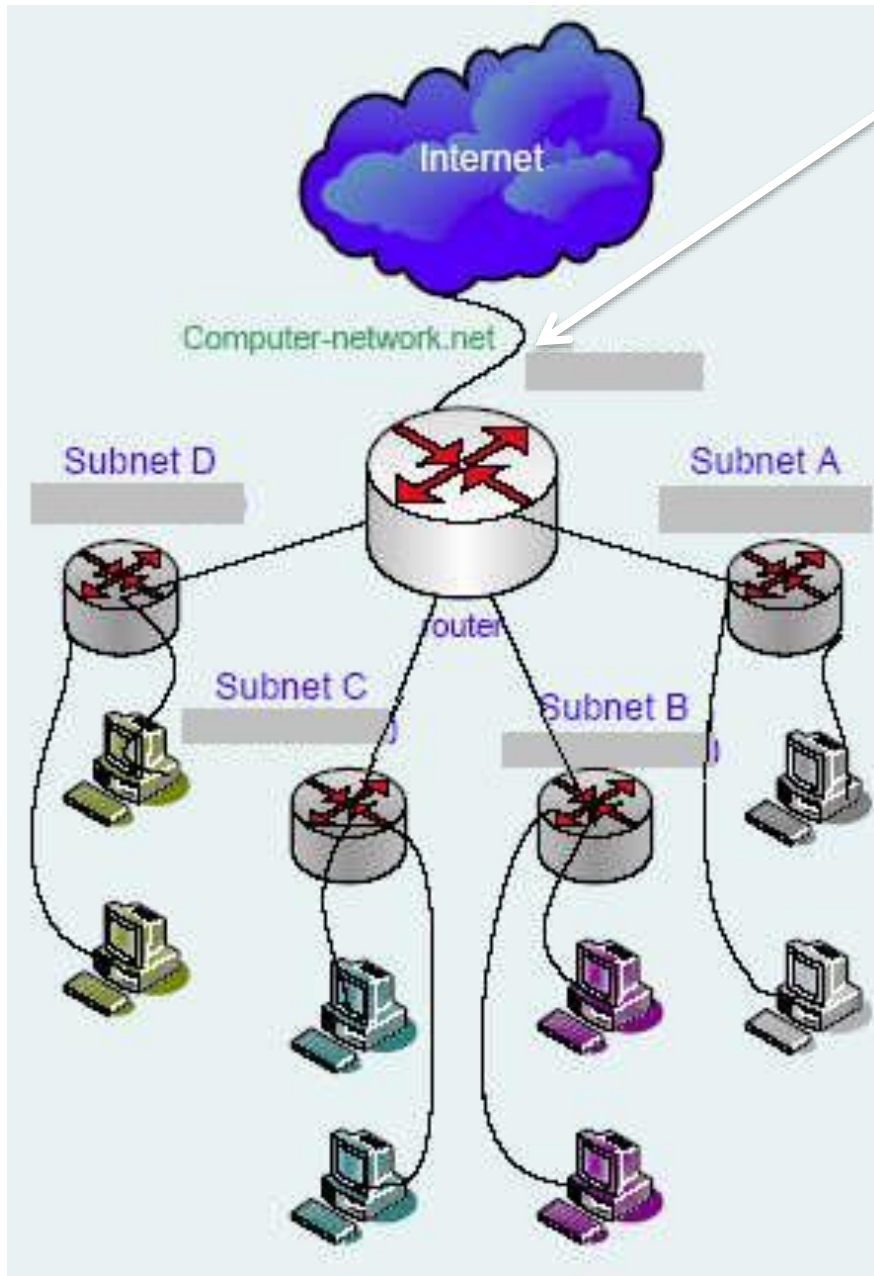
to 192.168.5.62/26

Subnet B -> 192.168.5.65/26

to 192.168.5.126/26







192.168.5.130 /24

Subnet mask = 255.255.255.0

192.168.5.0 = Network ID

Create 4 Subnets. How?

192.168.5.0 =>

11000000.10101000.00000101.00000000

Borrow 2 bits from host byte

11000000.10101000.00000101.11000000

Subnet Mask = 255.255.255.192

= /26

Subnet A -> 192.168.5.1/26

to 192.168.5.62/26

Subnet B -> 192.168.5.65/26

to 192.168.5.126/26

Subnet C -> 192.168.5.129/26

to 192.168.5.190/26

Subnet D -> 192.168.5.193/26

to 192.168.5.254/26

# How to create subnets

- Determine the number of required network IDs:
  - One for each subnet
  - One for each wide area network connection
- Determine the number of required host IDs per subnet:
  - One for each TCP/IP host
  - One for each router interface
- Based on the above requirements, create the following:
  - One subnet mask for your entire network
  - A unique subnet ID for each physical segment
  - A range of host IDs for each subnet

# Advantages

- Allows a single shared network address to split it up into many smaller networks.
- Without subnets, organizations would require many network addresses
  - Limited number of Network addresses available
- Alleviates traffic
  - Smaller routing tables
  - Alleviates excessive packet collision and congestion
- Easier to manage and solve problems
- Better Security
  - Separating departments with highly sensitive material
    - Accounting and Administration

# Disadvantages

- Doesn't allocate IP address proportionately per subnet
- Limited by the number of IP address
- Need to buy hardware such as routers

# IP addresses: how to get one?

**Q:** How does a *host* get IP address?

# IP addresses: how to get one?

**Q:** How does a *host* get IP address?

- hard-coded by system admin in a file
  - Windows: control-panel->network->configuration->tcp/ip->properties
  - UNIX: /etc/rc.config
- **DHCP: Dynamic Host Configuration Protocol:** dynamically get address from as server
  - “plug-and-play”

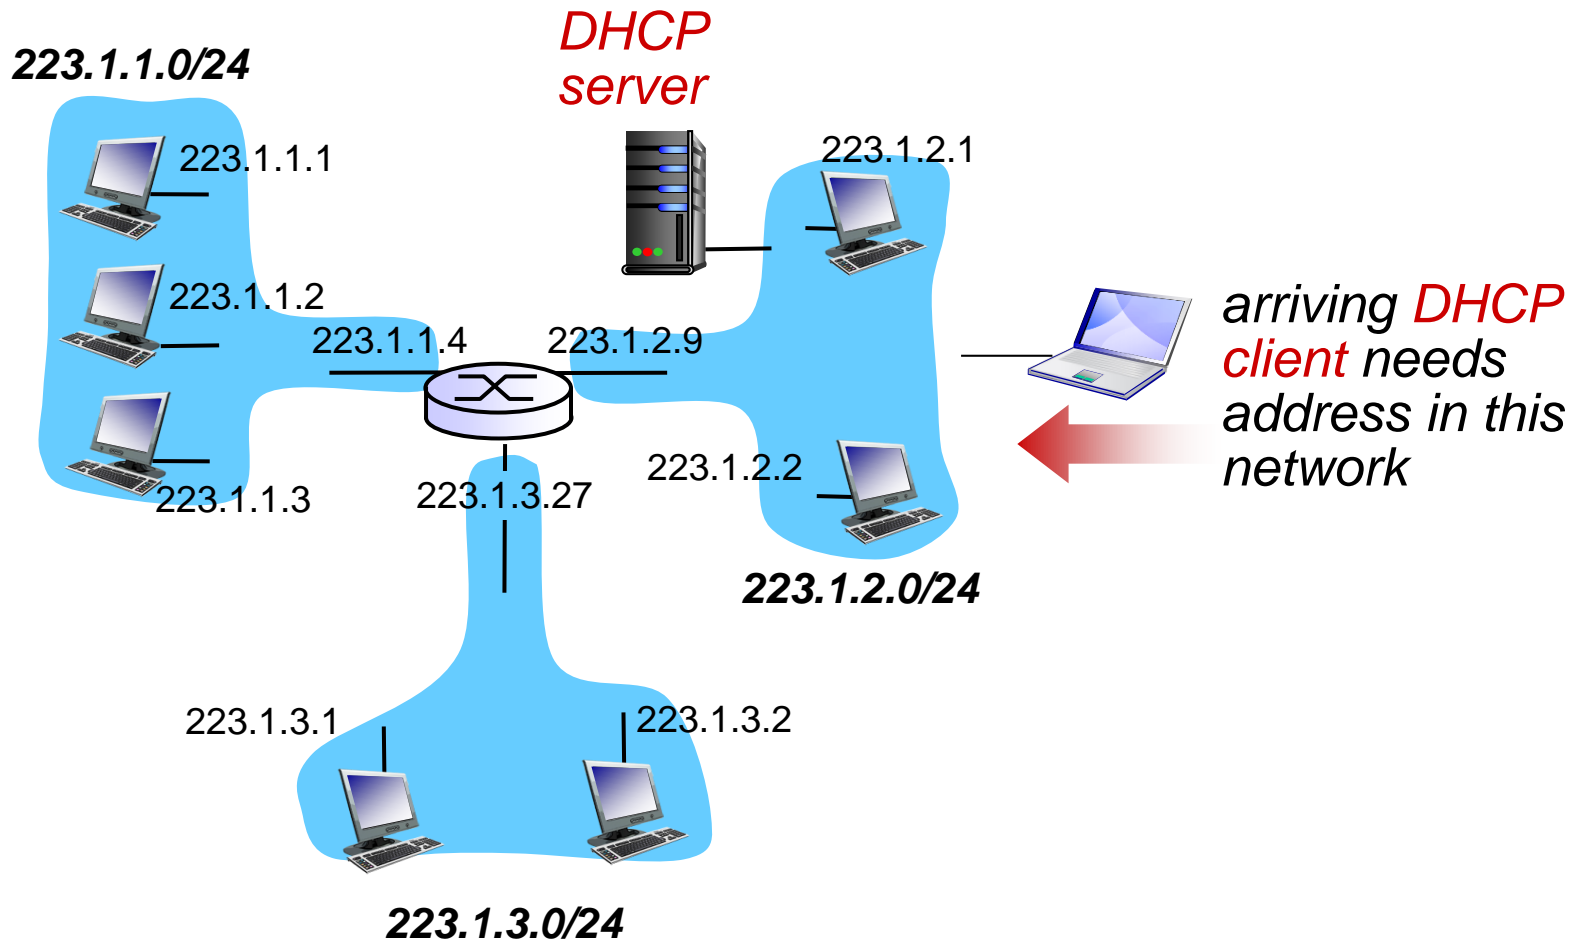


# DHCP: Dynamic Host Configuration Protocol

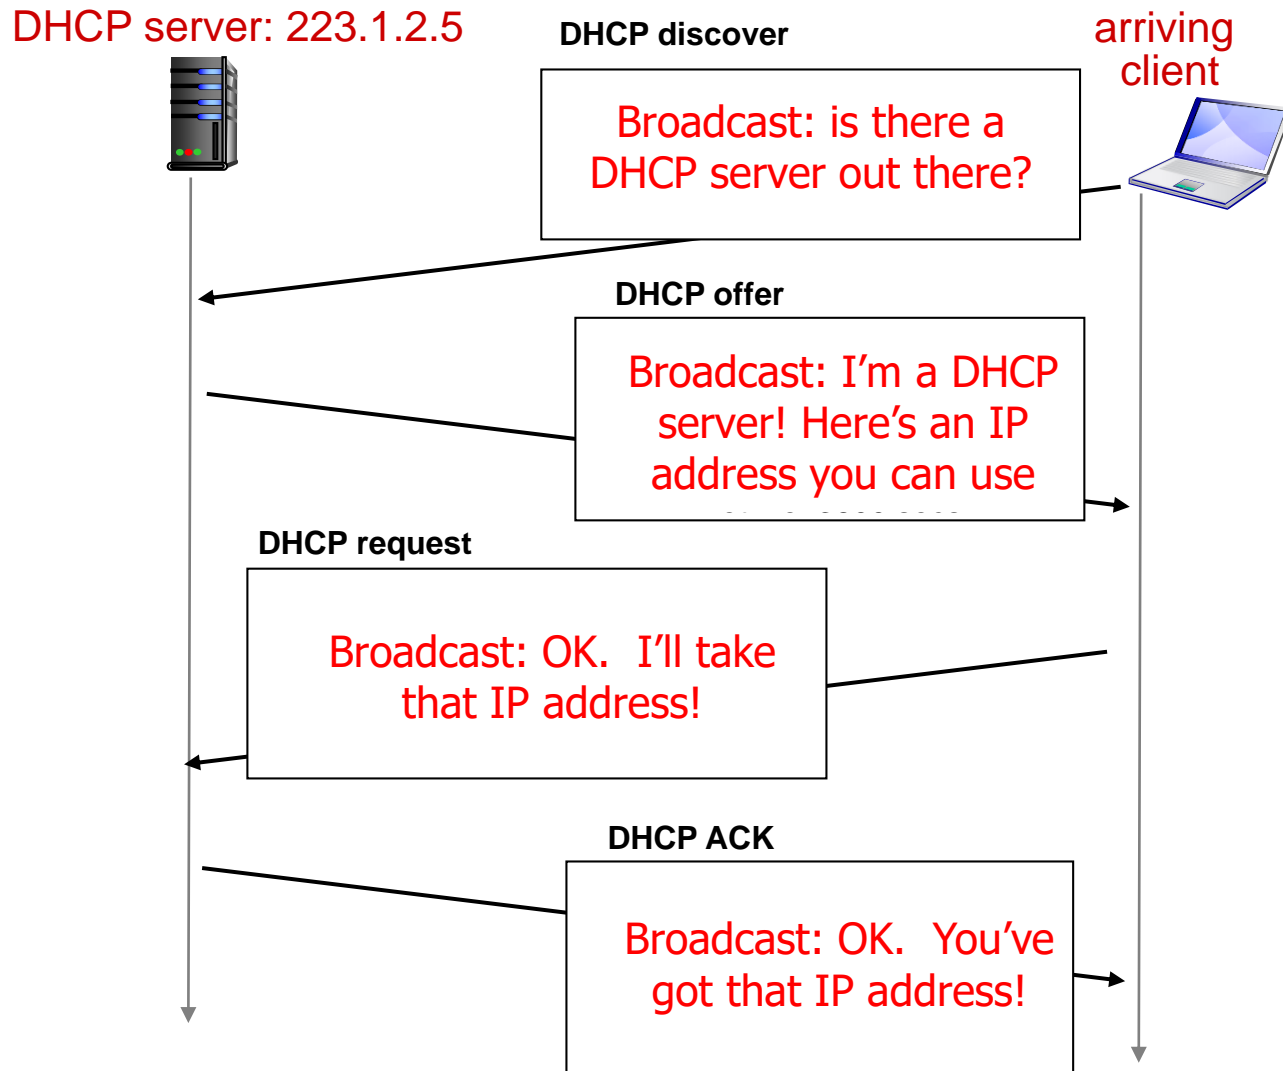
**Goal:** allow host to *dynamically* obtain its IP address from network server when it joins network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/“on”)
- support for mobile users who want to join network (more shortly)

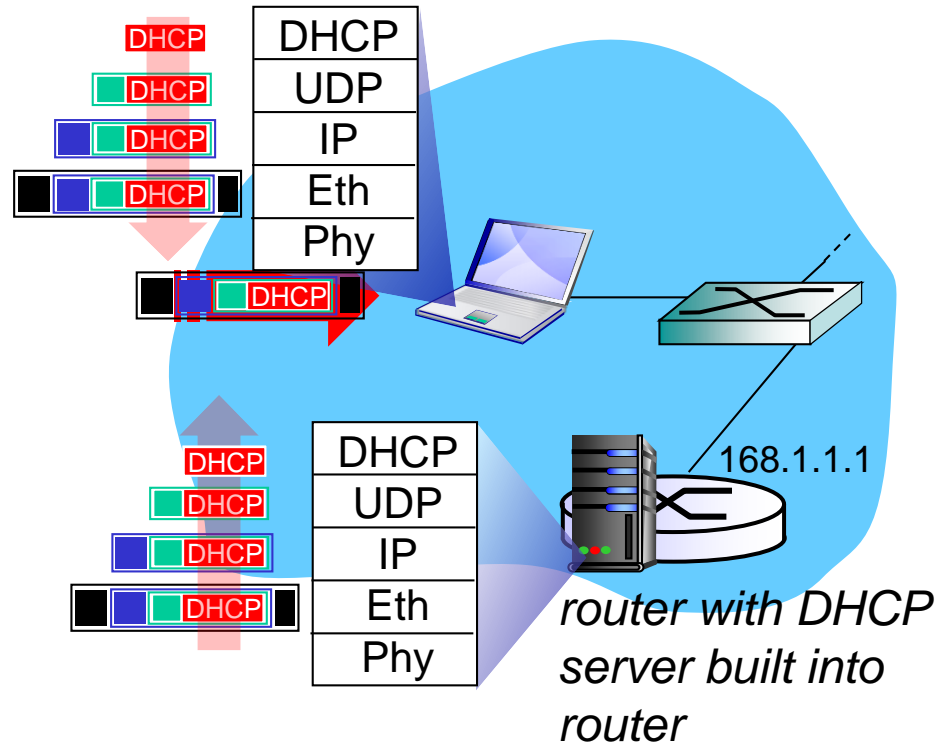
# DHCP client-server scenario



# DHCP client-server scenario

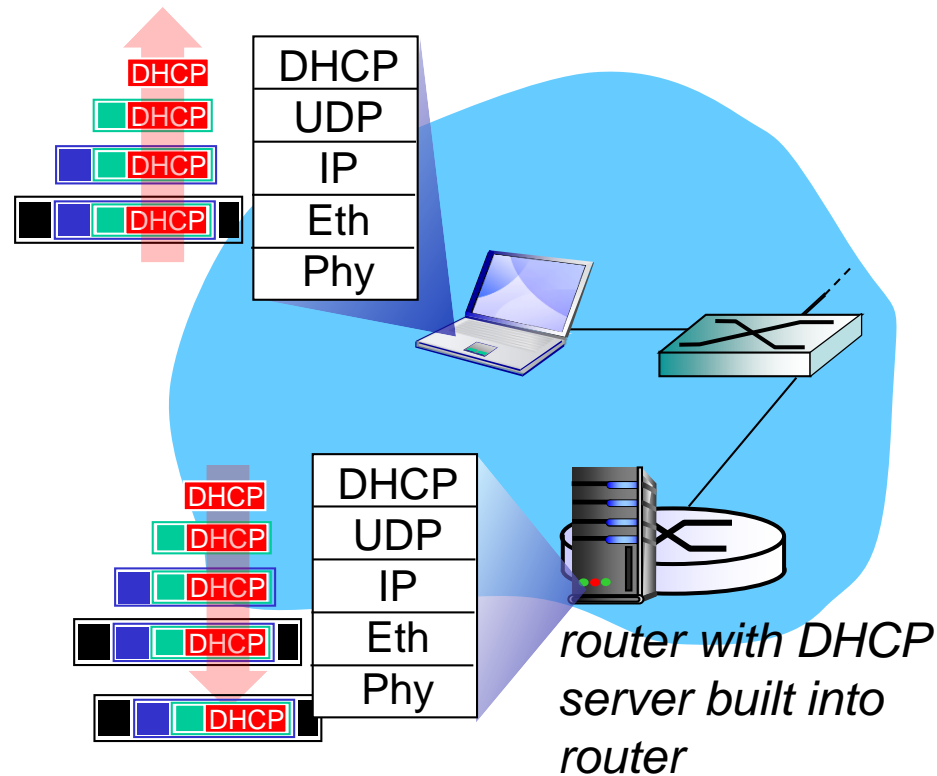


# DHCP: example



- connecting laptop needs its IP address, addr of first-hop router, addr of DNS server: use DHCP
- DHCP request encapsulated in UDP, encapsulated in IP.
- Ethernet frame broadcast (dest: FFFFFFFF) on LAN, received at router running DHCP server

# DHCP: example



- DHCP server formulates DHCP ACK containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
- encapsulation of DHCP server, frame forwarded to client, demuxing up to DHCP at client
- client now knows its IP address, name and IP address of DNS server, IP address of its first-hop router

# DHCP Wireshark exercise

- Refer to the document on Blackboard

# DHCP: more than IP addresses

DHCP can return more than just allocated IP address on subnet:

- address of first-hop router for client
- name and IP address of DNS sever
- network mask (indicating network versus host portion of address)

# IP addresses: how to get one?

**Q:** how does *network* get subnet part of IP addr?

**A:** gets allocated portion of its provider ISP' s address space

ISP's block	<u>11001000 00010111 00010000</u> 00000000	200.23.16.0/20
Organization 0	<u>11001000 00010111 00010000</u> 00000000	200.23.16.0/23
Organization 1	<u>11001000 00010111 00010010</u> 00000000	200.23.18.0/23
Organization 2	<u>11001000 00010111 00010100</u> 00000000	200.23.20.0/23
...	.....	....
Organization 7	<u>11001000 00010111 00011110</u> 00000000	200.23.30.0/23



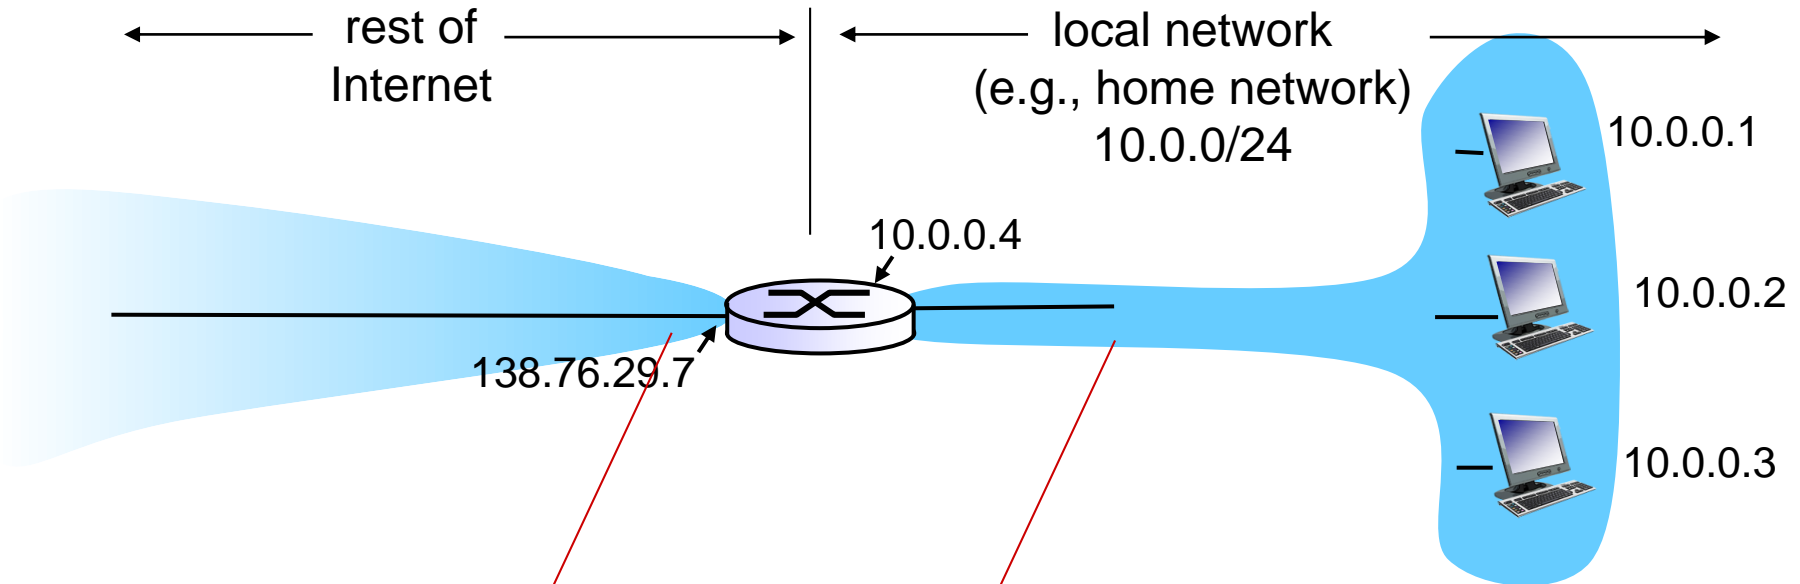
# IP addressing: the last word...

**Q:** how does an ISP get block of addresses?

**A: ICANN:** Internet Corporation for Assigned Names and Numbers <http://www.icann.org/>

- allocates addresses
- manages DNS
- assigns domain names, resolves disputes

# NAT: network address translation



*all* datagrams *leaving* local network have *same* single source NAT IP address: 138.76.29.7, different source port numbers

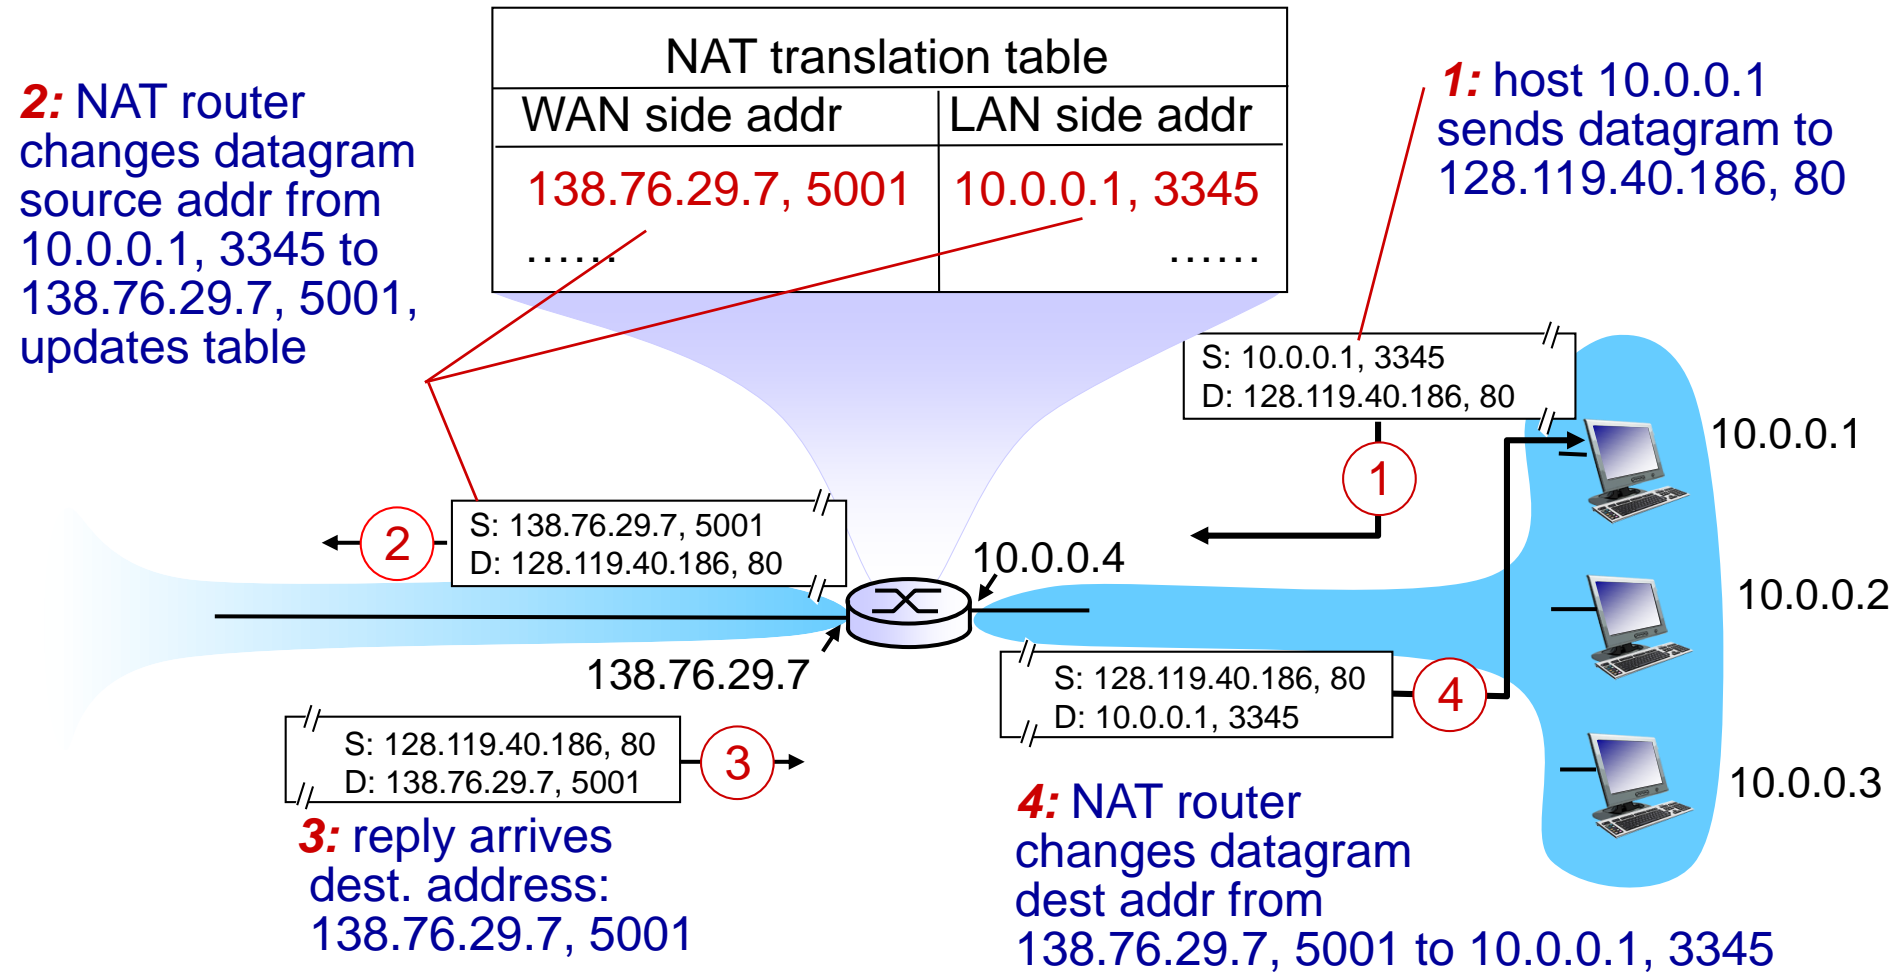
datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)

# NAT: network address translation

*Motivation:* local network uses just one IP address as far as outside world is concerned:

- range of addresses not needed from ISP: just one IP address for all devices
- can change addresses of devices in local network without notifying outside world
- can change ISP without changing addresses of devices in local network
- devices inside local net not explicitly addressable, visible by outside world (a security plus)

# NAT: network address translation



# NAT: network address translation

*Implementation:* NAT router must:

- *Outgoing datagrams: replace* (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #) . . . remote clients/servers will respond using (NAT IP address, new port #) as destination addr
- *Remember (in NAT translation table)* every (source IP address, port #) to (NAT IP address, new port #) translation pair
- *Incoming datagrams: replace* (NAT IP address, new port #) in dest fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

# NAT: network address translation

- 16-bit port-number field:
  - 60,000 simultaneous connections with a single LAN-side address!
- NAT is controversial:
  - routers should only process up to layer 3
  - address shortage should be solved by IPv6
  - violates end-to-end argument
    - NAT possibility must be taken into account by app designers, e.g., P2P applications
  - NAT traversal: what if client wants to connect to server behind NAT?

# Topics

- Overview of Network layer
  - data plane
  - control plane
- What's inside a router
- IP: Internet Protocol
  - datagram format
  - fragmentation
  - IPv4 addressing
  - network address translation
  - IPv6

# IPv6: motivation

- *Initial motivation:* 32-bit address space soon to be completely allocated.
- additional motivation:
  - header format helps speed processing/forwarding
  - header changes to facilitate QoS
- *IPv6 datagram format:*
  - fixed-length 40 byte header
  - **no fragmentation allowed**



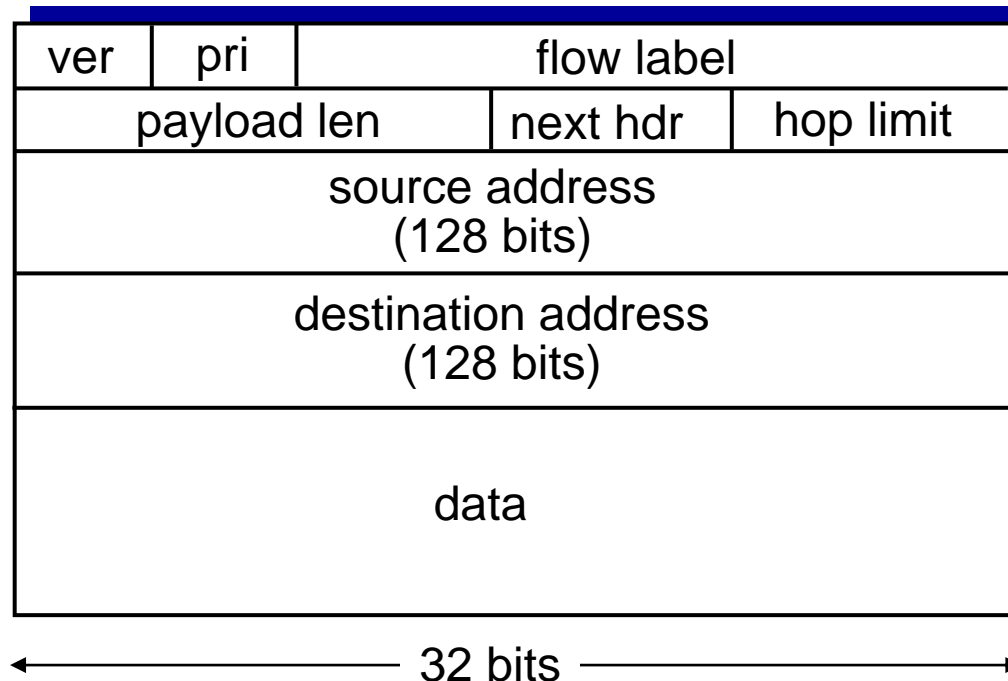
# IPv6 datagram format

*priority:* identify priority among datagrams in flow

*flow Label:* identify datagrams in same “flow.”

(concept of “flow” not well defined).

*next header:* identify upper layer protocol for data

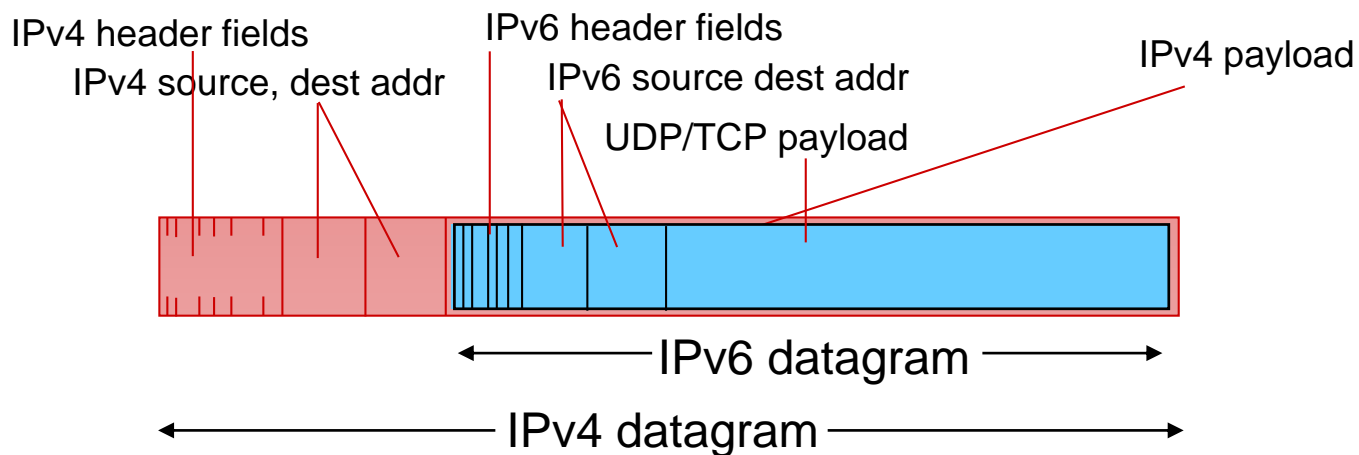


# Other changes from IPv4

- *checksum*: removed entirely to reduce processing time at each hop
- *options*: allowed, but outside of header, indicated by “Next Header” field
- *ICMPv6*: new version of ICMP
  - additional message types, e.g. “Packet Too Big”
  - multicast group management functions

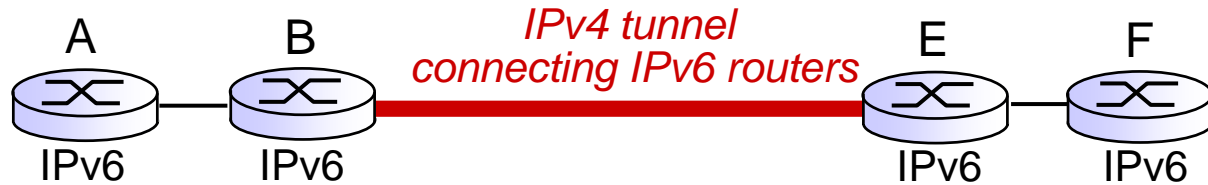
# Transition from IPv4 to IPv6

- not all routers can be upgraded simultaneously
  - no “flag days”
  - how will network operate with mixed IPv4 and IPv6 routers?
- **tunneling**: IPv6 datagram carried as *payload* in IPv4 datagram among IPv4 routers

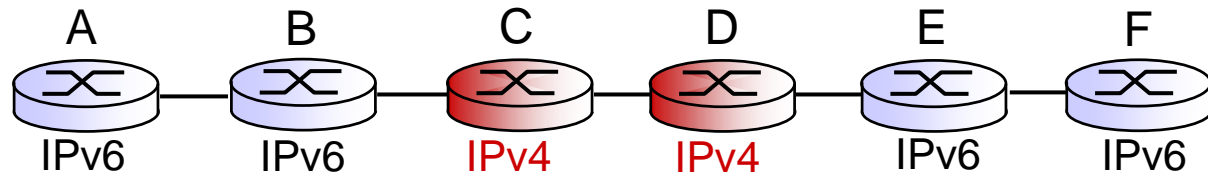


# Tunneling

logical view:

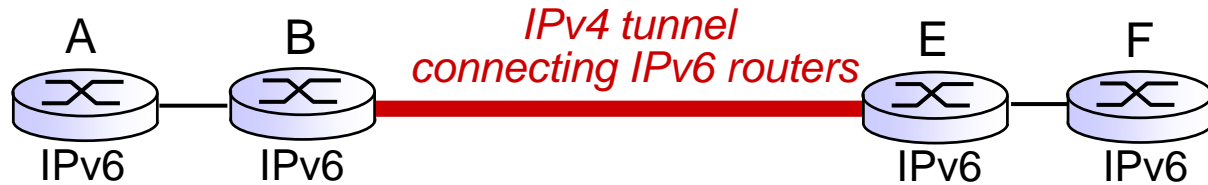


physical view:

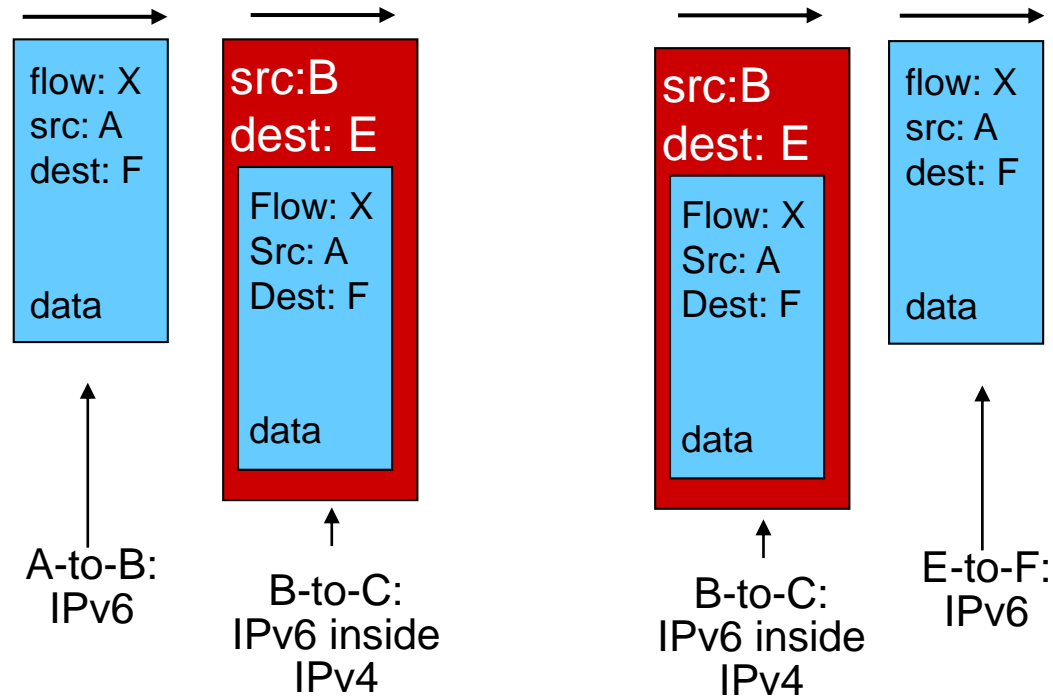
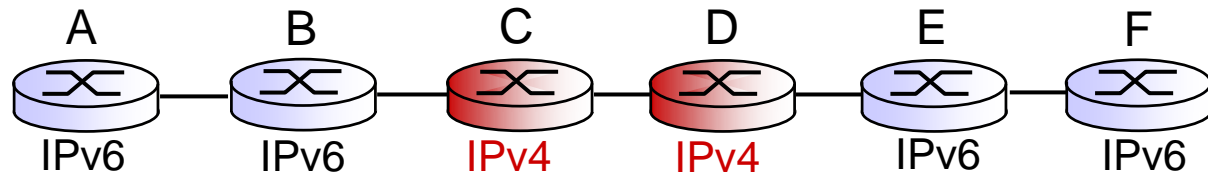


# Tunneling

logical view:



physical view:



# IPv6: adoption

- Carrier networks and ISPs with mobile networks leading the charge.
  - T-Mobile USA has more than 90% of its traffic going over IPv6
  - Verizon Wireless close behind at 82.25%.
  - Comcast and AT&T have its networks at 63% and 65%, respectively
- *Long (long!) time for deployment, use*
  - 20 years and counting!
  - *Why?*